

Almost all zeros of the Riemann zeta function are simple and on the critical line*

Hongyi Yang[†] Shihua Yang[‡]

August 23, 2026

Abstract

Let $N(T, 2T)$ count the nontrivial zeros of $\zeta(s)$ with ordinates in $(T, 2T]$, and N_0^s, N_d those simple and on the critical line, respectively distinct. Subject to the verification grading stated in the title note and in §14, we establish, unconditionally and with ineffective constants,

$$\lim_{T \rightarrow \infty} \frac{N_0^s}{N} = \lim_{T \rightarrow \infty} \frac{N_d}{N} = 1 :$$

almost all zeros of ζ are simple and lie on the critical line, in the density sense; the same holds for any fixed primitive Dirichlet L -function (Theorem 12.7). The proof runs a moment tower to its supremum: the trace-moment sequence of the Gabor-compressed Weil matrix of [5] is evaluated in exact rational arithmetic through order fourteen, is determinate (moment growth at most $(b!)^2$, Stieltjes–Carleman), and its limiting spectral measure has no atom at

*Verification status (read this first). *Unconditionality*: every statement in this paper is unconditional in the number-theoretic sense — no unproved hypothesis about ζ or any L -function enters at any tier (the arithmetic inputs are Montgomery’s band-limited pair-correlation theorem, Siegel–Walfisz, and the averaged shifted-convolution theorem of Matomäki–Radziwiłł–Tao; constants are ineffective). The grading below concerns the *verification provenance of six rational constants*, not any mathematical hypothesis. The grading is stated per constant in §14 and §10. The $k \leq 8$ subtower (headline 0.8493/0.9246) is at the grade of the predecessor chain: every consumed constant proved in exact rational arithmetic by two independent methods. The $k = 10, 12, 14$ rungs (headlines 0.8771/0.9385, 0.8964/0.9482, **0.9104/0.9552**) rest on the Toeplitz–trace representation of §8.6(vii), whose foundations are now theorems (§10): the model-transfer lemma, the branch-equality theorem ($m_b(N)N^{b+1}$ is a single polynomial, every $N \geq 1$) and the parity theorem ($L(-N) = (-1)^{b+1}L(N)$) together determine each new moment as an exact rational from finitely many exact finite- N evaluations — themselves proved rationals, confirmed bitwise on pre-registered held-out points, by two independent implementations, and by three independent 10^7 -sample measurement tiers. The consumed constants of every rung through $k = 14$ are thereby at proof grade, and the density-one capstone (Theorem 12.7) consumes only proved soft inputs — determinacy, the exact logarithmic law, and the order-uniform devices. The chain grade of the paper as a whole is *certified-candidate*, not record: each connected constant is proved RATIONAL (a finite signed sum of rational polytope volumes, §8.6); the six consumed constants of the $k \leq 6$ tower are *proved as exact rationals* by symbolic evaluation of those very volumes ($C_5 = \frac{1}{36}$, $\{4, 2\} = -\frac{23}{420}$, $\{6\} = -\frac{1}{126}$, $\{2, 2, 2\} = \frac{32}{105}$, $\{5, 2\} = \frac{1}{8}$, $\{2^4\} = \frac{1661}{3780}$, $C_7 = -\frac{17}{360}$, $C_8 = \frac{157}{4032}$, §8.5, §9; two independent exact methods where stated); every consumed constant now carries two independent confirmations (§14(d)) — the Lean formalization covers the identity and certificate layers only (kernel-compiled, core Lean 4, zero `sorry`; §14(f)), the analytic layer is unformalized, and no external review has taken place. The three write-out items W1–W3 are written out (§8.4), the pairing-layer constants are certified as exact rationals (§8.3, Appendix C), and the consumption step is backed by an exact rational dual certificate valid on unbounded support (Lemma 11.4). The precise grading of every layer, the discrepancy register, and the program’s verification gate (formalization and external review precede any record claim) are in §14 and Appendix B.

[†]University of Toronto. *E-mail*: hongyi.yang@mail.utoronto.ca

[‡]The University of Hong Kong. *E-mail*: u3590485@connect.hku.hk

zero (an exact finite- N logarithmic law), so the consumed Christoffel values decrease to zero. No effective rate is claimed at the limit; every instantiated rung is effective, and the deepest (Theorem 10.8) gives

$$\liminf_{T \rightarrow \infty} \frac{N_0^s}{N} \geq 1 - 2\lambda_7 = 0.910460\dots, \quad \liminf_{T \rightarrow \infty} \frac{N_d}{N} \geq 1 - \lambda_7 = 0.955230\dots,$$

with λ_7 an explicit rational with forty-digit numerator and denominator: more than *nine tenths* of the zeros are simple and on the critical line, and more than *nineteen twentieths* are distinct. The arithmetic inputs are Montgomery’s band-limited pair-correlation theorem, the Siegel–Walfisz theorem and Vaughan’s estimate; two structural theorems close the moment engine — *branch equality* ($m_b(N)N^{b+1}$ is a single polynomial in N , proved through a totally unimodular lattice-polytope representation and Ehrhart theory) and *parity* ($L(-N) = (-1)^{b+1}L(N)$, by Ehrhart–Macdonald self-reciprocity) — so every consumed constant through $k = 14$ is a proved exact rational, and every further order costs polynomially many exact finite- N evaluations. Claim grading is stated per constant in §14: the constant layers are proved; the analytic chain is a certified candidate pending external review; the consumption certificates are kernel-checked in core Lean 4. A standalone reproduction and certification package accompanies the paper (<https://github.com/JoshuaHKU/zeta-density-one-reproduction>); a falsifier registry (571 pre-registered checks, 61 fired and converted, all forward-recorded) documents the research trail.

MSC 2020: 11M26 (primary); 11M06, 15B52, 52B20, 44A60 (secondary).

Keywords: Riemann zeta function; simple zeros; critical line; pair correlation; Weil explicit formula; Christoffel function; moment problem; Ehrhart theory; totally unimodular matrices.

Contents

1	Introduction	5
2	Framework and notation, following [5]	7
2.1	Weil’s form, the test family, and $\widehat{G}$	7
2.2	The zero side: inertia, tail, and the counting inequality	9
2.3	The first two moments, and notation for the ladder	9
2.4	The re-centring identity	10
3	The certificate layer	11
3.1	The Chebyshev–Markov certificate and the exact exchange rates	11
3.2	The third moment, unconditionally	12
3.3	The fourth moment: the 197/60 ledger and the class R	13
3.4	The two walls of R , and its direct measurement	14
3.5	Model calibration	14
4	The one-sided fourth-moment route: 0.6916	14
4.1	Toolkit: singular series, arc identity, tails	15
4.2	Covered zone, middle band, bridge and aggregation	15
4.3	Deterministic mains and the continuum limit	16
4.4	Assembly	17
5	The truncation discharge	18
6	The local closure	18
7	Lemma D and the exact fourth moment	19
7.1	The frame and the Parseval identity	19
7.2	The three leaves	20
7.3	The theorem	20
8	The fifth and sixth moments	21
8.1	The product identities	21
8.2	Pair layers collapse onto certified anchors	22
8.3	The pairing layer certified exactly	22
8.4	The write-outs $W1$ – $W3$	23
8.5	The connected values	24
8.6	The convolution calculus, rationality, and the midpoint protocol	26
9	The seventh and eighth moments	29
9.1	The ledgers and the central-moment collapse	29
9.2	The joint constants, certified exactly	29
9.3	The transport meta-theorem W_∞	30
10	The ninth through fourteenth moments: the Toeplitz–trace route	31
11	Consumption and the proof of Theorem 1.1	35

12 The density-one limit: the tower’s supremum is one	38
12.1 The order schema: every rung is a theorem	38
12.2 Determinacy, the atom, and the Christoffel limit	40
13 Numerical verification	41
14 Verification status	43
15 Outlook: rates, structure, and the wall	45
A The comb identification, with a worked pair	48
B Provenance, research record, and discrepancy register	49
C The exact certification layer	54

1 Introduction

That a positive proportion of the nontrivial zeros of $\zeta(s)$ lie on the critical line is due to Selberg [18]; Levinson [11] obtained the proportion $\frac{1}{3}$, Conrey [6] reached $\frac{2}{5}$, and the record along Levinson’s method stands at $\kappa > 5/12$ [15] (see also [4, 8]). Very recently the two-thirds theorem [5] proceeded by a different route: a linear-algebraic reading of Montgomery’s pair-correlation sum, in which spectral information is read from trace moments of a Gabor compression $\tilde{G}$ of Weil’s Hermitian form, and the first two moments alone — both computable unconditionally from Montgomery’s prime-side evaluation at bandwidth $\lambda \leq 1$ [13, 2] — already yield the proportion $2/3 - o(1)$ (and 0.6725 with the Montgomery–Taylor window [14]). The present paper extends that framework to the third through sixth moments. By [5, Rem. 1.1] no certificate consuming only the bandwidth-one data can exceed 0.68185; the route past the ceiling identified there (§7.5(f) of [5]) is the fourth moment, $\text{tr } \tilde{G}^4 = d\ell_1^4(13/4 + o(1))$, worth 13/18.

This paper climbs that ladder in a single text with one notation, one bibliography, one verification ledger, and one claim grading: the certificate calculus and the third moment (§3), a logically independent one-sided fourth-moment route reaching 0.6916 (§4), the exact fourth moment worth 13/18 (§§5–7), the fifth and sixth moments with their consumption (§§8–11), reaching the theorem below; then the seventh through fourteenth moments (§§9–10) with the $k = 14$ certificate 0.9104/0.9552 (Theorem 10.8); and finally the title theorem (§12): the tower’s supremum is 1 — determinacy of the moment sequence, no atom at zero, and the Christoffel limit give $\lim N_0^s/N = \lim N_d/N = 1$ (Theorem 12.7), with no effective rate claimed at the limit and every finite rung effective.

Theorem 1.1 (Main theorem; grade: certified-candidate, see §14). *With ineffective constants,*

$$\liminf_{T \rightarrow \infty} \frac{N_0^s(T, 2T)}{N(T, 2T)} \geq \frac{138058329}{162540559} = 0.849377\dots, \quad \liminf_{T \rightarrow \infty} \frac{N_d(T, 2T)}{N(T, 2T)} \geq \frac{150299444}{162540559} = 0.924688\dots,$$

and the same holds for the zeros of any fixed primitive Dirichlet L-function. En route, each rung consuming exactly what its chain establishes: $m_4 \rightarrow 13/4$ two-sidedly (hence $13/18 - \varepsilon$ and $31/36 - \varepsilon$); by the logically independent one-sided route of §4, 0.6916/0.8458; and by the $k \leq 6$ subtower alone — whose every constant is symbolically certified — the exact unconditional rung

$$\liminf \frac{N_0^s}{N} \geq \frac{2025}{2519}, \quad \liminf \frac{N_d}{N} \geq \frac{2272}{2519} \quad (\text{Corollary 11.5}).$$

(An intermediate fourth-and-fifth-moment rung considered at an earlier stage of this research is repaired and demoted in §11: without sixth-moment information it is not supported on unbounded spectral support — register item D6.)

The dependency structure of Theorem 1.1 is graded explicitly in §14: the 13/18 rung rests on the classical-technology proof of Lemma D (§7) and the consumption layer of [5] (recalled in §2); the 2025/2519 rung additionally consumes the symbolically certified constants $C_5 = \frac{1}{36}$, $\{4, 2\} = -\frac{23}{420}$, $\{6\} = -\frac{1}{126}$, $\{2, 2, 2\} = \frac{32}{105}$ (§8.5; the pairing layer is exact, §8.3); the 0.8493 headline additionally consumes the seventh- and eighth-moment layer of §9 — the symbolically certified joints $\{5, 2\} = \frac{1}{8}$, $\{2^4\} = \frac{1661}{3780}$, $\{2, 2, 4\} = -\frac{127}{840}$, $\{4, 4\} = \frac{23}{4536}$, $\{6, 2\} = -\frac{563}{11340}$, together with $C_7 = -\frac{17}{360}$ and $C_8 = \frac{157}{4032}$ (both now symbolically certified by orbit-reduced exact runs, §9.2) — the remaining graded items are the second-path ladders of the three mixed joints and the Lean extension (§14(d)), and the 2025/2519 rung remains the most conservatively graded

statement. Nothing in this paper bears on the Riemann Hypothesis itself: proportion statements are insensitive to $o(N)$ off-line zeros: the inputs (functional equation, explicit formula, mean values of Dirichlet polynomials of length $\leq T$) are satisfied by objects for which the analogue of RH is false, and the argument produces lower bounds only [5, §1.5].

Relation to [5], and what is new

This paper is built on [5] and consumes its framework without repetition: the compressed Weil matrix, the zero-side block structure, the tail estimate, and the first two trace moments are used exactly as established there (§2). The advance is fourfold. (i) *The moment ladder*: [5] consumes only $\mu_1 \rightarrow 1$, $\mu_2 \rightarrow \frac{4}{3}$; this paper evaluates $\mu_3 \rightarrow 2$ and $\mu_4 \rightarrow \frac{13}{4}$ exactly, and the μ_5 – μ_8 layers at exact rational model values $(\frac{101}{18}, \frac{640}{63}, \frac{3439}{180}, \frac{747361}{20160})$, their connected Ursell constants proved as exact rationals by symbolic polytope-volume integration, and the rational layer of every order collapsed in closed form by the re-centring identity (§§3–9). (ii) *The truncation technology*: the ℓ_1 -truncated cell ledger with Siegel–Walfisz/Vaughan closure (§5), absent from [5], is what makes the higher moments computable with classical tools. (iii) *The consumption*: the rank–trace inequality of [5, Lem. 3.2] — the matrix form of $m^2 \geq 2m - 1$ — is replaced by the Christoffel function of the pinned moment data (§11): the optimal dual polynomial is the normalized kernel polynomial, whose coefficients are *rational*, so the consumption is a closed-form exact certificate at every degree, valid on unbounded support. (iv) *The verification apparatus*: pre-registered falsifier checks, a forward-only discrepancy register, and a kernel-compiled Lean layer (§14), replacing the traditional single-grade presentation of [5]. The constants rise from $2/3$ and $\frac{5}{6}$ (distinct) to 0.8493 and 0.9246 (at the stated grades; 2025/2519 and 2272/2519 at full symbolic certification).

Relation to the pair-correlation route

Baluyot, Goldston, Suriajaya and Turnage-Butterbaugh [2] showed that Montgomery’s pair-correlation method yields two-thirds-type conclusions under hypotheses weaker than RH, and Goldston–Suriajaya [3] reach the two-thirds constant itself conditionally; those are conditional statements about the *zeros*’ pair correlation. The present route is disjoint, and unconditional: it consumes the average behaviour of $2k$ -point correlations of $\Lambda \star \Lambda$ — prime-side objects, controlled by the Siegel–Walfisz theorem and Vaughan’s estimate alone — through the compressed-matrix framework, with no hypothesis on the zeros. The two-thirds threshold, previously reached only under such hypotheses, is crossed here unconditionally with margin ($2025/2519 > 4/5 > 2/3$), and the additional information consumed relative to the bandwidth-one data is exactly the third-through-eighth moment chain of §§3–9.

The method in one page

The moments m_k of the spectral measure of $\tilde{G}$ are carried by k -cycle ledgers: cells indexed by prime-power modulus tuples and lock (offset) vectors, weighted $1/(d\ell_1^k)$, with window-overlap volumes supplying a walk geometry. The consumption is an explicit Chebyshev–Markov certificate (§3): mass near the origin of the spectral measure is bounded by moment data alone, with exact exchange rates. Three observations reduce the moment evaluations to classical tools.

Truncation (§5). The $1/\ell_1^k$ weights discharge all cells with $\ell_1 > (\log X)^B$ trivially and two-sidedly; only logarithmic moduli remain.

The spectral frame (§7, §8). The lock variance of a k -cycle obeys exact identities — at $k = 4$ a discrete Parseval identity; at every k the *product identity* $\sum_j N_k(j)^2 = \sum_t A_1(t) \cdots A_k(t)$ (two chains share a lock vector iff they differ by a common shift) — which exhibit it as a spectral L^2 -distance whose every factor is a *windowed single prime sum*. At logarithmic moduli every main-term frequency is a rational of logarithmic denominator: the Siegel–Walfisz theorem gives an arbitrary log-power saving pointwise on the doubly-major bins, Vaughan’s estimate closes the minor and mixed bins, and the singular-series comb is identified exactly by the local law below. The apparent “twin hardness” of locked correlations dissolves: no per-lock evaluation is ever attempted.

The local law and the anchor collapse (§6, §8). The mod- p structure of every cycle length is an affine lock chain, whose density obeys a coincidence-counting closed form proved in four lines and verified exhaustively. The Bell(5) = 52 and Bell(6) = 203 class ledgers of the fifth and sixth moments then collapse onto the certified four-cycle anchors through machine-exact frozen-slot identities; the only genuinely new constants above m_4 are the connected Ursell values.

Verification discipline and contributions

The research program behind this paper operates falsifier-first: every structural claim carries a pre-registered numerical check; corrections are recorded forward, never edited away; claims are graded. Sixty-one of the registry’s checks have fired and been converted — several of them shaping this paper directly, including the 38th, which corrected a consumed constant and *raised* the theorem (Appendix B, D19). The chain’s grade, its named write-out items, and the results of the certification audit are stated in §14; per the program’s verification gate, Lean formalization and external review precede any record claim. The mathematical development of this paper was carried out by Claude (Anthropic), a large language model, within a research program directed by the authors, who specified objectives, verification discipline, and publication decisions, and who take full academic responsibility for the content. In accordance with prevailing journal authorship policies, the model is not listed as an author; its role is stated here and in the Acknowledgements. The theorem claims of this paper are advanced *at the stated grades*: the constant layers as proved exact rationals, the analytic layer as a certified candidate pending external review — not as established records.

2 Framework and notation, following [5]

The construction of the compressed Weil matrix $\tilde{G}$ — Weil’s Hermitian form, the test family, the zero-side block structure, the tail estimate, and the first two trace moments — is exactly that of the two-thirds paper [5] (§§2–5 there). This section recalls the notation and states the five results of [5] that the moment ladder consumes; proofs are in [5], whose own analytic imports are the published literature [20, 10, 19, 13, 2]. Our contribution begins with the third moment (§3).

2.1 Weil’s form, the test family, and $\tilde{G}$

Throughout, $l = \log(T/2\pi)$, $\ell_1 = l + 2 \log 2 - 1$ (so $N(T, 2T) = T\ell_1/2\pi + O(l)$), and for fixed $\lambda \in (0, 1]$: $L = \lambda l$, $X = e^L$, $d = \lfloor LT/2\pi \rfloor$. For $f \in C_c^2(\mathbb{R})$ put $h_f(z) = \hat{f}(z) = \int f(u)e^{izu} du$; two

integrations by parts give

$$|h_f(x + iy)| \leq e^{|y|\Lambda_f} \min(\|f\|_1, \|f''\|_1 |x + iy|^{-2}), \quad \text{supp } f \subset [-\Lambda_f, \Lambda_f]. \quad (2.1)$$

Weil's Hermitian form is $W(f, g) = \sum_{\rho} m_{\rho} h_f(\gamma_{\rho}) \overline{h_g(\gamma_{\rho})}$, the sum over distinct nontrivial zeros $\rho = \beta + i\gamma_{\rho}$ (multiplicity m_{ρ} explicit; absolutely convergent by (2.1) and $\sum m_{\rho} |\gamma_{\rho}|^{-2} < \infty$). The zero multiset is invariant under $\rho \mapsto 1 - \bar{\rho}$, so W is Hermitian. With $s = \frac{1}{2} + i\tau$ define the prime-side densities

$$\begin{aligned} \mu(\tau) &= \frac{1}{2\pi} \text{Re} \frac{\Gamma'}{\Gamma} \left(\frac{1}{4} + \frac{i\tau}{2} \right) - \frac{\log \pi}{2\pi}, \\ \Pi_X(\tau) &= \frac{1}{2\pi(\frac{1}{4} + \tau^2)} + \frac{1}{\pi} \text{Re} \frac{X^s - 1}{s}, \\ P_X(\tau) &= -\frac{1}{\pi} \sum_{n \leq X} \frac{\Lambda(n)}{\sqrt{n}} \cos(\tau \log n), \end{aligned}$$

and $\nu_X = \mu + \Pi_X + P_X$. Weil's explicit formula [20] (in the normalization of [10, Thm. 5.12], applied to $k = f * \tilde{g}$, $\tilde{g}(u) = g(-u)$, whose transform $h_k = h_f h_{\tilde{g}}$ is entire with the decay (2.1)) gives the spectral form

$$W(f, g) = \int_{\mathbb{R}} h_f(\tau) \overline{h_g(\tau)} \nu_X(\tau) d\tau, \quad \text{supp } f, g \subset [-\frac{L}{2}, \frac{L}{2}]. \quad (2.2)$$

We record $|\Pi_X(\tau)| \leq 3\sqrt{X}/(1 + |\tau|)$, $|P_X(\tau)| \ll \sqrt{X}$, $\int_T^{2T} \mu = N(T, 2T) + O(l)$ and $\int_T^{2T} \mu^2 = T\ell_1^2/4\pi^2 + O(l^2)$ (Stirling).

Fix a nondecreasing $\varrho \in C^3(\mathbb{R})$ with $\varrho = 0$ on $(-\infty, 0]$, $\varrho = 1$ on $[1, \infty)$, a ramp width $1 \leq w \leq L/8$, and the taper $\phi(u) = \varrho((L/2 - |u|)/w)$: even, $0 \leq \phi \leq 1$, $\text{supp } \phi = [-L/2, L/2]$, $\phi = 1$ on $[-L/2 + w, L/2 - w]$. Put $a = \frac{1}{L} \int \phi^2$, $b = \frac{1}{L} \int \phi^4$ (so $1 - 2w/L \leq b \leq a \leq 1$), $\Phi = \widehat{\phi^2}$, and note the window calculus

$$(L - 2w - |y|)_+ \leq g(y) := (\phi^2 \star \phi^2)(y) \leq (L - |y|)_+. \quad (2.3)$$

The Gabor system at critical density is $f_k(u) = \phi(u) e^{-i\tau_k u}$, $\tau_k = T + 2\pi k/L$, $0 \leq k < d$, and the compressed Weil matrix is

$$G_{kl} = W(f_k, f_l) = \sum_{\rho} m_{\rho} \widehat{\phi}(\gamma_{\rho} - \tau_k) \widehat{\phi}(\gamma_{\rho} - \tau_l) = \int_{\mathbb{R}} \widehat{\phi}(\tau - \tau_k) \widehat{\phi}(\tau - \tau_l) \nu_X(\tau) d\tau, \quad \tilde{G} = G/L,$$

real symmetric by either expression; its normalized trace moments are $\mu_k = \text{tr } \tilde{G}^k / (d\ell_1^k)$.

Lemma 2.1 (Poisson summation for the Gabor system). *For all $\tau, \tau' \in \mathbb{R}$, $\sum_{k \in \mathbb{Z}} \widehat{\phi}(\tau - \tau_k) \widehat{\phi}(\tau' - \tau_k) = L \Phi(\tau - \tau')$; in particular $\sum_{k \in \mathbb{Z}} \widehat{\phi}(\tau - \tau_k)^2 = aL^2$.*

Proof. This is [5, Lem. 2.2]: the Gabor system at critical density $h = 2\pi/L$ has a translation-invariant frame kernel with no aliasing error, for any window supported in an interval of length L . $\square$

2.2 The zero side: inertia, tail, and the counting inequality

Fix $D_0 = T^{1/2}$, $I = (T, 2T]$, $I' = (T - D_0, 2T + D_0]$, and split $G = A + E$ over distinct zeros with $\gamma \in I'$ (matrix A) and the rest (E); both are real symmetric. Classify the distinct zeros with $\gamma \in I'$: S_1 (simple, on-line; s_1), S_2 (multiple, on-line; s_2), and off-line pairs $\{\rho, 1 - \bar{\rho}\}$ (p of them), so $\#Z(I') = s_1 + s_2 + 2p$ and $N(I') \geq s_1 + 2s_2 + 2p$.

Proposition 2.2 (Block structure). $n_+(\tilde{A}) \leq s_1 + s_2 + p$ and $\text{rank } \tilde{A} \leq \#Z(I')$.

Proof. This is [5, Prop. 4.1]: Sylvester inertia under pull-back — an on-line zero contributes a positive 1×1 block, an off-line pair a hyperbolic block of signature $(1, 1)$, and no independence of the evaluation functionals is needed. $\square$

Proposition 2.3 (Tail). *Let A_0 be an absolute constant with $N(t+1) - N(t) \leq A_0 \log(t+3)$ [19, Thm. 9.2]. Then for $T \geq T_0$,*

$$\|\tilde{E}\|_{\text{op}} \leq \theta_0 := 4A_0 C_1^2 X^{1/2} \frac{\log(4T)}{D_0^2} \ll l T^{\lambda/2-1}, \quad C_1 = \|\phi''\|_1 = \frac{2\|\varrho''\|_1}{w},$$

and the trace norm satisfies $\|\tilde{E}\|_1 \leq 2\theta_0/L$.

Proof. This is [5, Prop. 4.2] (the taper's r^{-2} -decay (2.1) against the zero-density bound); we restate the explicit θ_0 because the truncation discharge of §5 and the third-moment tail step consume it. $\square$

Proposition 2.4 (Counting inequality). *For every $\theta \geq \theta_0$,*

$$s_1 \geq 2n_+^\theta(\tilde{G}) - N(I'), \quad \#Z(I') \geq n_+^\theta(\tilde{G}),$$

hence, with $N(I' \setminus I) \ll D_0 l = o(N)$,

$$N_0^s(T, 2T) \geq 2n_+^\theta(\tilde{G}) - N(T, 2T) - 2N(I' \setminus I), \quad N_d(T, 2T) \geq n_+^\theta(\tilde{G}) - N(I' \setminus I). \quad (2.4)$$

Proof. This is [5, Prop. 4.5]: Weyl's inequality (at threshold $\theta \geq \theta_0$, Proposition 2.3) reduces $n_+^\theta(\tilde{G})$ to $n_+(\tilde{A})$, and Proposition 2.2 with the multiplicity count $s_2 + p \leq \frac{1}{2}(N(I') - s_1)$ gives both displays. $\square$

These inequalities hold for any locally finite multiset in the strip invariant under $\rho \mapsto 1 - \bar{\rho}$ with $N(t+1) - N(t) \ll \log(t+3)$; they contain no arithmetic. All the arithmetic is in the trace moments.

2.3 The first two moments, and notation for the ladder

Proposition 2.5 (First two moments). $\text{tr } \tilde{G} = aL N(T, 2T) + O(L\sqrt{X})$, and

$$\text{tr } \tilde{G}^2 = 2\pi bL \int_T^{2T} \mu^2 + \frac{T}{\pi} \sum_{n \leq X} \frac{\Lambda(n)^2}{n} g(\log n) + O(Ll \log l (l^2 + X)).$$

Consequently, in the endpoint regime $\lambda \rightarrow 1^-$, $w/L \rightarrow 0$: $\mu_1 \rightarrow 1$ and $\mu_2 \rightarrow \frac{4}{3}$.

Proof. This is [5, Prop. 5.3 and Thm. 5.8]: the trace by the Gabor–Poisson collapse, the second moment by Montgomery’s prime-side evaluation of the pair-correlation second moment at bandwidth ≤ 1 , unconditional in the band [13, 2]. We restate it because the diagonal $\sum_{n \leq X} (\Lambda(n)^2/n) g(\log n) = \frac{L^3}{6}(1 + O(1/L))$, with the same window autocorrelation (2.3), is the calibration consumed by the third-moment ledger (Theorem 3.4) and by the model gates of §3.5. $\square$

Remark 2.6 (Endpoint conventions). The endpoint $\lambda = 1$ is admissible because every secondary term above is $o(1)$ relative to the mains even at $X \asymp T$; a reader preferring a power saving throughout may keep $\lambda < 1$ fixed and let $\lambda \rightarrow 1^-$, with the same limits [5, Rem. 6.1]. The taper tail (the $2w/L$ deficit in (2.3)) vanishes as $w/L \rightarrow 0$; all endpoint statements in this paper take the double limit in this order (cf. [5, §7.1]).

By Proposition 2.5 and the ledger of §3, in the endpoint regime,

$$\mu_1 \rightarrow 1, \quad \mu_2 \rightarrow \frac{4}{3}, \quad \mu_3 \rightarrow 2, \quad \mu_4 = \frac{197}{60} + R + o(1), \quad (2.5)$$

where R is the shifted-correlation class (§3.3), of model value exactly $-\frac{1}{30} = \frac{13}{4} - \frac{197}{60}$. The zero-side consumption is the Christoffel-function count of Proposition 2.4: a lower bound on the number of eigenvalues of $\tilde{G}$ exceeding the vanishing threshold θ_0 of Proposition 2.3 converts into the zero counts (2.4), and the mass bound near the origin comes from the moment data alone (§3). At the moment sequence $(1, 1, \frac{4}{3}, 2, \frac{13}{4})$ the count gives 13/18 and 31/36 exactly; it is monotone in each even moment and decreasing in the odd-moment deficits.

2.4 The re-centring identity

In the endpoint limit the flat taper places the Gabor grid exactly at the window’s Nyquist rate, so the frame kernel collapses on the diagonal ($\Phi(\tau_k - \tau_l) = L\delta_{kl}$, Lemma 2.1) and the mean-field block of $\tilde{G}$ is exactly scalar. Writing

$$P := \tilde{G}/(L\ell_1) - I,$$

I commutes with P , so the binomial theorem applies to traces verbatim:

$$m_k = \sum_{j=0}^k \binom{k}{j} \Sigma_j, \quad \Sigma_j := \lim \mathbb{E} \operatorname{tr} P^j/d, \quad (2.6)$$

an *exact* bookkeeping identity: the central moments Σ_j are precisely the fully-joint (singleton-free) class sums of the Bell ledgers of §§8–9, and every singleton-containing class of every order is absorbed by the binomial coefficients. The chain of §§3–8 inverts to the exact central moments

$$\Sigma_1 = 0, \quad \Sigma_2 = \frac{1}{3}, \quad \Sigma_3 = 0, \quad \Sigma_4 = \frac{1}{4}, \quad \Sigma_5 = \frac{1}{36}, \quad \Sigma_6 = \frac{61}{252}, \quad (2.7)$$

with three nontrivial coincidences verified exactly: $\Sigma_5 = C_5$; $\Sigma_6 = \{2, 2, 2\} + \{4, 2\} + \{6\}$; and the anchor-free rational parts $\frac{67}{12}, \frac{39}{4}$ of m_5, m_6 — obtained in §8 by full Bell-ledger assembly — drop out of (2.6) by a one-line binomial inversion. Consequently the rational layers of the seventh and eighth moments are *predicted* in closed form,

$$m_7 = \frac{685}{36} + \Sigma_7, \quad m_8 = \frac{217}{6} + 8\Sigma_7 + \Sigma_8, \quad (2.8)$$

and the Bell(7)/Bell(8) ledger assemblies of §9 must land on these exact rationals — a pre-registered gate (F-PRED-RAT) which they pass: $\frac{685}{36} + C_7 = \frac{6833}{360}$ and $\frac{217}{6} + 8C_7 = \frac{3221}{90}$ reproduce the independently assembled ledger values exactly. Two unrelated derivations, four exact agreements.

Windows and cells: windows $I_m = [X, 4X]$ and $I_n = (b_2/b_1)I_m$ of length $Y \asymp X$ (more generally $Y \asymp X^\vartheta$, $\vartheta \in (\frac{1}{2}, 1]$); prime-power moduli b_1, b_2 with $g = (b_1, b_2)$, $r = b_1/g$, $q = b_2/g$; shifts $k \leq K \asymp Y/\max(r, q)$ and locks $|j| \leq J \asymp X$; $e(x) = e^{2\pi i x}$; windowed prime sums on the position lattices

$$S_m(\beta) = \sum_{m \in I_m} \Lambda(m) e(b_2 m \beta), \quad S_n(\beta) = \sum_{n \in I_n} \Lambda(n) e(b_1 n \beta).$$

$N_4(k, j)$ denotes the Λ^4 -weighted count of locked quadruples $(m, m - rk, n, n - qk)$ with $b_1 n - b_2 m = j$, and $\mathfrak{S}_4(k, j) V(k, j)$ its singular-series model, where $V(k, j)$ is the window-overlap volume of the cell — the archimedean count of quadruples in the joint window under the lock constraint, of size $\asymp Y^2/(rq\ell_1)$ after consumption normalization — with the conventions of the fourth-moment cell ledger (§4.2), frozen in the reproduction suite. The five-case local table $\kappa_p(v_p(j); v_p(b_1), v_p(b_2))$ defining $\mathfrak{S}_{b_1, b_2}(j) = \prod_p \kappa_p$ is verified digit-exact against residue counting, with $\mathbb{E}_j[\mathfrak{S}] = 1$ (§4.1).

Import status. [5] is a publicly posted preprint; its framework is consumed exactly as stated in this section, and its own analytic imports are the published literature [20, 10, 19, 13, 2]. The averaged shifted-convolution input of the one-sided route is Matomäki–Radziwiłł–Tao [12], used as published. The formalization status of the paper’s layers is graded in §14(f): the identity, local-law and certificate layers are kernel-compiled in core Lean 4; the analytic layer is not formalized.

3 The certificate layer

This section presents the certificate that carries the ladder beyond the rank–trace inequality of [5, Lem. 3.2], and the program’s unconditional moment bookkeeping, machine-verified throughout. It supplies: the Chebyshev–Markov certificate with its exact $\kappa(c)$ laws; the third-moment theorem; the 197/60 decomposition of the fourth moment with the two-bookkeeping reconciliation; and the two-wall structure of the remaining class R , with its audit apparatus.

3.1 The Chebyshev–Markov certificate and the exact exchange rates

Lemma 3.1 (Certificate; machine-verified exactly). *Let ν be a probability measure on $\mathbb{R}$ with moments m_k . Put*

$$Q(x) = \frac{(x^2 - \frac{21}{8}x + \frac{3}{2})^2}{(3/2)^2}.$$

Then $Q \geq 0$, $Q(0) = 1$, $Q \geq 1$ on $(-\infty, 0]$, Q is decreasing on $[0, \frac{21}{16}]$, and

$$\int Q d\nu = \frac{4}{9} \left(m_4 - \frac{47}{16} \right) + R_Q, \quad R_Q := \frac{4}{9} \left[-\frac{21}{4} (m_3 - 2) + \frac{633}{64} (m_2 - \frac{4}{3}) - \frac{63}{8} (m_1 - 1) \right], \quad (3.1)$$

so $|R_Q| \leq 11\varepsilon$ if $|m_1 - 1|, |m_2 - \frac{4}{3}|, |m_3 - 2| \leq \varepsilon \leq 1$. For $0 < \theta \leq \frac{1}{100}$,

$$\nu((-\infty, \theta]) \leq (1 + 4\theta) \left[\frac{4}{9} \left(m_4 - \frac{47}{16} \right) + 11\varepsilon \right].$$

If $m_1 = 1$, $m_2 = \frac{4}{3}$, $m_3 = 2$, $m_4 \leq \frac{13}{4}$, then $\nu((-\infty, \theta]) \leq \frac{5}{36} + O(\theta)$, and this is sharp: the three-atom measure supported on $\{0, a, b\}$, with $a = 0.840635\dots$, $b = 1.784365\dots$ the roots of $x^2 - \frac{21}{8}x + \frac{3}{2}$, attains it. Consequently, under $m_4 \leq \frac{13}{4} + c$ the counts of §2 give the exact linear laws

$$\liminf \frac{N_0^s}{N} \geq \frac{13}{18} - \frac{8c}{9}, \quad \liminf \frac{N_d}{N} \geq \frac{31}{36} - \frac{4c}{9}.$$

Proof. Nonnegativity is clear and $Q(0) = 1$; for $x \leq 0$ the inner quadratic is $\geq \frac{3}{2}$, so $Q \geq 1$; on $[0, a]$ the inner quadratic decreases from $\frac{3}{2}$ through 0 and $a > \frac{1}{100}$; a direct computation gives $Q(\theta)^{-1} \leq 1 + 4\theta$ for $\theta \leq \frac{1}{100}$. Expanding Q and taking expectations,

$$\left(\frac{3}{2}\right)^2 \int Q d\nu = m_4 - \frac{21}{4} m_3 + \left(\frac{441}{64} + 3\right) m_2 - \frac{63}{8} m_1 + \frac{9}{4},$$

and substituting $m = (1, \frac{4}{3}, 2, m_4)$ gives $\frac{4}{9}(m_4 - \frac{47}{16})$; the displayed R_Q collects the first-order deviations, and $\frac{4}{9}(\frac{21}{4} + \frac{633}{64} + \frac{63}{8}) \leq 11$. For sharpness, the three-atom measure with moments $(1, \frac{4}{3}, 2, \frac{13}{4})$ exists by the Hankel recurrence $m_{k+1} = \frac{21}{8}m_k - \frac{3}{2}m_{k-1}$ on the atom set, and Q vanishes doubly at a, b . The linear laws follow by inserting the mass bound into the counts of §2 exactly as in the proof display of (3.1): mass $\leq \frac{5}{36} + \frac{4c}{9} + O(\theta)$, and $1 - 2(\frac{5}{36} + \frac{4c}{9}) = \frac{13}{18} - \frac{8c}{9}$, $1 - (\frac{5}{36} + \frac{4c}{9}) = \frac{31}{36} - \frac{4c}{9}$. $\square$

Remark 3.2. Every identity of Lemma 3.1 is verified symbolically (sympy, exact rationals) in the reproduction package (`certificate_verification.py`) and re-run in the certification audit; the roots, the $5/36$ evaluation, the $\kappa(c)$ laws, and the Hankel floors $m_5^{\text{ext}} = \frac{177}{32} = 5.53125$, $m_6^{\text{ext}} = \frac{2469}{256} = 9.6445\dots$ reproduce exactly. Re-optimizing the certificate polynomial for each c (a linear program) improves the record-breaking threshold from $c < 0.0559$ to $c < 0.0721$ (`moment_lp_reopt.py`: $M^* = 3.3221 = \frac{13}{4} \times 1.0222$); the moment-pinned LP of §11 is the same mechanism run at the full moment sequence.

3.2 The third moment, unconditionally

Lemma 3.3 (Exact 3-cycle resonances). *Let $\lambda \leq 1$ and let $n_1, n_2, n_3 \leq X = T^\lambda$ be prime powers with $|\log(n_1 n_2 / n_3)| \leq C/T$. Then for $T \geq T_0(C)$, $n_1 n_2 = n_3$; the analogous statement holds for every sign pattern of the 3-cycle.*

Proof. $n_1 n_2$ and n_3 are positive integers with ratio $e^{O(1/T)}$; if distinct, their ratio differs from 1 by at least $1/n_3 \geq T^{-\lambda}$, incompatible for large T unless $\lambda = 1$; the borderline is handled by $|n_1 n_2 - n_3| \geq 1$ versus $\leq C n_3 / T \leq C$: for $n_3 \leq T/(2C)$ equality is again forced, and the range $n_3 \in (T/2C, T]$ contributes only through the taper tail, which vanishes in the normalized limit (Remark 2.6). $\square$

Theorem 3.4 (Third moment). *For every fixed $\lambda \in (0, 1]$ and admissible taper, μ_3 converges as $T \rightarrow \infty$, and its limit tends to 2 in the endpoint limit $\lambda \rightarrow 1^-$, $\eta \rightarrow 0^+$.*

Proof (nine-class ledger; full details in the archived memo *V1_completion.md*). Expand $\text{tr } \tilde{G}^3$ against the 3-cycle $\Phi(\tau_1 - \tau_2)\Phi(\tau_2 - \tau_3)\Phi(\tau_3 - \tau_1)$ after the Gabor–Poisson collapse (Lemma 2.1); the tail correction is controlled by the trace–Hölder splitting $\text{tr}(A + E)^3 - \text{tr } A^3$ with $\|E\|_{\text{op}} \leq \theta_0 \ll lT^{\lambda/2-1}$ (Proposition 2.3) and $\|A\|_F^2 = O(d\ell_1^2)$: $o(d\ell_1^3)$ in total. Of the nine classes of $\nu_X^{\otimes 3}$, $\nu_X = \mu + \Pi_X + P_X$: (i) every class containing Π_X is $O(T^{\lambda/2-1})$ pointwise on the window, negligible; (ii) single- P classes are $O(\sqrt{X} l^2 L)$ by $\sum \Lambda(n) n^{-1/2} \leq 3\sqrt{X}$; (iii) the class μ^3 equals

$d\ell_1^3(1 + O(1/l))$ (the ℓ_1 -normalization is exact to $O(1/l)$; measured 1.0048/1.0014/1.0004 at $T = 600/38400/10^8$); (iv) each of the three μP^2 classes reduces, by the $k = 2$ case of Lemma 3.3 plus the Montgomery–Vaughan Hilbert inequality for the kernel tails, to the exact diagonal $\sum_{n \leq X} (\Lambda(n)^2/n) w(\log n)$ with the same $(L - u)_+$ -shaped edge weight as at $k = 2$ (structure constants of the cycle expansion are local to edges and slots, hence pinned by the $k = 2$ case, where the total is Montgomery’s evaluation), and $\sum_{n \leq X} (\Lambda^2/n)(L - \log n) = L^3/6(1 + O(1/L))$: each class contributes $\frac{1}{3}d\ell_1^3(1 + o(1))$ at the endpoint, the three positions summing to 1; (v) the class P^3 is forced by Lemma 3.3 onto exact relations $n_1 n_2 = n_3$, i.e. powers of a single prime, and $\sum_p \sum_{a+b=c} \log^3 p p^{-c} = O(1)$ makes it $O(Tl)$. Summing: $\mu_3 \rightarrow 1 + 1 = 2$. Cross-checks: the $k = 2$ instance of the same constant chain reproduces Montgomery’s $\frac{1}{3}$ to four digits; the model gate $\int O_3 C_3 = 0$ is exact (§3.5); the finite-height shape $m_3 m_1 / m_2^2$ matches the model to 2%. Two nonstandard steps — the $k = 3$ tail adaptation of Proposition 2.3 and the locality-plus-calibration fixing of the constant chain — are written out in the archived memo and are consumed at that rigor. $\square$

3.3 The fourth moment: the 197/60 ledger and the class R

Proposition 3.5 (Evaluable classes of the fourth moment). *By the identical ledger at $k = 4$ (81 classes; Π_X - and single- P classes negligible; paired sign patterns forced exact by Lemma 3.3; $\{2, 2\}$ -pairing weights explicit window integrals),*

$$\mu_4 = \underbrace{1}_{\mu^4} + \underbrace{2}_{\mu^2 P^2 \text{ diag}} + \underbrace{\frac{17}{60}}_{\{2,2\} \text{ bare pairings}} + R + o(1) = \frac{197}{60} + R + o(1)$$

in the endpoint regime, where $\frac{17}{60} = 2\kappa$ with $\kappa = \int_{[0,1]^2} xy \Sigma_{\text{pair}} O_4 = \frac{17}{120}$ (exact quadrature of the certified 4-cycle overlap geometry), and R is the shifted-correlation class: pairs of restricted products $N_1 = b_1 m_1 \neq N_2 = b_2 m_2$ (all legs prime powers $\leq X$), weight $\Lambda^4 / \sqrt{N_1 N_2}$, window kernel $\widehat{W}(\Delta) = [\sin 2T\Delta - \sin T\Delta] / \Delta$ at $\Delta = \log N_1 / N_2$, and the explicit overlap geometry O_4 (row/column asymmetric; unordered-cell weight $Q(\beta_1, \beta_2) + Q(\beta_2, \beta_1)$); R -units $= (1/\pi) \Sigma / (d\ell_1^4)$. The model value of R is exactly $-\frac{1}{30} = \frac{13}{4} - \frac{197}{60}$.

Remark 3.6 (Bookkeeping reconciliation; audit rounds 25–26). Two bookkeeping conventions coexist in the program archive and must not be mixed. The *bare-diagonal* convention above books the $\{2, 2\}$ pairings at $\frac{17}{60}$ and carries the pair-level Montgomery plateau (the $\min(u, 1)$ saturation of C_2) inside R . The *sine-model cumulant* convention books $\{2, 2\}$ as $2t_{\text{adj}} + t_{\text{opp}}$ with the connected part separate. With the anchors now certified as exact rationals (§8.3: $t_{\text{adj}} = \frac{7}{60}$, $t_{\text{opp}} = \frac{1}{30}$), the reconciliation is exact and symmetric:

$$\underbrace{\frac{17}{60}}_{\text{bare}} = \underbrace{\frac{4}{15}}_{2t_{\text{adj}} + t_{\text{opp}}} + \underbrace{\frac{1}{60}}_{\text{plateau}}, \quad R_{\text{model}} = \underbrace{-\frac{1}{60}}_{\text{plateau (pair-of-pairs Wick)}} + \underbrace{-\frac{1}{60}}_{\text{connected core}} = -\frac{1}{30},$$

and $1 + 2 + \frac{17}{60} = \frac{197}{60}$, $2 \cdot \frac{17}{120} = \frac{17}{60}$ exactly. The archived grid values (0.2677, plateau 0.0156, connected -0.0177 ; cross-convention measurement -0.01567 in `m1_suite.py`) carried a small grid bias in t_{adj} (quoted 0.11717, exact $\frac{7}{60} = 0.11667$; register item D7); the exact connected model value $-\frac{1}{60} = -0.01667$ agrees with the cumulant engine’s independent grid measurement -0.01663 to 0.2%. The totals consumed anywhere in this paper were and remain exact ($\frac{197}{60}$, $-\frac{1}{30}$, $\frac{13}{4}$); the split is now exact too. All ledgers below state which convention they use.

3.4 The two walls of R , and its direct measurement

Everything in this subsection is unconditional structure or finite-height measurement; it is the audit apparatus against which the proofs of §§5–7 were developed and checked. (i) *Two ghost lemmas*. The mean-field quadruple term is $O(l^{-3})$: the u -Jacobian cancels $1/\sqrt{N_1 N_2}$ exactly and the inner kernel integral is $-2[\text{Si}(2T/N) - \text{Si}(T/N)]$, a boundary layer at $N \asymp T$; and the smooth part of the cluster-squared sector vanishes identically, since $\int_{\mathbb{R}} \widehat{W}_{\text{in}} d\Delta = 0$. (ii) *Two walls*. Consequently $R = R_{\text{pp}} + R_{\text{conn}}$: an edge layer carried by Poisson-dual lattice terms at top scale (some leg within $O(1)$ of X ; the explicit-formula regime) and a bulk object (model bands at $s \sim 1.5$). (iii) *Separation observable, certified at three heights*: binning by $e = l - \max u_i$ gives, at $l = 5.95/6.64/7.33$: edge $-0.0061/-0.0068/-0.0060$ (saturating toward the model -0.0156) and bulk $+0.0004/-0.0012/-0.0015$ (sign turn at $l \approx 6$, monotone toward -0.0177). (iv) *Direct measurement*. Meet-in-middle enumeration of all resonant quartic prime pairs (exact kernels, class-resolved, flat window, $\lambda = 1$) gives

l	4.56	5.94	7.33	8.72
R (per $d\ell_1^4$)	-0.0008	-0.0057	-0.0074	-0.0101

— negative at every height, trending at the $1/l$ -rate of all calibration gates; a two-code-path audit showed the sharp cutoff $C = 40$ understates $|R|$ (full-kernel -0.0012 and -0.0098 at the first two heights), so these are magnitude lower bounds. The same audit certified the ledger’s constant chain end-to-end at $k = 2, 3, 4$ to 0.05–0.8%. The $k=1$ Poisson stationary term of the edge layer is evaluated in closed form (zone-1 master formula; validated at three heights and two tapers to 5–18%): a positive decaying transient $\sim cl/\ell_1^4$ explaining the late sign turn of the near band.

3.5 Model calibration

The free-fermion cumulant engine computes the Fourier-field cumulants of the sine process by the ordered-set-partition expansion of $\text{Tr} \log(1 + (e^f - 1)K)$; gates: $C_2(v) = \min(|v|, 1)$ exactly, $\int O_3 C_3 = 0$ (whence $m_3(1) = 2$), and the $m_4(1)$ assembly 3.25103 against $\frac{13}{4}$ (grid error 0.03%). The value $m_4(1) = \frac{13}{4}$ is pinned three independent ways (cumulant assembly; GUE simulation $m_{1.4} = 1.0000, 1.3313, 1.9913, 3.2213(27)$; Hankel forcing from the extremal constants). GUE double-window Richardson extrapolation measures the model values

$$m_5(1) = 5.573 \pm 0.113, \quad m_6(1) = 10.02 \pm 0.27, \quad (3.2)$$

(sine-model/GUE measurement — see the discrepancy register, Appendix B, item D2, for the provenance correction of an earlier attribution), and the six-moment certificate at the central values gives 0.7328/0.8664: the ladder climbs only through the odd moments’ strict-positivity gaps, since the $k \leq 4$ extremal already has $m_5 = \frac{177}{32}$, $m_6 = 9.6445$. The assembled values of §8 lie inside the bands (3.2).

4 The one-sided fourth-moment route: 0.6916

This section establishes a one-sided bound on R sufficient to cross the pair-correlation ceiling 0.68185, by classical technology plus the averaged shifted-convolution theorem of [12]. It is logically independent of the exact evaluation of §§5–7 and is retained both as the program’s first

crossing of the ceiling and because its architecture (cells, bridge, aggregation, discharge ledger) is consumed by the Lemma D proof.

Theorem 4.1 (One-sided fourth-moment bound; grade: certified-candidate with the numerical-analysis step of Lemma 4.10). *With $R = R(1)$ the class of Proposition 3.5, $\limsup_{T \rightarrow \infty} R(1) \leq \bar{R} := 0.0066$. Consequently, by the count of §2 at $(1, 1, \frac{4}{3}, 2, \frac{13}{4} + \Delta)$, $\Delta = \bar{R} + \frac{1}{30} = 0.039933$,*

$$\liminf \frac{N_0^s}{N} \geq 1 - 2\Lambda_2(0) = 0.691615, \quad \liminf \frac{N_d}{N} \geq 1 - \Lambda_2(0) = 0.845807,$$

where $\Lambda_2(0) = \frac{5/108 + \Delta/3}{1/3 + 4\Delta/3} = 0.154193$, and the same for any fixed primitive Dirichlet L -function. Any $\bar{R} \leq 0.0222$ crosses the 0.68185 ceiling; the margin is 0.0156 in R -units.

(An earlier draft printed $\Lambda_2(0) = 0.156293$; the certification audit recomputed the closed form and found the stated theorem constants 0.6916/0.8458 correct and that figure stale — discrepancy register D1, Appendix B.) The constant $\bar{R}$ decomposes as $[\text{core} \leq -0.0055] + [\varepsilon_{\text{tail}} \leq 0.0111] + [\varepsilon_{CL} = 0.0010]$, where ε_{CL} is a convergence-slack allowance for the core sequence (Lemma 4.10).

4.1 Toolkit: singular series, arc identity, tails

For $b_1, b_2 \geq 1$, $j \neq 0$ the local density of $b_1 m_1 - b_2 m_2 = j$ is $\mathfrak{S}_{b_1, b_2}(j) = \prod_p \kappa_p$ with the five-case table verified digit-exact; $\mathbb{E}_j[\mathfrak{S}] = 1$. Its arc expansion:

Proposition 4.2 (Arc identity). $\mathfrak{S}_{b_1, b_2}(j) = \sum_{q \geq 1} \beta_q c_q(j)$, $\beta_q = \frac{\mu(q_1)\mu(q_2)}{\varphi(q_1)\varphi(q_2)}$, $q_i = \frac{q}{(q, b_i)}$.

At $b_1 = b_2 = 1$ this is the classical evaluation used in [12, (66)–(67)]; the general case follows by the same p -local comparison against the κ_p table. The truncation tail obeys, for $g = (b_1, b_2) \leq (\log M)^C$,

$$\sum_{|j| \leq H} (\mathfrak{S} - \mathfrak{S}^{(P)})^2(j) \ll H (\log)^{-B+O(1)}, \quad P = (\log)^B, \quad (4.1)$$

by the (q_1, q_2) -grading and $\sum_{|j| \leq H} c_q(j)^2 \ll (H + q)\varphi(q)$. *Discharge ledger* (kept, per the program's audit discipline): cells with $g > (\log)^C$ are discarded at certified cost 0.71% weight share, $\leq 3 \times 10^{-4}$ R -units. The trivial envelope $\sum_{n \leq x} \Lambda(n)\Lambda(n+h) \ll \mathfrak{S}(h)x$ uniformly for $|h| \leq x$ is [9, Cor. 3.14].

4.2 Covered zone, middle band, bridge and aggregation

Write $S_i(\alpha) = \sum_{m \sim M_i} \Lambda(m)e(b_i m \alpha)$. The two-modulus structure is discharged by four archive lemmas (rounds 38–40): the major-arc law $S_i(\alpha) = \frac{\mu(q_i)}{\varphi(q_i)} v_i(\eta) + O(M_i \exp(-c\sqrt{\log M}))$ with $q_i = q/(q, b_i)$ (b -uniformity free); the classification of joint near-major points ($q \leq \min(P^2 g, P b_1, P b_2)$), a finite b -independent family whose content is the $(q_1, q_2 \leq P)$ -graded truncation $\mathfrak{S}^{(P)}$; the dilation transfer ($\gamma_i = b_i \alpha$ obeys Vaughan's estimate in *its own* approximation quality); and a smoothing-collar lemma aligning the Gallagher width. After this discharge the $b = 1$ machinery of [12] applies verbatim:

Proposition 4.3 ([12, Thm. 1.3(i)], variance form). *For $X^{8/33+\varepsilon} \leq H \leq X^{1-\varepsilon}$, any center $h_0 \leq X^{1-\varepsilon}$, and any A :*

$$\sum_{|h-h_0| \leq H} \left| \sum_{X < n \leq 2X} \Lambda(n)\Lambda(n+h) - \mathfrak{S}(h)X \right|^2 \ll_{A,\varepsilon} HX^2(\log X)^{-A}.$$

The covered share of the class at the 8/33 exponent is 0.885 (anchored at the archived 0.797 at exponent $\frac{1}{3}$). For the middle band, with $A(b_2, j) = \sum_m \Lambda(m)\Lambda((b_2m+j)/b_1)\mathbf{1}[b_1 \mid b_2m+j]$ and $MT = \mathfrak{S}_{b_1, b_2}(j) \text{len}/(b_1b_2)$:

Lemma 4.4 (Dispersion swaps; exact). *$\sum_{b_2, j} A^2$, $\sum A \cdot MT$ and $\sum MT^2$ are, exactly, weighted twin-sum families over the structured shifts $h = qk$ (outer pairs $m - m' = rk$), the $\mathfrak{S}$ -against- Λ correlation on the strip, and $2J E_2(b_1, b_2)(\text{len}/b_1b_2)^2(1+o(1))$ respectively, where $E_2 = \prod_p \mathbb{E}_v[\kappa_p^2]$ (anchor: $E_2(1, 1) = 2 \prod_{p>2} (1 + (p-1)^{-3})$).*

Lemma 4.5 (Bridge). *For any window W and modulus q : $\sum_{k \neq 0} \sum_{n, n-qk \in W} \Lambda(n)\Lambda(n-qk) = \sum_{a \bmod q} \left[\psi_W(q, a)^2 - \sum_{n \in W, n \equiv a} \Lambda(n)^2 \right]$.*

Lemma 4.6 (Structured-to-full aggregation). *Let W have length Y , $\text{dev}(h)$ the (smoothly weighted) twin deviation on W , Q a set of moduli with shift counts $K_q \leq H/q$, $H \leq Y^{1-\varepsilon}$. If $\sum_{h \leq H} |\text{dev}|^2 \ll_{A'} HY^2(\log)^{-A'}$ for every A' , then for every A*

$$\sum_{q \in Q} \sum_{k \leq K_q} |\text{dev}(qk)|^2 \ll_A HY^2(\log Y)^{-A},$$

uniformly in Q — no factor of the moduli is lost.

Proof. The left side is $\sum_{h \leq H} \nu(h) |\text{dev}(h)|^2$ with $\nu(h) \leq \tau(h)$. Split at $\nu \leq (\log Y)^C$: the generic part costs $(\log)^C$ against the arbitrary- A' input; on the rare part use the envelope $|\text{dev}(h)| \leq 9\mathfrak{S}(h)Y$ with $\sum_{h \leq H, \tau(h) > V} \tau(h) \ll H(\log)^3/V$ at $V = (\log)^C$. Choose $C = A + 4$, $A' = 2A + 4$. $\square$

The band consumer applies a single Cauchy–Schwarz *in the shift variable only* (any Cauchy–Schwarz across the family pays its Poisson floor, measured 50–60 $\times$ over budget, and is never used). The glue layer (welding weight $w_k(n) = \sum_{m \in I(n)} \Lambda(m)\Lambda(m-rk)$) closes by the exact factorization of main terms, Abel summation against the glue for the major arcs, and the divisor-bounded envelopes of [12, Prop. 5.4, App. A] for the minor arcs.

4.3 Deterministic mains and the continuum limit

Write $\mathfrak{S}_{b_1, b_2} = \bar{D} + g_P + g_{\text{tail}}$ ($\bar{D}$ the exact finite local mean; g_P the $q \leq P$ oscillating piece; $P = 40$).

Lemma 4.7 (Partial-sum identity). *For every $M \geq 1$, $\sum_{j \leq M} (\mathfrak{S}_{b_1, b_2} - \bar{D})(j) = - \sum_{d \geq 1} d \left\{ \frac{M}{d} \right\} \gamma_d^{\text{free}}$,*

with explicit Euler-local γ , absolutely convergent against the ℓ^1 weights $d \min(1, M/d)$.

Lemma 4.8 (Finite- J evaluation and uniform tail). *For every finite J and θ , $\sum_{j \leq J} \text{tail}_P(j) \text{row}_j(\theta) = \sum_{d \geq 1} \gamma_d^{>P} S_0(d\theta; \lfloor J/d \rfloor)$ with $S_0(x; K) = \sum_{k \leq K} \frac{\sin 2xk - \sin xk}{k}$; hence the θ -uniform bound $\leq C_{\text{saw}} \sum_d |\gamma_d^{>P}|$. No $J \rightarrow \infty$ interchange is performed anywhere.*

Lemma 4.9 (Certified sawtooth constant). $|S_0(x; K)| \leq C_{\text{saw}} := 2.10$ for all x and $K \geq 1$.

Proof. For $K \leq 64$: rigorous grid verification with the Lipschitz bound $|\partial_x S_0| \leq 3K$ certifies $\max |S_0| \leq 1.8675$. For $K > 64$: the Dirichlet representation with the sine-integral comparison near the singular set and the second mean value theorem away from it gives ≤ 2.05 ; samples at $K = 128, 512, 4096$ read 1.842, 1.849, 1.851, approaching $\text{Si}(\pi) = 1.85194$ from below. (Re-run in the certification audit: identical values.) $\square$

Zone-integrating the uniform tail against the certified geometry gives $\varepsilon_{\text{tail}} \leq 0.0115$ ($T = 2400$), 0.0111 ($T = 4800$), decreasing; the evaluated core (exact finite forms) is $\text{core}(2400) = -0.00549$, $\text{core}(4800) = -0.0064$, certified against the pipeline anchor (-0.00545 vs the suite's -0.00544).

Lemma 4.10 (Continuum limit of the core; grade: proof shape with computed constants — see the grading note below). *$\text{core}(T)$ converges as $T \rightarrow \infty$ to $C_{\text{core}} = \sum_{b_1 \leq b_2} \frac{\Lambda(b_1)\Lambda(b_2)}{b_1 b_2} F_{\infty}(b_1, b_2)$, an absolutely convergent series of computable continuum sawtooth integrals, with $|\text{core}(T) - C_{\text{core}}| \ll (\log T)^{-1}$; the ensemble collapses onto the universal (1,1)-class (Mertens mass concentrates on pairs of large primes), whose slope and constant are computed, not fitted: partial-sum slopes $0.99684 \rightarrow 0.99962 \rightarrow 0.99996$ over successive decades (classical law, exactly 1 — now proved: the Dirichlet series is $\zeta(s+1)E(s)$ with $2C_2E(0) = 1$, Proposition C.1), and the constant $c_0 = -1.4228$ is now identified in closed form: $c_0 = \gamma - 2 = -1.4227843\dots$ (Euler–Mascheroni γ ; Proposition C.1, verified to 10^{-5} at $M = 2 \times 10^7$); θ -ladder $c^* = -0.5433$, decade-stable, matching the archive's independent $-0.55(2)$ ($c^* = \log \pi + \gamma - 2 - \overline{\text{osc}}$ in the pin of the archived F -QC5 gate, so the remaining computed piece of c^* is the mean oscillation $\overline{\text{osc}} \approx 0.2652$, not identified here). The continuum quadrature with Richardson extrapolation in $1/l$ ($l = 40, 60, 100, 160$) yields $C_{\text{core}} = -0.0209 \pm 0.0026$, independently confirming the archive ladder $-0.020(2)$ — two disjoint routes at 4%.*

Grading note. The convergence structure (scaled-variable Riemann sum, universality collapse) is an analytic argument; the values of the collapsed series' slope and constant, and the quadrature of C_{core} , are computed constants with stated bands, not certified enclosures. Theorem 4.1 charges only the conservative core $\leq -0.0055 + \varepsilon_{CL}$, $\varepsilon_{CL} = 0.0010$, against the measured decreasing grid; the certification write-out of the quadrature is a named open item (§14). Substituting the computed C_{core} would give $\bar{R} = -0.0098 \pm 0.0026 < 0$, hence 0.7031/0.8515; that upgrade is *not* consumed anywhere in this paper. More strongly: *no result outside this section consumes any constant of this section* — Theorems 1.1, 10.8 and 12.7 proceed by the exact route of §§5–7 onward and are independent of Lemma 4.10; what Lemma D borrows from §4 is its certified *architecture* (cells, bridge, aggregation, discharge ledger), no numerical value.

4.4 Assembly

At $\lambda = 1$, uniform full-kernel convention:

$$R(1) \leq \text{core}(T) + \varepsilon_{\text{tail}} + o(1) \leq (-0.0055 + 0.0010) + 0.0111 + o(1) = 0.0066 + o(1),$$

the $o(1)$ collecting the covered-zone, band and low-zone deviations of §4.2 (the low-zone cap's certified finite-height price at the operating point is 0.0163, inside the LP budget 0.0387+0.0177). Feeding $\bar{R} = 0.0066$ into the count gives Theorem 4.1. $\square$

5 The truncation discharge

Lemma 5.1 (Cell-mass bound). *Uniformly in the lock and shift ranges, the Λ^{2k} -weighted, consumption-normalized mass of the k -cycle cells at parameter ℓ_1 is $\ll Y^2(\log Y)^c \ell_1$.*

Proof. For a fixed cell, each locked configuration is determined by (m, k) together with the residue data of the lock; the count of (m, k) pairs in the windows is $\ll YK \ll Y^2/\max(r, q)$, each carrying weight $\ll (\log Y)^{2k}$, and the lock constraint selects a $1/\ell_1$ -proportion of the n -range up to divisor-bounded multiplicity; the consumption normalization (division by the ledger volumes $V \asymp Y^2/(rq\ell_1)$) leaves at most one factor ℓ_1 . The numerical ledger face measures the aggregate mass essentially flat in ℓ_1 (fitted exponent -0.11), well inside the bound. $\square$

Lemma 5.2 (Multiplicity). *The number of unordered prime-power tuples with cell parameter $\ell_1 = \ell$ is $\ll_\varepsilon \ell^\varepsilon$ (divisor-bounded).*

Proposition 5.3 (Truncation–consumption compatibility). *Let F denote the consumed k -cycle cell aggregate, $k \geq 4$, and $F^{(P)}$ its restriction to $\ell_1 \leq P$. Then*

$$|F - F^{(P)}| \ll_\varepsilon Y^2(\log Y)^c P^{-(k-2)+\varepsilon} \quad \text{two-sidedly.}$$

Hence for every A_0 there is B (any fixed $B > (c + A_0)/(k - 2)$ suffices) such that proving the variance bound for moduli $\leq P = (\log X)^B$ proves it in the consumed form, with total truncation cost $\ll (\log X)^{-A_0}$: only moduli $\leq (\log X)^B$ are ever consumed.

Proof. Weight $\times$ mass $\times$ multiplicity $= \ell^{-k} \cdot \ell \cdot \ell^\varepsilon$ summed over $\ell > P$. The bound is on absolute values, hence two-sided. The archived D1 layer independently certifies the squared-remainder grading at $k = 4$ ($\mathbb{E}_j[(F - F^{(P)})^2] \ll P^{-1-\delta}$; tail locals $t_p^2 \sim 4/p^2$). Two truncations coexist in the program and must not be confused: the sawtooth/frequency truncation inside the C_{full} -quadrature machinery of §4.3, whose tail carries that limit, and the present ℓ_1 -cell truncation, whose tail is trivially small; the audit trail (archive rounds 91–93, item S1) verified the disentanglement. An earlier draft stated the multiplicity as ℓ_1 ; the divisor bound of Lemma 5.2 is what the displayed exponent uses. $\square$

6 The local closure

The local (mod p) structure of every moment order is governed by affine lock chains. Fix a cycle length $b \geq 2$, a prime power $P_0 = p^s$, unit coefficients $a_1, \dots, a_b$, and free locks $j_1, \dots, j_{b-1} \in \mathbb{Z}/P_0$; define points by $n_{t+1} \equiv a_t n_t + j_t$, the last lock induced by cyclic closure. Write $n_t = A_t(n_1 + d_t)$ where $A_{t+1} = a_t A_t$, $A_1 = 1$, $d_1 = 0$; then $d_{t+1} - d_t = j_t/A_{t+1}$ is a unit multiple of j_t . Let $N(j)$ count the $n_1 \in \mathbb{Z}/P_0$ with all b points prime to p .

Proposition 6.1 (Coincidence-counting closure). *Pointwise in the lock vector j ,*

$$N(j) = P_0 \frac{p - D(j)}{p}, \quad D(j) = \#\{d_t \bmod p : t = 1, \dots, b\},$$

and for any class $\mathcal{C}$ of locks,

$$\kappa_b(\mathcal{C}) = \frac{1 - \mathbb{E}[D \mid \mathcal{C}]/p}{(1 - 1/p)^b}, \quad \mathbb{E}_j[\kappa_b] = 1 \text{ exactly.}$$

Forbidden-class merges are pair events of weight $1/(p - 1)$; triple merges enter at $O(p^{-2})$.

Proof. Since A_t is a unit, $n_t = A_t(n_1 + d_t)$ is prime to p iff $n_1 \not\equiv -d_t \pmod{p}$; the all-unit condition excludes exactly the $D(j)$ distinct residues $\{-d_t\}$, each of density $1/p$ in $\mathbb{Z}/P_0$, giving the count. The class formula is its average against $(1 - 1/p)^b$. For the mean-one identity: under uniform locks the increments are independent uniform, so $(d_2, \dots, d_b)$ is uniform on $(\mathbb{Z}/p)^{b-1}$ and $\mathbb{E}[D] = p(1 - (1 - 1/p)^b)$, whence $1 - \mathbb{E}[D]/p = (1 - 1/p)^b$. A merge across a gap of lock-sum σ requires $p \mid \sigma$: impossible for a single unit lock; for a sum of two units it occurs for exactly one value of the second, weight $1/(p - 1)$. $\square$

Machine verification (pipeline, `face_local_closure`; re-run in the certification audit with identical output): zero pointwise violations over full enumerations and closed-form match to 2×10^{-16} for every lock class at $b = 4, 5, 6$, $p \in \{5, 7, 11, 13\}$, $s \in \{1, 2\}$, with the mean-one identity exact. The κ_4 closed forms of the fourth-moment table are instances ($\kappa_4(0, 0) = 1 - (4 - \frac{2}{p-1})/p$). The proposition supplies the singular-series combs at every cycle length and the j -averaged factorization identities used below. (A first draft of the corresponding Lean seed omitted the normalising inverse in $d_t = B_t A_t^{-1}$; the exact-enumeration proxy fired — registry r103, the 29th conversion — and the seed was corrected: Appendix B.)

7 Lemma D and the exact fourth moment

The fourth-moment evaluation $m_4 \rightarrow 13/4$ closes the estimate isolated by the program (“Lemma D”: the averaged, lock-resolved quadruple-correlation variance) in its consumed, truncated form. In its *unrestricted* form — moduli up to a power of X — it remains open, one structural step beyond the published diagonal cases of the Barban–Davenport–Halberstam/dispersion lineage. Two observations remove the difficulty from the consumed form: the truncation discharge of §5, and the spectral frame below.

7.1 The frame and the Parseval identity

By the archived, machine-verified bridge and glue identities (§4.2), the (k, j) -aggregated variance for a modulus pair is carried by the *dilated density* $\rho(\tau)$: with A_m, A_n the autocorrelations of the two position-lattice sequences, on the analysis grid of period N (a power of two exceeding four times the largest position, so no correlation wraps),

$$\rho(\tau) = \sum_t A_m(t) A_n(t - \tau) = \text{irfft}(|S_m|^2 |S_n|^2)(\tau),$$

and the lock-variance obeys the *exact* discrete Parseval identity

$$\sum_{\tau=0}^{N-1} \rho(\tau)^2 = \frac{1}{N} \left[\text{dens}_0^2 + 2 \sum_{0 < \beta < N/2} \text{dens}_\beta^2 + \text{dens}_{N/2}^2 \right], \quad \text{dens} = |S_m|^2 |S_n|^2, \quad (7.1)$$

verified bitwise in the pipeline (10^{-16} relative; re-run in the certification audit). The same identity applied to ρ –Model turns the variance into the spectral L^2 -distance of dens from the model comb. We classify the frequencies β by the arc quality of the pair $(\gamma_m, \gamma_n) = (b_2\beta, b_1\beta) \bmod 1$ at cutoff $P^* = (\log X)^{B^*}$: both major (MM), mixed, or both minor.

7.2 The three leaves

Lemma 7.1 (MM leaf). *Let β be a doubly major bin: $\gamma_m = a/Q + \eta$, $Q \leq P^*$, $|\eta| \leq \text{arc width}$, and likewise γ_n . Then, for every A' ,*

$$S_m(\beta) = \frac{\mu(Q)}{\varphi(Q)} I_m(\eta) + O_{A', B^*}(Y(\log Y)^{-A'}), \quad I_m(\eta) = \sum_{m \in I_m} e(m\eta),$$

uniformly, and the same for S_n ; the constants are ineffective. Consequently, on the MM bins, $|\text{dens} - \text{comb}| \ll \text{main} \cdot Y(\log Y)^{-A'} + Y^2(\log Y)^{-2A'}$ pointwise, where comb is the product of the squared main terms.

Proof. $S_m(\beta) = \sum_m \Lambda(m) e(m\gamma_m)$ is a windowed single prime sum at $\gamma_m = a/Q + \eta$. Splitting into progressions mod Q and applying the Siegel–Walfisz theorem in the form [7, §22] (or [10, Cor. 5.29]) with partial summation in the window against the phase $e(m\eta)$ gives the display; the constants depend only on (A', B^*) and are ineffective through Siegel’s theorem. The pointwise comb comparison follows by expanding the product. The per-bin face (`face_sw`) measures the absolutely normalised error across six families: RMS 3.6×10^{-3} at $T = 2400$ falling to 3.6×10^{-4} at $T = 153600$. $\square$

Remark 7.2 (Ineffectivity enters here). This is the single point where ineffective constants enter the chain: Siegel’s theorem gives no computable $T_0(A')$, so every theorem constant downstream is ineffective, while every finite-height face measured in §13 carries effective constants. The two kinds of statement are graded separately throughout (see the closing remark of §13); nothing else in the paper — the identities, the certificate, the local law, the window calculus — consumes an ineffective input.

Lemma 7.3 (Minor and mixed leaves). *For γ minor at cutoff P^* (denominator range $((\log X)^{B^*}, Y(\log X)^{-B^*})$ by Dirichlet approximation), $|S(\gamma)|^2 \ll Y^2(\log Y)^{-B^*/2+c_0}$ by Vaughan’s estimate. Hence the doubly-minor spectral mass carries a $(\log)^{-B^*+2c_0}$ saving squared, and each mixed bin at least the one-sided saving; the mixed and minor contributions to the spectral L^2 -distance are $\ll (\log X)^{-B^*/2+c_1}$ relative to the natural size. The size bookkeeping (well-spacing of image families and the hybrid large-sieve/minor-sup stratum budget) is classical and is face-checked at $24\text{--}31\times$ tightness (`face_arcs`).*

Lemma 7.4 (The comb is the singular series). *The spectral main comb of Lemma 7.1 equals the truncated singular-series model: the surviving frequencies are exactly the CRT progression $t \equiv 0 \pmod{b_2}$, $t \equiv \tau \pmod{b_1}$ in the τ -frame, and the local factors match the coincidence-counting closed form of Proposition 6.1; the truncation tail of the singular series beyond P^* is $\ll (\log X)^{-B^*+c}$ by the standard $\sum_{q > P^*} \mu^2(q)(q, h)/\varphi(q)^2$ bound. The identification is displayed explicitly, with a worked modulus pair, in Appendix A; the τ -resolved model tracks the measured profile at 0.2–4% across sixteen families and eight heights (`face_dispersion`).*

7.3 The theorem

Theorem 7.5 (Lemma D, truncated form). *Let B be fixed, $b_1, b_2 \leq (\log X)^B$ prime powers, windows of length $Y \asymp X^\vartheta$, $\vartheta \in (\frac{1}{2}, 1]$. Then for every $A > 0$,*

$$\sum_{k \leq K} \sum_{|j| \leq J} |N_4(k, j) - \mathfrak{S}_4(k, j)V(k, j)|^2 \ll_{A, B} KJV^2(\log X)^{-A}.$$

Consequently, with Proposition 5.3, $m_4 \rightarrow 13/4$ two-sidedly.

Proof. Choose $B^* = 2(A + c_1 + B)$ and apply (7.1) to ρ -Model. The spectral L^2 -distance splits over the arc classification. On MM bins, Lemma 7.1 with $A' = A + 2B^*$ gives a pointwise relative saving $(\log)^{-A'}$, hence $\sum_{\text{MM}} |\text{dens} - \text{comb}|^2 \ll (\log)^{-2A'+c} \sum \text{main}^2$; by Lemma 7.4 the comb agrees with the model to $(\log)^{-B^*+c}$ relative, absorbed by the choice of B^* . The mixed and doubly-minor bins contribute $\ll (\log)^{-B^*/2+c_1}$ relative by Lemma 7.3. The (k, j) -aggregation constants of the bridge and glue layer are archived and certified (§4.2); the normalization match to KJV^2 is the consumer interface there. All bounds are on absolute values, hence two-sided; all constants depend only on (A, B) and are ineffective. Choosing B^* then A' in that order involves no circularity. $\square$

Remark 7.6 (The progression-frame route; falsifier r95). The archived reduction of the variance to a locked two-modulus progression covariance suggests a shorter argument: expand both progression errors in additive characters and evaluate by Siegel–Walfisz. Executed at the moduli (r, q) separately this produces *incomplete* character sums (the lock couples residues modulo r to residues modulo q , and the a -sum over one modulus has no orthogonality against the other); the expansion must be lifted to the joint modulus $M \asymp rq \leq (\log X)^{2B}$, where it becomes a complete-sum argument parallel to the one above. An earlier draft presented that sketch without the lift; it was replaced by the spectral proof, in which the τ -transform runs over the full period and the orthogonality (Lemma 7.4) is exact. The slip and its correction are falsifier r95 on the forward record.

Corollary 7.7 (The 13/18 rung). *With ineffective constants,*

$$\liminf_{T \rightarrow \infty} \frac{N_0^s(T, 2T)}{N(T, 2T)} \geq \frac{13}{18}, \quad \liminf_{T \rightarrow \infty} \frac{N_d(T, 2T)}{N(T, 2T)} \geq \frac{31}{36},$$

and the same for any fixed primitive Dirichlet L -function.

Proof. Theorem 7.5 and Proposition 5.3 give $m_4 = 13/4 + o(1)$ two-sidedly; Theorem 3.4 gives $\mu_3 \rightarrow 2$; the certificate (Lemma 3.1, equivalently the Christoffel-function count of Proposition 2.4 at $(1, 1, \frac{4}{3}, 2, \frac{13}{4})$) gives 13/18 and 31/36 exactly, monotonicity in the fourth moment converting $m_4 \rightarrow 13/4$ into the two liminf bounds. The Dirichlet L -function case is as in [5, Thm. E]: the framework uses only the functional-equation symmetry $\rho \mapsto 1 - \bar{\rho}$, the zero-density bound, and the explicit formula, all available for any fixed primitive Dirichlet L -function. Constants are ineffective through Lemma 7.1. $\square$

8 The fifth and sixth moments

8.1 The product identities

Proposition 8.1. *For a b -cycle of position lattices sharing a window, with any per-lattice densities f_i, g_i : $\sum_j N_f(j)N_g(j) = \sum_t \prod_i C_i(t)$, $C_i(t) = \sum_x f_i(x)g_i(x+t)$. In particular the lock variance is the pointwise product of the lattice autocorrelations, summed over the common shift.*

Proof. Two chains share a lock vector iff they differ by a common shift; the chain sum then factorizes over the positions. $\square$

This identity ($O(N \log N)$ to evaluate) transports the spectral argument of §7 *factor by factor* to every cycle length: each factor is a windowed single prime sum, so the three-leaf closure

(per-factor Siegel–Walfisz on all-major bins, Vaughan on any minor factor, comb by Proposition 6.1, truncation by Proposition 5.3 with stronger exponents) applies verbatim. Machine faces: the twin-singular-series product model tracks the empirical five-fold product at 0.28–2.5% (six modulus patterns, $T = 4800$ to 76800, improving with height) and the six-fold product at 0.11–1.29%.

8.2 Pair layers collapse onto certified anchors

The $\text{Bell}(5) = 52$ and $\text{Bell}(6) = 203$ class ledgers are machine-enumerated (class tables 1/10/15/10/10/5/1 and 1/15/45/15/20/60/10/15/15/6/1; pentagon chords $10+5$; hexagon chords $30+15$; matching crossing spectrum $5/6/3/1$). Frozen-slot identities — a singleton block repeats a walk value and leaves the overlap range unchanged; machine-exact — reduce every pair class to the four-cycle geometry, whose anchors are now *exact rationals* (§8.3): $t_{\text{adj}} = \frac{7}{60}$, $t_{\text{opp}} = \frac{1}{30}$, $\Phi_4 = \frac{1}{4} - 2t_{\text{adj}} - t_{\text{opp}} = -\frac{1}{60}$ (sine-model cumulant convention, cf. Remark 3.6). Odd-block classes vanish by the local law. Hence

$$m_5 = 1 + \underbrace{10t_{\text{adj}} + 5t_{\text{opp}} + 5\Phi_4}_{= 67/12 \text{ exactly}} + C_5, \quad m_6 = \frac{39}{4} + 6C_5 + \{2, 2, 2\} + \{4, 2\} + \{6\},$$

the rational parts being anchor-free (the t -geometry cancels: $10t_{\text{adj}} + 5\Phi_4 + 5t_{\text{opp}} = \frac{5}{4}$ and $30t_{\text{adj}} + 15t_{\text{opp}} + 15\Phi_4 = \frac{15}{4}$ identically — verified symbolically in the certification audit and kernel-compiled in the Lean package, §14(f)), with $\{2, 2, 2\} = \frac{32}{105}$ now *certified exactly by dihedral placement class* (Proposition 8.2; the superseded bundle $\frac{131}{420}$ and its correction are register item D19), $\{4, 2\} = -\frac{23}{420}$ (the four-cycle connected object consumed jointly with a spectator pair) and $\{6\} = -\frac{1}{126}$ (both likewise certified exactly, §8.5). The former write-out items W1–W3 are written out in §8.4.

8.3 The pairing layer certified exactly

Proposition 8.2 (Exact pairing constants). *The two-chord anchors are the exact rationals $t_{\text{adj}} = \frac{7}{60}$, $t_{\text{opp}} = \frac{1}{30}$. The 15 three-chord matchings of the hexagon fall into five dihedral placement classes, of sizes 2/3/6/3/1 and crossing numbers 0/0/1/2/3, with exact values*

$$T_0 = \frac{3}{70} \text{ (sequential)}, \quad T'_0 = \frac{17}{420} \text{ (nested)}, \quad T_1 = \frac{1}{90}, \quad T_2 = \frac{1}{180}, \quad T_3 = \frac{1}{70},$$

whence

$$\{2, 2, 2\} = 2T_0 + 3T'_0 + 6T_1 + 3T_2 + T_3 = \frac{32}{105}.$$

The crossing number is not a complete invariant: the two crossing-zero classes carry different values ($T_0 \neq T'_0$), and every bundle in this paper is therefore evaluated placement class by placement class (register D19).

Proof. Each constant is the integral over a box of $(1 - \text{spread})_+ \cdot \prod \min(|x_i|, 1)$ for an explicit walk (the chord-pattern walks of the Bell ledgers): a piecewise-polynomial integrand all of whose kink hypersurfaces are $\{0, \pm 1\}$ -combinations of the variables at levels $\{0, \pm 1\}$. Since 0 is a walk position and every variable is a position or a difference of two positions, the support lies in $[-1, 1]^d$ and the min-caps never bind. The integral is then computed *exactly* in rational arithmetic by iterated integration with full breakpoint enumeration: inner breakpoints have unit leading coefficients, so all collision and boundary resultants remain small-integer forms, and on every

kink-free piece the integrand is a polynomial of degree at most $2/4/8$ (in the inner/middle/outer variable), integrated exactly by closed Newton–Cotes rules of order $5/7/9$ at rational nodes. Every piece is integrated twice (whole piece, and the two halves): for polynomials of degree below the rule order the two exact rationals must coincide, and any missed kink forces a discrepancy; no discrepancy occurred. t_{opp} and T_0 additionally admit short closed-form evaluations by orthant decomposition (Appendix C), agreeing with the algorithmic values; every class value is reproduced by a second, independent exact method (polytope decomposition sharing no arithmetic with the iterated integration), and orbit invariance is machine-checked by evaluating extra members of each dihedral class. Code and outputs: `certification/exact_t222.py` and the `phase0` audit of the reproduction package. $\square$

Remark 8.3 (The classification correction; register D19). An earlier stage of this research consumed $\{2, 2, 2\} = 5T_0 + 6T_1 + 3T_2 + T_3 = \frac{131}{420}$, classifying the 15 matchings by crossing number and evaluating one representative per crossing class; the nested crossing-zero geometry was never separately evaluated, and the archived ladder value $0.31170(60)$ was assembled from the same four representatives — a circular confirmation. The error was caught by a pre-registered calibration gate of the eighth-moment pairing run (the 38th conversion), the nested class was closed exactly ($T'_0 = \frac{17}{420}$, three independent paths), and two model-side estimators that had failed their anchors by $55\text{--}110\sigma$ and $11\text{--}21\sigma$ against the old value — and read $\sim 1\sigma$ against the corrected one — are retro-recorded as independent detections. The correction *lowers* M_6 and *raises* the theorem: the previously published bounds $(0.7962/0.8981)$ consumed a weaker-but-true upper bound and remain valid.

Remark 8.4 (The anchor correction). The archived anchor $t_{\text{adj}} = 0.11717$ (rounds 69–71 grid) is corrected by Proposition 8.2 to $\frac{7}{60} = 0.11\bar{6}$: a $+0.43\%$ rectangle-rule bias at kinks. Nothing consumed downstream changes: the m_5 and m_6 rational layers are anchor-free (above), and the reconciliation of Remark 3.6 becomes exact. The formerly conjectural rational identifications of the three-chord classes are now theorems. Register items D7, D8.

8.4 The write-outs W1–W3

The three items carried as named write-outs in the audit ledger are discharged here; their pre-registered faces are in `certification/writeouts_verification.py` (all green; one face fired during construction and was converted — Appendix B).

Lemma 8.5 (W2: free-slot factorization). *In any b -cycle cell ledger, a slot whose lock is unconstrained factorizes exactly: the cell aggregate equals the contracted $(b - 1)$ -cycle aggregate times the window mass $\sum_{m \in W} \Lambda(m)$ of the free slot, and $\sum_{m \in W} \Lambda(m) = |W|(1 + O(e^{-c\sqrt{\log X}}))$.*

Proof. The lock sum over the free slot’s offset is unrestricted, so the double sum factorizes by Fubini; on the walk-geometry side this is the frozen-slot identity (a singleton block repeats a walk value and leaves the overlap range unchanged — machine-exact, `face_o5_gate/face_b6_frame`). The mass evaluation is the prime number theorem with the classical error term [7, §18]; the consumption normalization divides it out. The face verifies both the exact Fubini identity on the live family and the PNT mass at height $(1.00041 \text{ at } X = 2 \times 10^5)$. $\square$

Lemma 8.6 (W3: triple-pair factorization). *The $\{2, 2, 2\}$ class aggregate of the sixth-moment ledger equals the product of the three pair main terms against the joint overlap volume, up to a deviation aggregate $\ll_A K J' V'^2 (\log X)^{-A}$ in the consumed normalization.*

Proof. The six positions form three disjoint affine pair-chains. The p -local density factorizes over the three chains by the CRT and Proposition 6.1 (disjoint slots; the merger corrections are the pair events already booked in the local law), so the singular-series model is the product of the three twin series. Archimedean coupling enters only through the joint window overlap (the walk geometry O_6), which is the certified exact layer of Proposition 8.2. The deviation from the product of pair mains is controlled by three applications of the $k = 2$ instance of the spectral transport: by Proposition 8.1 the resolved variance factorizes over the three chains, each factor a windowed single prime sum, and the three-leaf closure of §7 applies to each factor with the exponents of Proposition 5.3. The face tracks the product model at 1.9% (median over four shift patterns) at $X = 6 \times 10^4$. $\square$

Proposition 8.7 (W1: the resolved multi-dimensional comb). *For a b -cycle with $b - 1$ lock coordinates, the resolved dilated density $\rho(\tau_1, \dots, \tau_{b-1})$ obeys the multi-dimensional comb identification: the model comb is supported on the product CRT progressions (one congruence per coordinate, as in Lemma 7.4), with local factors the κ_b -values of Proposition 6.1 and twin-series factors off the moduli; and the resolved spectral L^2 -distance from the comb obeys the bound of Theorem 7.5 with the same exponents.*

Proof. By Proposition 8.1 the resolved variance is carried by the product $\prod_i |S_i|^2$ on the joint analysis grid, and the multi- τ transform is complete in each coordinate, so the $(b-1)$ -dimensional Parseval identity holds coordinatewise (verified exactly in the face, 1.2×10^{-16}). The frequency classification of §7 applies per factor: on all-major bins each factor obeys Lemma 7.1; any minor factor is closed by Lemma 7.3; the comb identification in each coordinate is Lemma 7.4, whose CRT progression now runs per coordinate (complete orthogonality in each, since each transform covers its full period). The argument of Theorem 7.5 then repeats verbatim with $b - 1$ classification indices; the truncation exponents strengthen with k (Proposition 5.3). The previously stated bracketing (aggregate faces above, resolved four-cycle case below) is subsumed: the aggregate instances are the τ -summed cases, and the resolved four-cycle case is Lemma 7.4 itself. The resolved-comb face (worked 5-cycle family) shows the CRT class contrast directly. $\square$

8.5 The connected values

$C_5 = \int O_5 C_5^{\text{Urs}} = \frac{1}{36}$, $\{4, 2\} = -\frac{23}{420}$ and $\{6\} = -\frac{1}{126}$, proved in exact rational arithmetic and independently confirmed at machine precision.

The symbolic certification. The direct route — the piecewise exact integration of Proposition 8.2, generalized from three to four and five dimensions — closes $\{4, 2\}$ term by term, but not $\{6\}$: the piece count multiplies across integration levels, and single five-dimensional terms project to centuries of exact arithmetic. The closure is a change of representation. Since

$$\text{ov}(P) = \left(1 - (\max_i p_i - \min_i p_i)\right)_+ = \text{Leb}\{t \in \mathbb{R} : t \leq p_i \leq t + 1 \ \forall i\},$$

each term of the signed sums of §8.6 is the volume of one bounded rational polytope,

$$\int_{\mathbb{R}^d} \text{ov}_A(x) \text{ov}_B(x) dx = \text{Vol}_{d+2}\{(x, t, s) : t \leq A_i(x) \leq t + 1, s \leq B_j(x) \leq s + 1\},$$

with every facet coefficient in $\{0, \pm 1\}$ (boundedness because 0 lies in both position families, forcing $t, s \in [-1, 0]$; the weight $\min(|v|, 1)$ of the $\{4, 2\}$ variant lifts one further dimension); the cost migrates from the *product* of per-level piece counts to the size of the face lattice.

Each volume is evaluated by the exact facet recursion $\text{Vol}_n(P) = \frac{1}{n} \sum_i \text{dist}(c, H_i) \text{Vol}_{n-1}(F_i)$, arranged so that projecting each facet to a coordinate hyperplane cancels the irrational facet normalizations — the computation stays in rational arithmetic throughout — and memoized on tight constraint sets. All 1310 terms (150 for C_5 , 78 for $\{4, 2\}$, 1082 for $\{6\}$) are evaluated as exact fractions, archived per term in the reproduction package, and their signed sums are

$$C_5 = \frac{1}{36}, \quad \{4, 2\} = -\frac{23}{420}, \quad \{6\} = -\frac{1}{126}$$

exactly. Gates, all passed: the pure four-cycle through the same code path returns $\Phi_4 = -\frac{1}{60}$ (by both exact methods); the three-dimensional regression reproduces every exact value of Proposition 8.2 ($\frac{3}{70}, \frac{1}{90}, \frac{1}{180}, \frac{1}{70}, \frac{7}{60}, \frac{1}{30}$); the two symbolic methods agree fraction by fraction on every term computed by both; the term construction (position families, masks, signs, weights) is cross-checked numerically against the ladder engine; and the three placement classes of $\{4, 2\}$ come out as $-\frac{2}{315}, -\frac{1}{1260}, -\frac{1}{252}$, recombining with multiplicities 6/6/3 as $-\frac{23}{420}$.

The numerical confirmation. Seven-rung grouped midpoint ladders (dv down to 0.00125 for C_5 , 0.0125 for $\{6\}$, 0.0015625 for the four-dimensional spectator-pair variant of $\{4, 2\}$; GPU-accelerated engine with support pruning, term trie and negation symmetry, mathematically identical to the reference engines and cross-validated bitwise on three architectures) converge, with Romberg adjacent-row drift 2×10^{-16} , to the proved fractions; within the error band the rational reconstruction is *unique by the $Q^2 \geq 1/(2\varepsilon)$ criterion* and lands on the same values. Credibility anchors: the identical pipeline reproduces, in the same runs, the exact values 0 ($b = 3$) and $\Phi_4 = -\frac{1}{60}$ ($b = 4$) at 10^{-18} ; the raw-ladder error ratios against the certified fractions converge to the pure h^2 value 4.000; and $\{4, 2\}$ *hit its pre-registered candidate* ($-\frac{23}{420}$, registered from three coarse reference rungs before the fine ladder ran), componentwise. The seven-cycle constant is identified at one rung less as $C_7 = -\frac{17}{360}$ (identification grade only; its symbolic closure is delegated, and it is consumed in the $k = 7, 8$ layer at that stated grade, §9). The mixed odd-block class $\{3, 3\}$ *vanishes*: identically zero — exact floating-point zeros, not small numbers — on every grid tested (two protocols, four resolutions, all three placement classes), the model-side face of the odd-block vanishing proved on the arithmetic side in §8; the residual allowance it formerly carried in the m_6 budget is retired (register D17). An observation, not consumed: the certified and identified constants through level 6 — $\frac{1}{3}, \frac{7}{60}, \frac{1}{30}, -\frac{1}{60}, \frac{32}{105}, \frac{1}{36}, -\frac{23}{420}, -\frac{1}{126}$, the $\{4, 2\}$ components and the five hexagon chord classes — have denominators dividing 1260, and with $C_7 = -\frac{17}{360}$ all divide $2520 = \text{lcm}(1..9)$; but the law *fails at level 8*: the $\{2^4\}$ orbit values of §9 carry denominators up to $90720 = 2^5 \cdot 3^4 \cdot 5 \cdot 7$, and the bundle $\frac{1661}{3780}$ has $3780 \nmid 2520$ — which is precisely why small-denominator reconstruction found no candidate there (register D21). (Earlier endpoint-grid evaluations are retired — register item, 32nd conversion.) Fluctuations are controlled by the factor-wise variance lemmas (the transports of Theorem 7.5; the resolved multi-dimensional comb exhibits are bracketed by the aggregate faces and the resolved four-cycle case — now Proposition 8.7). End-to-end faces: the assembled $m_5 = \frac{101}{18} = 5.6\bar{1}$ and $m_6 = \frac{640}{63} = 10.15873\dots$ both lie inside the sine-model bands (3.2).

Grading note. By the dimension theorem of §8.6 every connected constant is rational; the symbolic certification above computes the three consumed constants exactly, and they enter §11 at *proof* grade — exact-rational computer-assisted, the grade of Proposition 8.2 (register D18). The former identification proviso is discharged; the numerical identification (machine-precision Romberg limit, uniqueness-guaranteed reconstruction, double exact-anchor calibration, the pre-registration hit for $\{4, 2\}$) stands as an independent confirmation, agreeing bitwise with the symbolic values. The precise grading is §14(d).

8.6 The convolution calculus, rationality, and the midpoint protocol

The connected constants admit a second, structural representation that both fixes their arithmetic nature and dictates the numerical protocol. Write W for the unit-square window autocorrelation and B for the sine-kernel box; the cumulant integrands assemble from iterated convolutions of these two profiles. Three results (derivations and verification gates in the reproduction package, the `cyclic_cumulant` and `midpoint_ladder` engines):

(i) *The trace/convolution identity.* At $b = 2$ the connected integral evaluates in closed form, $\int O_2 C_2 = 1 - \|W \star \widehat{B}\|^2 = \frac{1}{3}$ exactly — the calibration gate (F-TRACE) for everything below; the $b = 3, 4$ instances reproduce the exact pairing values of Proposition 8.2 to 3×10^{-14} .

(ii) *The partition-cyclic cumulant formula (F-CYC).* For the b -cycle connected class,

$$C_b(v) = \sum_{P \vdash [b]} (-1)^{|P|-1} \sum_{\text{cyclic orders of } P} \text{ov}(\text{merged-}w \text{ walk}),$$

the inner sum running over the distinct cyclic arrangements of the blocks and ov the overlap volume of the merged walk — validated against the direct Ursell integrands at $b = 3, 4, 5$ to 3×10^{-14} .

(iii) *Dimension and rationality.* In every (P, σ) -term of (ii) the walk constraints eliminate one integration variable ($(m - 1)$ -dimensionality), and the resulting domain is a finite union of polytopes cut by the unit-coefficient hyperplane family $\mathcal{H}$ of Appendix C(A1) with polynomial integrands of rational coefficients. Each term is therefore a rational number, and so are C_5 , $\{4, 2\}$ and $\{6\}$ — the same argument that produced $t_{\text{adj}} = \frac{7}{60}$ and $\{2, 2, 2\} = \frac{32}{105}$ in closed form. What separates the remaining constants from Proposition 8.2 is only the piece count of the exact integration, not its kind; the polytope route of §8.5 tames that piece count and carries the exact integration through all three constants, closing item N1 (§14).

(iv) *The prolate-trace identity.* Let A denote the sine-kernel integral operator $K(x, y) = \text{sinc}(\pi(x - y))$ restricted to $[0, 1]$ (the classical prolate operator at unit density). Then for every $m \geq 1$,

$$\text{tr } A^m = \int_{\mathbb{R}^{m-1}} (1 - \text{span}(0, p_1, \dots, p_{m-1}))_+ \prod_{i=1}^m \text{sinc}(\pi v_i) dv,$$

the right-hand side being exactly this paper's overlap-weighted pure S -cycle polytope integral.

Proof: write $\text{tr } A^m = \int_{[0,1]^m} \prod K(x_i, x_{i+1}) dx$, substitute the increments $v_i = x_{i+1} - x_i$, and note $\int_{\mathbb{R}} \prod_k \mathbf{1}[p_k \in [t, t + 1]] dt = (1 - \text{span})_+$ — the overlap weight *is* the translation average of the hard window. $\square$ (Verified numerically at $m = 3$ to all printed digits.) Thus the entire pure-moment cycle layer, at every order, is carried by a single classical spectrum; the connected constants differ from it only through the coincidence (Ursell) dressing, which a dedicated probe shows is *not* a spectral-polynomial functional of that spectrum.

(v) *The non-crossing closed-form calculus.* For pairing classes whose chord diagram has *no crossing* (nesting allowed), the class value collapses to one-dimensional convolution algebra: with $g(t) = \int_t^{t+1} |x| dx = t^2 + t + \frac{1}{2}$ on $[-1, 0]$, the sequential $\frac{b}{2}$ -chord orbit equals $\int_{-1}^0 g(t)^{b/2} dt$, and nested chords compose as iterated one-dimensional convolutions. Four exact confirmations: $b = 6$ sequential = $\frac{3}{70}$ and nested = $\frac{17}{420}$ (the two crossing-zero classes of Proposition 8.2); $b = 8$ sequential = $\frac{83}{5040}$ (the archived $\{2^4\}$ orbit); and the $b = 10$ sequential value $\frac{73}{11088}$, *predicted by the closed form before consulting the archive* and found there exactly (the largest $\{2^5\}$ orbit). A companion *span decomposition* extends the elementary reach across crossings at low width: reducing each orthant's overlap factor to an explicit piecewise-linear span turns all

five $b = 6$ classes — the crossing values $\frac{1}{90}, \frac{1}{180}, \frac{1}{70}$ included — into Dirichlet-type polytope integrals, re-deriving the corrected bundle $\{2, 2, 2\} = \frac{32}{105}$ of Proposition 8.2 in closed form: a third, fully independent path for the D19 constant. The calculus also identifies the source of the denominator phenomenology: $\int \text{sinc}^{2m}$ and its relatives are central B-spline values, whose factorial-family denominators generate the level-dependent lcm-patterns observed throughout (§15(e)). The intrinsic complexity parameter of the general (crossing) case is the *crossing width* of the diagram — non-crossing sectors collapse to closed form, while fully crossing orbits retain the ambient dimension; this is quantified in §15(c).

(vi) *The log-determinant identity: no atom at the origin.* The tower's measurements come from the exact finite model: the Gram matrix $\hat{G}_N = WW^*/N$ of the Fourier system $W_{jm} = e^{im\theta_j}$ ($m = 0, \dots, N-1$) at the eigenangles θ_j of a Haar-random $N \times N$ unitary, whose expected spectral moments reproduce the sine-model tower. Its log-determinant is exactly computable at every finite N : $\det \hat{G}_N = |\Delta(e^{i\theta})|^2/N^N$ with Δ the Vandermonde, and Morris's integral $\mathbb{E} |\Delta|^{2s} = \Gamma(1 + (1+s)N)/(\Gamma(2+s)^N \Gamma(1+N))$ gives, on differentiating at $s = 0$,

$$\mathbb{E} \left[\frac{1}{N} \sum_i \log \lambda_i \right] = H_N - 1 - \log N, \quad \text{Var} \left[\frac{1}{N} \sum_i \log \lambda_i \right] = \psi'(N+1) - \frac{\psi'(2)}{N},$$

exactly for every N (at $N = 1$ both sides vanish). Consequently the limiting spectral measure ν has

$$\int \log x d\nu = \lim_N (H_N - 1 - \log N) = \gamma - 1$$

— the Euler constant's second exact appearance in the program (the continuum quadrature constant of item N2 is $\gamma - 2$) — and, quantitatively, *no atom at the origin*: since $\log^+ x \leq (x-1)_+$ and, by Cauchy-Schwarz against the exact second central moment $\Sigma_2 = \frac{1}{3}$ of (2.7), $\frac{1}{N} \sum_i (\lambda_i - 1)_+ \leq \sqrt{\Sigma_2}$, the negative part obeys $\frac{1}{N} \sum_i \log^- \lambda_i \leq \frac{1}{\sqrt{3}} + 1 - \gamma + o(1)$, whence by Markov

$$\nu([0, \varepsilon]) \leq \frac{1/\sqrt{3} + 1 - \gamma}{\log(1/\varepsilon)} \implies \nu(\{0\}) = 0.$$

Both identities are enforced as standing acceptance gates on every spectral pool of the program (gate F-LOGDET): on the three archived pools ($N = 128, 256, 512$; 800/250/60 samples) the mean identity holds at $-0.08/+0.24/-0.48$ standard errors and the variance identity, with kurtosis-corrected error bars, at $-1.05/+1.53/+0.10$. The same Morris integral in fact yields the *entire* cumulant generating function exactly at every finite N : for $s > -1$,

$$\log \mathbb{E} \exp \left(s \sum_i \log \lambda_i \right) = \log \Gamma(1 + (1+s)N) - N \log \Gamma(2+s) - \log \Gamma(1+N) - sN \log N,$$

whence every cumulant in closed form, $\kappa_j(\sum_i \log \lambda_i) = N^j \psi^{(j-1)}(N+1) - N \psi^{(j-1)}(2)$, and the limiting free energy $\Lambda(s) = \lim_N \frac{1}{N} \log \mathbb{E} e^{s \sum \log \lambda_i} = (1+s) \log(1+s) - s - \log \Gamma(2+s)$. All cumulants of order $j \geq 2$ are extensive, so the log-determinant obeys a central limit theorem with the exact variance constant $\Lambda''(0) = 2 - \pi^2/6$ and concentrates exponentially — the concentration form of the no-atom statement (gate F-CGF: tilted means at $s = \pm 0.1$ and third cumulants reproduce the closed forms on the archived pools within one standard error). The consequence for the program's horizon is drawn in §15(f).

(vii) *The Toeplitz-trace representation: exact finite- N moments.* Since WW^* and W^*W share their spectrum, $\hat{G}_N$ of (vi) is *isospectral* to the random Toeplitz matrix $G'_{mn} = t_{m-n}$ built

from the CUE traces $t_k = \frac{1}{N} \text{tr } U^k$. Expanding $\text{tr}(G')^b$ over index walks gives, exactly at every finite N ,

$$m_b(N) = \frac{1}{N^{b+1}} \sum_{k_1 + \dots + k_b = 0} (N - \text{span}(0, S_1, \dots, S_{b-1}))_+ N^{\#\{k_i=0\}} \mathbb{E} \prod_{k_i \neq 0} \text{tr } U^{k_i},$$

S_j the partial sums — the discrete twin of the overlap weight $(1 - \text{span})_+$ that runs through this paper — and the trace moments are exact at every N by Schur orthogonality, $\mathbb{E}[p_\mu \bar{p}_\nu] = \sum_{\lambda \vdash K, \ell(\lambda) \leq N} \chi^\lambda(\mu) \chi^\lambda(\nu)$, the Diaconis–Shahshahani range being the no-long-column sector. Two consequences carry two-line proofs. *Vanishing lemma*: a nonzero weight forces $\text{span} < N$, hence $|k_i| < N$ stepwise and $K < N$ for every $(+\dots+, -)$ pattern, landing those configurations in the orthogonality sector where mixed cycle types integrate to zero. *Closed forms at every finite N* :

$$m_2(N) = \frac{4}{3} - \frac{1}{3N^2}, \quad m_3(N) = 2 - \frac{1}{N^2}, \quad \Sigma_3(N) \equiv 0,$$

by $\mathbb{E}|\text{tr } U^k|^2 = \min(k, N)$ and an elementary lattice sum. Exact evaluation of the full sum (rational arithmetic; two independent implementations, direct summation against multiset-aggregated dynamic programming, bitwise-matched) delivers $m_b(N)$ through $b = 10$, and the values organize into polynomials in $1/N^2$ of degree $\lfloor b/2 \rfloor$ whose constant terms *re-derive, independently of every position-space method in this paper*, the tower $\frac{13}{4}, \frac{101}{18}, \frac{640}{63}, \frac{3439}{180}, \frac{747361}{20160}$ — in particular a fourth independent path to the D19-corrected sixth moment, the fit being agnostic and explicitly excluding the superseded $\frac{4297}{420}$ — with every identification confirmed bitwise on held-out values of N (e.g. the 12-digit-numerator rational $m_9(7)$ and the held-out $m_{10}(8)$, both predicted before computation). The $1/N^2$ coefficients are exact finite-size laws, e.g. $\Sigma_2(N) = \frac{1}{3} - \frac{1}{3N^2}$ through $\Sigma_8(N) = \frac{633}{2240} - \frac{3371}{2016}N^{-2} + \frac{12247}{2880}N^{-4} - \frac{5545}{1008}N^{-6} + \frac{277}{105}N^{-8}$, and they reproduce the systematic deficits observed in the 10^7 -sample calibration runs at $N = 128$ and $N = 192$ across all six known central moments (twelve comparisons, all within statistics): measurement bias converted into theorem. Whether $m_b(N)N^{b+1}$ is a polynomial in N for every b (observed through $b = 10$; eight held-out confirmations, plus a vanishing top coefficient at $b = 9$ confirming the degree law) is the representation’s one open structural question at the time it was first observed; §10 now proves it outright (Theorem 10.4): the answer is yes, at every b and every $N \geq 1$, with the parity $L(-N) = (-1)^{b+1}L(N)$ (Theorem 10.5). A sector-isolation experiment had sharpened its shape: writing the sum as Wick sector plus character-correction sector (Schur orthogonality), each sector is separately a *period-two* quasi-polynomial from $b = 5$ on (even and odd branches verified with their own held-out points through $N = 17$ at $b = 5$), and the odd-parity discrepancies of the two sectors cancel exactly in the total — a cancellation that the polytope representation of §10 now explains: the sector split is an artifact of the character basis, and the assembled object has no residue classes at all.

The midpoint protocol. The numerical ladders for (i)–(iii) revealed a grid pathology recorded as the registry’s 32nd conversion (Appendix B, D10): endpoint grids on the kinked integrands sign-flip at $b \geq 6$ on coarse steps, while midpoint grids — which never sample on the kink set $\mathcal{H}$ — settle monotonically. All connected constants of this paper are evaluated by grouped midpoint-grid ladders, calibrated at $b = 4$ against the exact values of Proposition 8.2 (agreement to the ladder’s own band), and the earlier endpoint-grid adoptions are retired.

9 The seventh and eighth moments

9.1 The ledgers and the central-moment collapse

The $\text{Bell}(7) = 877$ and $\text{Bell}(8) = 4140$ class ledgers are machine-enumerated twice, by two independent implementations (one canonicalizing dihedral patterns directly, one by orbit decomposition under D_k with the divisibility check $|\text{orbit}| \mid |D_k|$ — a check that caught a deduplication error in one enumerator before any value was computed; register D20). Every class containing a 3-block vanishes (the odd-block vanishing of §8, verified classwise); every singleton-containing class collapses by the frozen-slot identities onto lower anchors and is absorbed, order by order, by the binomial coefficients of the re-centring identity (2.6). What survives is exactly the singleton-free layer:

$$\Sigma_7 = \{5, 2\} + C_7, \quad \Sigma_8 = \{2^4\} + \{2, 2, 4\} + \{4, 4\} + \{6, 2\} + C_8,$$

with placement-class inventories (all dihedral, all machine-audited twice): $\{5, 2\}$: 21 partitions in 3 orbits of size 7; $\{2^4\}$: 105 matchings in 17 orbits; $\{2, 2, 4\}$: 210 in 22; $\{4, 4\}$: 35 in 7; $\{6, 2\}$: 28 in 4 (sizes 8/8/8/4). The rational layers are the closed forms (2.8), and the ledger assemblies reproduce them exactly (gate F-PRED-RAT: two independent derivations, §2.4).

9.2 The joint constants, certified exactly

Proposition 9.1 (Exact joint constants of levels 7 and 8). *Evaluated placement class by placement class, in exact rational arithmetic, by the facet-recursion polytope integrator (second method: iterated Newton–Cotes breakpoint integration; the two share no arithmetic):*

$$\{5, 2\} = 7 \left(\frac{5}{504} + \frac{1}{360} + \frac{13}{2520} \right) = \frac{1}{8}, \quad \{2^4\} = \frac{1661}{3780},$$

$$\{2, 2, 4\} = -\frac{127}{840}, \quad \{4, 4\} = \frac{23}{4536}, \quad \{6, 2\} = -\frac{563}{11340},$$

the $\{2^4\}$ value being the size-weighted sum of 17 exact orbit values (archived per orbit; denominators up to 90720). The pure connected constants, by the same integrator under orbit reduction (685 D_7 -orbits of 9366 terms; 6027 D_8 -orbits of 94586 terms; signed orbit-weight bookkeeping verified against the raw term sums at $b = 5, \dots, 8$):

$$C_7 = -\frac{17}{360}, \quad C_8 = \frac{157}{4032}.$$

Hence, all inputs now symbolically certified,

$$\Sigma_7 = \frac{1}{8} - \frac{17}{360} = \frac{7}{90}, \quad \Sigma_8 = \frac{307}{1260} + \frac{157}{4032} = \frac{633}{2240},$$

and, by (2.8), $m_7 = \frac{685}{36} + \frac{7}{90} = \frac{3439}{180}$, $m_8 = \frac{217}{6} + \frac{28}{45} + \frac{633}{2240} = \frac{747361}{20160}$.

Verification pedigree. $\{5, 2\} = \frac{1}{8}$ hit its pre-registered candidate (registered from the model-side measurement 0.124944 before the symbolic run; deviation of the symbolic result from the candidate: zero), and has since acquired a full per-orbit GPU ladder as well (3/3 orbits within 1.1×10^{-6}) — a deliberate redundancy, since the C_8 episode below shows that a candidate hit is not by itself a guarantee. The $\{2^4\}$ orbit values agree with an independent three-rung midpoint ladder, computed before the symbolic run, at 17/17 classes (maximum deviation 1.3×10^{-7} from the h^4 extrapolations; the pre-registered ladder band was optimistic by a factor 3 — register

D21). $\{2, 2, 4\}$, $\{4, 4\}$ and $\{6, 2\}$ currently rest on the facet-recursion path alone (acceptance suite 22/22 against every previously certified constant, orbit-invariance checks green); their independent second-path GPU ladders have returned and passed (F-V-JOINTS, 33/33 orbits within 3×10^{-6} ; §14(d)). The delegated symbolic runs for the pure constants have *returned*: C_7 's run (685 orbits, 208s on 24 cores) landed *exactly* on the identification value $-\frac{17}{360}$ — a prediction registered two revisions earlier, now a third pre-registration hit — and C_8 's run (6027 orbits, 3811s on 200 processes) returned $\frac{157}{4032}$, *missing* the pre-registered candidate $\frac{7}{180}$ by 4.96×10^{-5} : inside the tolerance gate (3×10^{-3}), so the corner moved and the certificate re-optimized (§11); the candidate miss is forward-recorded (D22), and the denominator $4032 = 2^6 \cdot 3^2 \cdot 7$ repeats the level-8 break of the 2520 law. Model-side cross-checks on the returned values: $\Sigma_8 = \frac{633}{2240} = 0.28259$ against the measurement 0.28249 (deviation 9.9×10^{-5}); m_7, m_8 against the CUE model to 1.6×10^{-4} , 1.7×10^{-3} .

9.3 The transport meta-theorem W_∞

The arithmetic side of the seventh and eighth moments consumes the same four k -uniform devices as §§5–8 — the product identity (Proposition 8.1, any b), the coincidence-counting local law (Proposition 6.1, any cycle length), the truncation discharge (Proposition 5.3, whose exponent $P^{-(k-2)}$ *strengthens* with k), and the three-leaf spectral closure of §7. Rather than writing out the four classes of this paper one by one in analogy with W1–W3, we record the write-out *once*, uniformly in the block structure; every joint class of every moment order is an instance, and the marginal analytic cost of each further order is zero.

Theorem 9.2 (W_∞ : factor-wise transport of singleton-free classes). *Let $b \geq 2$ and let $\mathcal{B} = \{B_1, \dots, B_r\}$ be a partition of the b -cycle's slots into blocks of size ≥ 2 , none of size 3. Regard each block as an affine lock chain of the b -cycle cell ledger (locks free within a block; blocks interacting only through the shared window). Then the class aggregate equals the product of the chain main terms against the joint overlap volume O_b , up to a deviation aggregate, in the consumed normalization,*

$$\left| \text{Agg}(\mathcal{B}) - \prod_{i=1}^r \text{Main}(B_i) \cdot O_b \right| \ll_{A,B} K J' V'^2 (\log X)^{-A} \quad \text{for every } A > 0.$$

Proof. (i) *Local factorization.* The blocks' position lattices are disjoint, so the mod- p density factorizes across blocks by the CRT; within each block it is the coincidence-counting closed form (Proposition 6.1, any chain length), the cross-block merger corrections being the pair events already booked in the local law (weight $1/(p-1)$; triple mergers $O(p^{-2})$). The singular-series model is therefore the product of the per-chain combs. Classes containing a 3-block vanish identically by the odd-block vanishing of §8 (arithmetic side) and its model face, so their exclusion loses no generality.

(ii) *Archimedean coupling only through O_b .* The only archimedean interaction between blocks is the joint window overlap — the walk geometry O_b — which is the certified exact layer of §8.3, §9.2.

(iii) *Factor-wise variance transport.* By Proposition 8.1, valid for arbitrary per-lattice densities, the resolved variance is the product of the chain autocorrelations summed over the common shift; each coordinate of the multi- τ transform runs over its full period, so the $(r-1)$ -dimensional Parseval identity holds coordinatewise exactly as in Proposition 8.7. The frequency classification proceeds per factor: on all-major bins every factor is a windowed single prime sum obeying

Lemma 7.1; any minor factor is closed by Lemma 7.3; the comb identification in each coordinate is Lemma 7.4, its CRT progression running per coordinate with complete orthogonality.

(iv) *Truncation and constants.* The ℓ_1 -truncation discharges by Proposition 5.3 with exponent $P^{-(b-2)+\varepsilon}$, strengthening with b ; the (k, j) -aggregation constants of the bridge and glue layer and the $KJ'V'^2$ normalization interface are the archived certified layer of §4.2, independent of the block structure. All bounds hold for absolute values, hence two-sidedly; all constants depend only on (A, B) and are ineffective through Siegel's theorem. No step depends on b or on $\mathcal{B}$ beyond the parameters displayed. $\square$

Corollary 9.3 (Instances). *W3 (Lemma 8.6) is the case $\mathcal{B} = \{2, 2, 2\}$; the classes consumed in this paper are the further instances $W4 = \{5, 2\}$, $W5 = \{2^4\}$, $W6 = \{2, 2, 4\}, \{6, 2\}$, $W7 = \{4, 4\}$; and the ninth- and tenth-order ledgers ($\{7, 2\}, \{5, 4\}, \{5, 2, 2\}; \{8, 2\}, \{6, 4\}, \{5, 5\}, \{6, 2, 2\}, \{4, 4, 2\}, \{4, 2, 2, 2\}, \{2^5\}$) require no further write-out. Pure classes $\mathcal{B} = \{b\}$ are not routed through Theorem 9.2: their model values enter directly and their fluctuations are controlled by the factor-wise variance lemmas, as at $k = 5, 6$ (§8.5).*

Theorem 9.4 (The seventh and eighth moments). *In the endpoint regime, at the grades stated in Proposition 9.1,*

$$m_7 \rightarrow \frac{685}{36} + \Sigma_7 = \frac{3439}{180}, \quad m_8 \rightarrow \frac{217}{6} + 8\Sigma_7 + \Sigma_8 = \frac{747361}{20160},$$

two-sidedly at the grade of layer (b) of §14: the rational layers by the re-centring identity and the frozen-slot collapse (both exact), the singleton-free classes by Theorem 9.2 (instances $W4$ – $W7$, Corollary 9.3) with the certified joint constants, and the fluctuation control by the factor-wise variance transports of Theorem 7.5 with the truncation exponents of Proposition 5.3 at $k = 7, 8$.

Proof. The trace expansion at $k = 7, 8$ runs over the Bell ledgers of §9.1. Classes containing a 3-block vanish by the odd-block vanishing (§8); classes containing singletons factorize exactly by $W2$ (Lemma 8.5) and collapse onto lower anchors — the totality of these contributions is the binomial-rational layer, equal to $\frac{685}{36}$ resp. $\frac{217}{6} + 8\Sigma_7$ -corrections by (2.6), with the ledger assembly reproducing the closed forms exactly (F-PRED-RAT). The singleton-free classes are the W_∞ instances of Corollary 9.3 and the pure classes C_7, C_8 ; their model values are the certified constants of Proposition 9.1, and their deviations transport factor-wise as in the $k = 5, 6$ layer (§8.5, “Fluctuations”), with stronger truncation exponents. The tail and threshold steps are k -independent (Proposition 2.3). Cross-checks: the assembled $m_7 = 19.1055\dots$ and $m_8 = 37.0714\dots$ lie inside the model-measurement bands $(19.09 \pm 0.09, 37.1 \pm 0.3)$ and above/inside the corrected moment-cone vertex $(19.0712, 36.7622)$ over the pinned $m_1, \dots, m_6$. $\square$

10 The ninth through fourteenth moments: the Toeplitz–trace route

The six further central moments consumed by the $k = 10, 12, 14$ certificates are produced by the representation of §8.6(vii), which we now run as an engine. By the isospectrality of the Vandermonde–Gram and trace-Toeplitz matrices, $m_b(N) = \mathbb{E} \int x^b d\nu_N$ is an exact rational at every finite N (times N^{b+1} it is an *integer* — the lattice representation is a sum of integers — now a theorem, Corollary 10.7(i), confirmed on all eighty-two archived values — seventy-nine in the frozen table plus the three odd-order extension points), and the computed values organize, at every order through $b = 14$, into polynomials in $1/N^2$ of degree exactly $\lfloor b/2 \rfloor$. Each order is

determined by the exact values at $N = 2, \dots, \lfloor b/2 \rfloor + 2$ and confirmed at a *pre-registered held-out* $N = \lfloor b/2 \rfloor + 3$, predicted before computation and matched bitwise in every case; during the campaign the degree was measured, not assumed — it is now the theorem of Corollary 10.7(ii) — (at $b = 12$ the degree-5 fit was excluded by its own failed prediction before degree 6 was adopted, and at $b = 14$ the registered degree-6 prediction for $m_{14}(9)$ agreed with the true value to six decimal places yet differed as an exact fraction — a discrimination only exact arithmetic can make). Two independent implementations (direct summation; multiset aggregation with span dynamic programming, written from the mathematical specification alone) agree bitwise on every overlap point. The resulting exact layers, each satisfying its pre-registered re-centring anchor (F-PRED-RAT) and matching three independent 10^7 -sample measurement tiers at $N = 128/192/256$ through the exact finite-size laws $\Sigma_j(N)$ (thirty-nine comparisons, all within three standard errors):

$$\begin{aligned} \Sigma_9 &= \frac{52207}{302400}, & \Sigma_{10} &= \frac{1333891}{3326400}, & \Sigma_{11} &= \frac{128291}{340200}, & \Sigma_{12} &= \frac{1092211019}{1556755200}, \\ \Sigma_{13} &= \frac{45789263}{52390800}, & \Sigma_{14} &= \frac{27183066233}{18162144000}, \end{aligned}$$

equivalently $m_9 = \frac{11166011}{151200}$, $m_{10} = \frac{83443081}{554400}$, $m_{11} = \frac{852071287}{2721600}$, $m_{12} = \frac{1033020076559}{1556755200}$, $m_{13} = \frac{240004263497}{167650560}$, $m_{14} = \frac{85585542088667}{27243216000}$. The full $1/N^2$ -polynomials, the seventy-nine exact lattice values (plus the three odd-order extension points), and the eighteen-gate suite live in the reproduction package (`constants/tt/`).

Three structural facts sharpen the record; the first two are now theorems (Corollary 10.7 below), the third remains empirical. *Integrality and denominators*: $m_b(N)N^{b+1} \in \mathbb{Z}$ always (Corollary 10.7(i)), the limit denominator divides $(b+1)!$ for *every* b (Corollary 10.7(iii)), and the sharper Lagrange bound $D(b)$ is verified for $b \leq 13$ — the denominator phenomenology of §15(e) now has a proved ceiling at all orders. *Monotone pricing*: the exact Christoffel sequence satisfies $n\lambda_n(0)$ strictly increasing for $n \leq 7$ (0.2500, 0.2778, 0.2942, 0.3012, 0.3070, 0.3108, 0.3134), equivalently each new orthonormal value $p_n(0)^2$ lies below the running average — which is *why* every exact rung has landed slightly below its measured-band pricing. The route’s foundations are three statements — model transfer, branch equality, parity — which we now prove, resting on a single shared mechanism: by the cyclic cluster structure of determinantal processes [17] (exact at every finite N), each joint cumulant of the traces $\text{tr} U^{k_i}$ is a signed sum over cyclic orders of integrals $\int \prod S_N(\theta_i - \theta_{\sigma_i}) \prod e^{ik_i\theta_i}$; expanding each S_N in its Fourier sum, the angular integrals enforce conservation edge by edge and the integral collapses to the *lattice box count* $(N - \text{span}_\sigma(k))_+$ — exactly the span weights of (vii). Hence $m_b(N)N^{b+1}$ is a finite sum, over set partitions and cyclic orders, of lattice sums of products of piecewise-polynomial span counts.

Lemma 10.1 (Model transfer). *For every b , $\lim_{N \rightarrow \infty} m_b(N)$ exists and equals the sine-model moment m_b defined by the overlap-weighted cluster calculus of §8.6.*

Proof. Each (partition, cyclic-order) term above is a lattice Riemann sum: rescaling $k = Nv$, the counts $(N - \text{span})_+/N$ converge to $(1 - \text{span})_+$ pointwise and uniformly on compacts (piecewise-linear functions with the same walls), the overlap weight confines all difference variables to the compact polytope $\text{span} \leq 1$, and the integrands are uniformly bounded; dominated convergence plus Riemann summation give convergence of each term to the corresponding continuum overlap-weighted box-convolution integral, which is term by term the defining F-CYC expression of the sine model (the sine-kernel cycles are the same Fourier box convolutions in the continuum). The terms are finitely many. Eight instances ($b \leq 8$) agree bitwise with the independently proved position-space constants. $\square$

Lemma 10.2 (Shift-lift: each cluster term is one lattice polytope). *The cluster expansion indexes the terms of $m_b(N) N^{b+1}$ by: a set partition $P = \{B\}$ of the b cycle edges (r blocks); for each block a merge partition Q_B of its slots; and a cyclic order σ_B on the classes of Q_B . The coefficient is the pure sign $\prod_B (-1)^{|Q_B|-1}$ (no factorials; the per-block count is $\sum_m S(|B|, m)(m-1)!$, and the total number of terms is the Fubini number of b). Write $k_i = m_i - m_{i+1 \bmod b}$ and, for $S \subseteq \{0, \dots, b-1\}$, $\partial S = \sum_{i \in S} (e_i - e_{i+1 \bmod b})$; let $S_1 \subset S_2 \subset \dots \subset S_{|Q_B|} = B$ be the nested unions of classes along σ_B . Then the term equals its sign times the number of lattice points $(\vec{m}, \vec{c}) \in \mathbb{Z}^{b+r}$ satisfying*

$$0 \leq m_i \leq N-1, \quad \partial B \cdot \vec{m} = 0 \quad (\forall B), \quad 0 \leq c_B + \partial S_j \cdot \vec{m} \leq N-1 \quad (\forall j),$$

the empty prefix included. The chain $\{S_j\}$ is a subchain of a full flag of σ -prefixes, so every structural statement proved below for full flags (total unimodularity, dimension, parity) covers the merged terms verbatim: a subchain only removes inequality rows. The finest layer $Q_B =$ singletons recovers the slot-order boxes; the assembly identity was verified bitwise at $b = 3, 4, 5$ (13/75/541 terms) against the archived exact tables.

Proof. At fixed $\vec{m}$ the count factorizes over blocks. For one block, the conservation delta of the cyclic cumulant is the equality row $\partial B \cdot \vec{m} = 0$; given it, the shift-count identity $\#\{c_B \in \mathbb{Z} : 0 \leq c_B + P_S \leq N-1 \ \forall S\} = (N - \text{span}_{\sigma_B}(k))_+$ reproduces the box count exactly (c_B ranges over an integer interval of length $N - \text{span}$ when positive). Singleton blocks reduce under the same formula to $N\delta_{k_i,0}$, the exact first cumulant — no special case. Note what the lift buys: the argmax / argmin switching walls of span, which are *interval* boundary rows and break the laminar structure, are absorbed by the c -coordinates; no chamber decomposition remains. $\square$

Lemma 10.3 (Total unimodularity of the lifted wall system). *For every $(P, \vec{\sigma})$, the constraint matrix of Lemma 10.2 — rows e_i , ∂B , and $e_{c_B} + \partial S$ over the coordinates $(\vec{m}, \vec{c})$ — is totally unimodular.*

Proof (Ghouila–Houri, top-anchored alternating signing). By the Ghouila–Houri row criterion it suffices to sign any subset R of rows so that the signed sum has all entries in $\{0, \pm 1\}$. The boundary rows of R split by block into nested chains (the selected prefixes of σ_B , with the plain row ∂B counted as a second copy of the full block). *Within a chain*, sort by inclusion and sign alternately *from the largest set down*, the largest taking $+1$; at the tie between the plain row ∂B and the lifted row $e_{c_B} + \partial B$, the plain row takes the earlier sign. Each edge e lies in an upper part of the chain, so the signed indicator sum $f_B = \sum \varepsilon_S \chi_S$ evaluates on each edge to a prefix sum of $+1, -1, +1, \dots$, i.e. $f_B \in \{0, 1\}^B$. *Across chains* the blocks are disjoint, so $f = \sum_B f_B \in \{0, 1\}^b$ globally, and the $\vec{m}$ -part of the signed sum of boundary rows is the cyclic difference Df with entries $f_j - f_{j-1} \in \{0, \pm 1\}$. For the c_B -entry: the signs of the lifted rows form the whole alternating sequence minus, possibly, the sign taken by the plain row; the alternating total lies in $\{0, 1\}$ and the plain row, being first at its tie (or the top of the chain), carries $+1$, so the c_B -entry lies in $\{-1, 0, 1\}$. Finally each selected unit row e_i is signed $-\text{sign}((Df)_i)$ (arbitrarily if zero). All entries of the signed sum lie in $\{0, \pm 1\}$. $\square$

Theorem 10.4 (Branch equality: one polynomial, all $N \geq 1$). *For every b , $m_b(N) N^{b+1}$ is a single polynomial in N of degree at most $b+1$, valid at every integer $N \geq 1$.*

Proof. Fix a $(P, \vec{\sigma})$ -term. Every constraint of Lemma 10.2 has the form $0 \leq \langle a, x \rangle \leq N-1$ or $\langle a, x \rangle = 0$, so the term equals $\text{ehr}_Q(N-1)$, the Ehrhart counting function of the *fixed* polytope

$Q = \{x : 0 \leq \langle a, x \rangle \leq 1, \partial B \cdot x = 0\}$ at dilation $N - 1$. By Lemma 10.3 and integrality of the right-hand data, every vertex of Q is integral: Q is a lattice polytope, so ehr_Q is a polynomial at every dilation $n \geq 0$ — no residue classes, no threshold. Its degree is $\dim Q$: the equality rows $\{\partial B\}$ have rank exactly $r - 1$ (the disjointly supported $\{\chi_B\}$ meet the kernel of the cyclic difference only in the constants, and $\sum_B \partial B = 0$), and the point $(\frac{1}{2}, \dots, \frac{1}{2})$ satisfies every inequality strictly, so $\dim Q = (b + r) - (r - 1) = b + 1$. Summing the finitely many terms with their N -independent Soshnikov coefficients gives the claim. $\square$

Theorem 10.5 (Parity). *Each $(P, \vec{\sigma})$ -term of Theorem 10.4, and hence the total, satisfies $L(-N) = (-1)^{b+1}L(N)$ as polynomials.*

Proof. The all-ones vector w (in $\vec{m}$ and $\vec{c}$ alike) satisfies $\langle a, w \rangle = 1$ for every inequality row — $\partial S \cdot \mathbf{1} = 0$ for boundary vectors, so units give 1 and lifted rows give c -component 1 — and $\partial B \cdot w = 0$. Hence $x \mapsto x - w$ bijects the relative-interior lattice points of nQ with the lattice points of $(n-2)Q$: $\text{ehr}_Q^\circ(n) = \text{ehr}_Q(n-2)$ (self-reciprocity). Ehrhart–Macdonald reciprocity gives $\text{ehr}_Q(-n) = (-1)^{\dim Q} \text{ehr}_Q^\circ(n)$ with $\dim Q = b + 1$; composing, the term $T(N) = \text{ehr}_Q(N - 1)$ satisfies $T(-N) = \text{ehr}_Q(-(N + 1)) = (-1)^{b+1} \text{ehr}_Q(N - 1) = (-1)^{b+1}T(N)$. The block count r has cancelled in $\dim Q = b + 1$: parity is the same for every partition shape, which is exactly the per-term localization observed in the true-cumulant scans (15/15 terms at $b = 4$, 52/52 at $b = 5$, zero violations). $\square$

Corollary 10.6 (The moment tower is proved). *By parity, $m_b(N)N^{b+1}$ has at most $\lfloor (b+1)/2 \rfloor + 1$ unknown coefficients. The archived exact values (proved rationals: Schur orthogonality, dual implementations) number at least that many at every $b \leq 14$, with strict surplus at every even b and, after the surplus extension points landed, at the odd b as well; the polynomial is therefore uniquely determined, and Lemma 10.1 identifies its leading coefficient with the sine-model moment. The values $m_9, \dots, m_{14}$ displayed above are exact rationals at full proof grade, and with them every constant consumed by the $k = 10, 12, 14$ certificates.*

Corollary 10.7 (Integrality, degree law, denominators). *For every $b \geq 1$ and $N \geq 1$: (i) $m_b(N)N^{b+1} \in \mathbb{Z}$; (ii) $m_b(N)$ is a polynomial in $1/N^2$ of degree at most $\lfloor b/2 \rfloor$ (attained at every computed order $b \leq 14$); (iii) the limit denominator satisfies $\text{den}(m_b) \mid (b+1)!$.*

Proof. (i) By Lemma 10.2 the object is a finite signed sum of lattice-point counts. (ii) Write $L_b(N) = m_b(N)N^{b+1}$. By Theorem 10.5 the nonzero coefficients of L_b sit in degrees $\equiv b + 1 \pmod{2}$. Moreover $L_b(0) = 0$: per term, $T(0) = \text{ehr}_Q(-1) = (-1)^{b+1} \text{ehr}_Q^\circ(1)$, and the relative interior of Q at dilation one contains no lattice point (any strict unit constraint $0 < m_i < 1$ is integrally empty). Hence the least degree present is ≥ 1 , and ≥ 2 when b is odd (parity), so $m_b(N) = L_b(N)/N^{b+1}$ is a polynomial in $1/N^2$ of degree at most $b/2$ for even b and $(b-1)/2$ for odd b . (iii) By (i) and $L_b(0) = 0$, the polynomial L_b of degree $\leq b + 1$ is integer-valued at all nonnegative integers, so its coordinates in the binomial basis $\binom{N}{0}, \dots, \binom{N}{b+1}$ are integers and its leading coefficient — which is m_b — lies in $\frac{1}{(b+1)!}\mathbb{Z}$. $\square$

The period-two behaviour of the isolated sectors observed in (vii) is a basis artifact — the doubled-coefficient wall $2u = N$ of the character decomposition — which the assembled polytope representation simply does not possess; sector cancellation is thereby explained, not merely observed. The proofs above were subjected to the program’s falsifier-first discipline before entering this paper: the signing construction of Lemma 10.3 was verified mechanically on all

row subsets of all 587 wall systems with $b \leq 5$; full-minor scans found zero violations over 21,171 deduplicated $\vec{m}$ -space systems ($b \leq 7$ exhaustive), 197,357 systems at $b = 8$ (two-leg sampled coverage, stated in the archive), and 42,271 lifted systems of Lemma 10.3 ($b \leq 7$ exhaustive, the 37,633 systems at $b = 7$ by GPU screening with exact integer Bareiss adjudication of every flagged minor); the complete merged term families of Lemma 10.2 — Fubini sizes 75/541/4,683 at $b = 4/5/6$ — passed single-polynomial fits with held-out points from $N=1$, per-term parity, and the signed assembly reproducing $m_b(N)N^{b+1}$ bitwise, by two independent implementations written from the coefficient specification alone (one per side, with a CPU/GPU cross-check at $b = 4$); the true-cumulant Möbius layer passed per-term parity on all 15/52/203 partitions at $b = 4/5/6$ with its own bitwise assembly gate; and the odd- b surplus points were pre-registered as predictions of the already-determined fits and then computed — $m_9(8)$, $m_{11}(9)$ and $m_{13}(10)$ hit bitwise (3,243,531 terms at $b=13$, $N=10$), giving every $b \leq 14$ table a fit surplus of at least one point. Σ_9 's facet-recursion second method has landed: $C_9 = \frac{27649}{302400}$ hit its pre-registered target bitwise (60,739 signed term orbits over 1,091,670 F-CYC terms, three-machine shards, zero-sum and rebuild self-checks; the target was pinned in advance by the $\{9\}$ joint identity with the already-certified $\{7, 2\} = -\frac{4313}{12600}$), so Σ_9 now carries two independent proofs — the branch-equality chain and the facet recursion; the formerly outstanding delegated dual-implementation points $m_{13}(4)$, $m_{14}(3)$ landed bitwise. None of this involves any hypothesis about ζ .

Theorem 10.8 ($k \leq 14$ tower; proved — see Lemma 10.1, Theorems 10.4, 10.5, Corollary 10.6). *At the moment data $(1, m_1, \dots, m_{14})$ listed above, the 8×8 Hankel matrix is positive definite together with its shift, the kernel polynomial Q_7^* has strictly alternating rational coefficients (Remark 11.2), and the Christoffel value is the exact rational*

$$\lambda_7(0) = \frac{352633869846878511557783511830740995191}{7876602339133293193971616991853147607579} = 0.04476979 \dots,$$

with zero rationalization premium. Consequently,

$$\liminf_{T \rightarrow \infty} \frac{N_0^s}{N} \geq 1 - 2\lambda_7 = \frac{7171334599439536170856049968191665617197}{7876602339133293193971616991853147607579} = 0.9104604105 \dots,$$

$$\liminf_{T \rightarrow \infty} \frac{N_d}{N} \geq 1 - \lambda_7 = \frac{7523968469286414682413833480022406612388}{7876602339133293193971616991853147607579} = 0.955230205 \dots,$$

and the intermediate rungs give $1 - 2\lambda_5 = 0.877197 \dots$, $1 - \lambda_5 = 0.938598 \dots$ ($k = 10$) and $1 - 2\lambda_6 = 0.896403 \dots$, $1 - \lambda_6 = 0.948201 \dots$ ($k = 12$), with $\lambda_5 = \frac{46970100247159}{764967228211380}$ and $\lambda_6 = \frac{13166900320841109317259245}{254195527518153210548497708}$.

The certificates were rehearsed blind: a pre-registered window (headline ≈ 0.9101 , distinct ≈ 0.9551 , computed from banded Σ_{13}, Σ_{14} before their exact values existed) was opened after the fact and hit. The same rehearsal quantified that band-grade consumption is worthless at this depth — a ± 0.004 band on Σ_{13} sweeps the headline by 7.6×10^{-2} — so exactness is not a luxury of the route but its point.

11 Consumption and the proof of Theorem 1.1

With $(m_1, \dots, m_4) = (1, \frac{4}{3}, 2, \frac{13}{4})$ pinned two-sidedly (Theorem 3.4, Theorem 7.5, (2.5)) and the certified values of §8.5 and §9, the remaining constraints are exact rationals:

$$m_5 \geq M_5 := \frac{101}{18}, \quad m_6 \leq M_6 := \frac{119}{12} + \Sigma_6 = \frac{640}{63}, \quad m_7 \geq M_7 := \frac{3439}{180}, \quad m_8 \leq M_8 := \frac{747361}{20160}.$$

Every band contribution is retired by the exact certifications, and the consumption is closed-form:

Lemma 11.1 (The consumption step is a Christoffel function). *Let ν be a probability measure on $[0, \infty)$ with Hankel matrix $(H_n)_{ij} = m_{i+j}$, $0 \leq i, j \leq n$, nonsingular. Then*

$$\nu(\{0\}) \leq \lambda_n(0) := \min\{\int Q^2 d\nu : \deg Q \leq n, Q(0) = 1\} = \frac{1}{e_0^\top H_n^{-1} e_0},$$

attained by the normalized kernel polynomial $Q^ = H_n^{-1}e_0/(e_0^\top H_n^{-1}e_0)$, whose coefficients are rational whenever the moments are. Moreover (sign lemma) Q^* has all roots in the interior of the support, hence alternating coefficients, so $P = (Q^*)^2 = \sum y_j x^j$ has $\text{sign}(y_j) = (-1)^j$: the certificate always consumes odd moments from below and even moments from above — exactly the one-sided data the ledgers deliver, with zero premium at every degree.*

Proof. $\nu(\{0\}) = \int 1_{\{0\}} d\nu \leq \int Q^2 d\nu$ for any Q with $Q(0) = 1$, since $Q^2 \geq 0$ and $Q(0)^2 = 1$; minimizing the quadratic form $q^\top H_n q$ over q_0 -normalized coefficient vectors gives the stated closed form by one Lagrange step, with minimizer $H_n^{-1}e_0$ normalized — manifestly rational. If $g(x) = \prod(x - r_i)$ with all $r_i > 0$ then the coefficients of g alternate strictly, so every product $g_a g_b$ contributing to $(g^2)_j$ carries the sign $(-1)^j$ with no cancellation; the roots of Q^* interlace the support by orthogonal-polynomial theory (they are the zeros of the reproducing kernel $K_n(x, 0)$), hence are positive. $\square$

Remark 11.2 (Moment-data form of the sign lemma). The sign lemma needs only the moment data: if the Hankel matrix and its shift (m_{i+j+1}) are positive definite — a Stieltjes condition verified for the tower in exact rational arithmetic (leading shifted minors $1, \frac{2}{9}, \frac{23}{1296}, \frac{27857}{45722880}$; gate F-STIELTJES in the certification chain) — then the orthogonal polynomials have positive zeros, and the x^m -coefficient of each summand of $K_n(x, 0) = \sum_j p_j(0)p_j(x)$ carries the sign $(-1)^m$ independently of j : the alternation of Q_n^* and of $(Q_n^*)^2$ holds term by term, with no cancellation, at every degree at once. The per-degree sign verifications of the certification chain are thereby cross-checks, not load-bearing steps, at every present and future order of the tower.

Remark 11.3 (The conversion coefficients are tight). The passage from spectral mass to zero counts uses the multiplicity ledger: an m -fold zero contributes an exact $(m-1)$ -dimensional kernel, whence $\text{distinct} \geq 1 - \lambda$ and, since $(m-1)/m \geq \frac{1}{2}$, $\text{simple} \geq 1 - 2\lambda$ — the factor 2 coming from the all-double extremal. This factor cannot be improved from the information the framework carries. Beyond the moments, the Nyquist diagonal collapse (§2.4) makes the normalized diagonal of $\tilde{G}$ equal to one at every zero, so each double contributes an exact $\begin{pmatrix} 1 & 1 \\ 1 & 1 \end{pmatrix}$ block; compressing onto the (mutually orthogonal) pair-sum vectors and applying the power-trace inequality with Cauchy interlacing yields the constraint family $\sum_{\text{top } p} \lambda_i^j \geq 2^j p$ for all $j \geq 1$, p the number of doubles. At the Christoffel extremal measure (atom λ_6 at the origin plus six Gauss nodes, computed exactly from the $k \leq 12$ moments) every member of this family is inactive — the top-quantile tail means exceed the thresholds by +5.6% ($j = 1$) through +43% ($j = 6$), with margins widening in j and in k — so the all-double configuration is consistent with the entire information set and the coefficient 2 is sharp. Any future improvement of the headline at fixed moment data must therefore come from the measure side (a smaller λ_k), not from the conversion layer; a paired window-family scan (eight directions, including the first-order m_2 -descent direction) likewise finds the flat window λ -optimal, so the frame's two design layers carry no residual slack.

Lemma 11.4 (Exact dual certificate, degree six). *At $(m_1, \dots, m_4)$ pinned and $(M_5, M_6) = (\frac{101}{18}, \frac{640}{63})$,*

$$Q_3^*(x) = 1 - \frac{8232}{2519}x + \frac{7368}{2519}x^2 - \frac{1932}{2519}x^3,$$

and for any probability measure ν on $[0, \infty)$ — no support bound — with $m_1 = 1$, $m_2 = \frac{4}{3}$, $m_3 = 2$, $m_4 = \frac{13}{4}$, $m_5 \geq M_5$, $m_6 \leq M_6$:

$$\nu(\{0\}) \leq \int (Q_3^*)^2 d\nu \leq \lambda_3(0) = \frac{247}{2519}, \quad 1 - 2\lambda_3 = \frac{2025}{2519}, \quad 1 - \lambda_3 = \frac{2272}{2519},$$

all exact: the one-sided substitutions are valid because the coefficients of $(Q_3^)^2$ alternate (Lemma 11.1), and the corner evaluation is a finite rational computation (kernel-checked in core Lean 4, `lean/RhGate/Certificate.lean`, and independently in `certification/certify84.py`). No solver, grid, atom extraction or rationalization occurs anywhere; the historical atoms of the earlier revision are, to four decimal places, the roots of Q_3^* 's numerator at the superseded M_6 (Appendix C(A6)).*

Corollary 11.5 (The exhausted $k \leq 6$ tower). *With ineffective constants, unconditionally at the grade of the $k \leq 6$ chain (every consumed constant symbolically certified),*

$$\liminf_{T \rightarrow \infty} \frac{N_0^s}{N} \geq \frac{2025}{2519}, \quad \liminf_{T \rightarrow \infty} \frac{N_d}{N} \geq \frac{2272}{2519},$$

and these are the exact optima of the six-moment data: no certificate consuming only $m_1, \dots, m_6$ can do better.

Lemma 11.6 (Exact dual certificate, degree eight). *At $(m_1, \dots, m_6)$ pinned as above and $(M_7, M_8) = (\frac{3439}{180}, \frac{747361}{20160})$, the degree-4 kernel polynomial Q_4^* has rational coefficients with alternating signs, and for any probability measure ν on $[0, \infty)$ with the stated one-sided data,*

$$\nu(\{0\}) \leq \lambda_4(0) = \frac{12241115}{162540559}, \quad 1 - 2\lambda_4 = \frac{138058329}{162540559} = 0.849377\dots, \quad 1 - \lambda_4 = \frac{150299444}{162540559}.$$

The corner is interior to the moment cone (vertex $m_7 = 19.0712\dots$, $m_8 = 36.7621\dots$ over the pinned $m_1, \dots, m_6$), and the payoff is continuous in (M_7, M_8) : in the ledger variables, the headline needs only $\Sigma_7 \geq 0.074844$ (margin 3.8% at $\Sigma_7 = \frac{7}{90}$) and $\Sigma_8 \leq 0.290506$ (margin 2.8% at the certified $\Sigma_8 = \frac{633}{2240}$), Σ_7 entering monotonically increasing despite the $+8\Sigma_7$ term in m_8 — so the two new sums are needed only one-sidedly and to percent-level accuracy.

Theorem 1.1 follows through the consumption mechanism of §2: Lemma 11.6 with Theorem 9.4 gives the headline at the stated grades; Corollary 11.5 is the fully-certified rung; m_4 alone gives 13/18, 31/36 (Corollary 7.7, certificate Q of Lemma 3.1, leading coefficient positive — also support-free). The Dirichlet L -function case is verbatim. All constants ineffective. $\square$

Corollary 11.7 (Sign-only rung). *If the three connected constants merely satisfy $C_5 \geq 0$, $\{4, 2\} \leq 0$ and $\{6\} \leq 0$ (with $C_5 \leq 0.1$), then, by the re-optimized exact certificate of Appendix C(A5),*

$$\liminf \frac{N_0^s}{N} \geq 0.7403, \quad \liminf \frac{N_d}{N} \geq 0.8701.$$

Sign certification is a far coarser target than the exact certification now in place (§8.5); the full payoff curve is tabulated in Appendix C(A5).

Remark 11.8 (The repaired intermediate rung; register D6). An earlier stage of this research quoted an intermediate rung (0.7295/0.8647) from the fourth- and fifth-moment chain alone, “consuming no sixth-moment information”. The certification audit found this rung unsound as stated: with only a *lower* bound on m_5 and no m_6 (or support) information, the origin-mass problem on unbounded support is not improved at all — mass $\varepsilon = c/\Lambda^4$ at height Λ boosts m_5 arbitrarily at vanishing cost to $m_1, \dots, m_4$, and the audit’s LP at support ranges 8/20/60/200 degraded $0.73144 \rightarrow 0.72580 \rightarrow 0.72340 \rightarrow 0.72257 \rightarrow \frac{13}{18}$. Any dual certificate for a lower fifth-moment bound needs a negative x^5 -coefficient dominated by a positive x^6 -coefficient, i.e. sixth-moment (or support) information is *necessary*. The rung is therefore demoted: it holds in the stated form only with the m_6 cap, in which case it is subsumed by Lemma 11.4. The full certificate is unaffected ($y_6 > 0$).

Merge-audit note on the band model. The audit first ran an *anti*-correlated variant (the δ of m_5 entering m_6 with opposite sign) and obtained 0.78972; the correlated form above is the correct accounting, since C_5 is a single unknown shifting both constraints together. Both runs are recorded in the register (D3) so the sensitivity of the headline to the band model is on the record: under the (incorrect, doubly conservative) anti-correlated accounting the headline would have read 0.7897/0.8949. With every band now retired and the inputs exact (§8.5), the accounting question no longer arises; the note is kept as part of the forward record.

12 The density-one limit: the tower’s supremum is one

Every even order $2k$ of the tower delivers the unconditional rung $\liminf N_0^s/N \geq 1 - 2\lambda_k(0)$, and Theorem 10.4 has made every further rung a polynomial-cost computation. This section proves that the supremum of the rungs is 1: the moment sequence (m_b) is determinate, its measure has no atom at zero, and the Christoffel values $\lambda_k(0)$ therefore decrease to zero. No rate is claimed — the statement is the capstone limit, not an effective schedule; the effective version is the hard-edge programme of §15. We first state the per-order input as a theorem in its own right (§12.1), then prove the soft trio and take the supremum (§12.2).

12.1 The order schema: every rung is a theorem

Theorem 12.7 consumes, for every k , the rung inequality at order k . Rather than leave this as a proof paragraph, we state it as a theorem and write out its proof from the order-uniform components, each of which is stated in this paper at general order; order enters only as a displayed parameter, and no uniformity in k is claimed or needed (Remark 12.8).

Lemma 12.1 (Pure-class transport at every order). *For every $b \geq 4$, the fully-locked b -cycle class aggregate transports to its model constant: in the consumed normalization,*

$$|\text{Agg}(\{b\}) - C_b \cdot O_b| \ll_{A,B} K J' V'^2 (\log X)^{-A} \quad \text{for every } A > 0,$$

where C_b is the b -chain model cluster constant — an exact rational at polynomial cost (Lemma 10.1, Theorem 10.4). For $b = 3$ the aggregate vanishes ($\Sigma_3 = 0$, the odd-block vanishing of §8).

Proof. This is the single-factor ($r = 1$) instance of the machinery of Theorem 9.2, with the factor the whole chain; we run the same four steps. (i) *Local law.* The mod- p density of the fully-locked b -chain is the coincidence-counting closed form of Proposition 6.1, stated for any

cycle length; no cross-block merger arises. (ii) *Archimedean layer*. The walk geometry O_b is the certified exact overlap layer (§8.3, §9.2), and the model value of the locked chain against it is the b -chain cluster (box-convolution) constant C_b of the F-CYC calculus — whose existence and exact rational value for *every* b is Lemma 10.1 with Theorem 10.4. (iii) *Variance transport*. By Proposition 8.1 (arbitrary per-lattice densities) the resolved variance is the chain autocorrelation summed over the common shift; the argument of Theorem 7.5 repeats verbatim with $b - 1$ classification indices (§8.4): the $(b - 1)$ -dimensional Parseval identity holds coordinatewise since each transform covers its full period, all-major bins are windowed single prime sums obeying Lemma 7.1, any minor factor is closed by Lemma 7.3, and the comb identification runs per coordinate by Lemma 7.4 with complete CRT orthogonality in each coordinate. (iv) *Truncation and constants*. The ℓ_1 -truncation discharges by Proposition 5.3 with exponent $P^{-(b-2)+\varepsilon}$, strengthening with b ; the (k, j) -aggregation constants and normalization interface are the archived certified layer of §4.2, independent of the chain length; the tail and threshold steps are order-independent (Proposition 2.3). All bounds hold two-sidedly; all constants depend only on (A, B) and the order b , and are ineffective through Siegel’s theorem. $\square$

Theorem 12.2 (Order schema). *For every $b \geq 2$, in the endpoint regime, the arithmetic trace moment converges two-sidedly to the model constant, $m_b(T) \rightarrow m_b$, at the grade of §14 layer (b). Consequently, for every $k \geq 2$,*

$$\liminf_{T \rightarrow \infty} \frac{N_0^s}{N} \geq 1 - 2\lambda_k(0), \quad \liminf_{T \rightarrow \infty} \frac{N_d}{N} \geq 1 - \lambda_k(0),$$

where $\lambda_k(0)$ is the Christoffel value of the moment data $(1, m_1, \dots, m_{2k})$. The implicit constants may depend on k and are ineffective.

Proof. Six components, each already stated at general order. (1) *Ledger decomposition (exact)*. The trace expansion of $\tilde{G}^b$ runs over the Bell ledger of coincidence classes; the re-centring identity (2.6) is an exact operator-binomial identity at every order, and every singleton-containing class factorizes exactly by W2 (Lemma 8.5, any b -cycle) and is absorbed, order by order, by the binomial coefficients. What remains are the singleton-free classes of orders $\leq b$. (2) *Three-blocks vanish*. A singleton-free class containing a 3-block carries, after the factorization of Theorem 9.2(i)–(ii), the 3-chain main term as a factor, and that constant is zero ($\Sigma_3 = 0$ exactly; arithmetic side classwise in §8, model face F-CYC). (3) *Multi-block classes* (all blocks of size ≥ 2 , none of size 3, $r \geq 2$): Theorem 9.2, proved at general b and $\mathcal{B}$. (4) *Pure classes*: Lemma 12.1. (5) *Model assembly*. The model-side class values assemble to Σ_b and, through (2.6), to m_b : this is the F-CYC cluster calculus, whose every constant exists and is an exact rational for every b (Lemma 10.1, Theorem 10.4); steps (2)–(4) equate the arithmetic aggregates with exactly these values, so the arithmetic moment converges to m_b , two-sidedly. (6) *Consumption*. The zero-side count is the Christoffel-function count of Proposition 2.4 at the moment data $(1, m_1, \dots, m_{2k})$, monotone in each even moment, with the order-independent tail and threshold devices of Proposition 2.3; at each fixed k it yields the displayed inequalities, exactly as instantiated at $k \leq 7$ in Theorems 1.1 and 10.8. No step depends on k beyond the displayed parameters. $\square$

Grading. Every analytic step above is the layer-(b) chain of §14 — the schema introduces no new grade and no new device; its content is that the set of consumed devices is finite and fixed, with the order entering only as a parameter. The instantiated rungs $k \leq 7$ remain the effective, fully written-out cases.

12.2 Determinacy, the atom, and the Christoffel limit

Lemma 12.3 (Moment growth). *For every $b \geq 1$ and every N , $m_b(N) \leq (b!)^2$; hence $m_b \leq (b!)^2$.*

Proof. Expand $\mathbb{E} \prod_i \text{tr } U^{k_i}$ over set partitions P of the b cycle slots into joint cumulants of the blocks. For a determinantal process the cumulant of a block B of trace statistics is a sum, over merge partitions Q of the slots of B , of cluster integrals whose truncated correlation functions are signed sums over the cyclic permutations of kernel chains [17]; for trace test functions each chain integral collapses, by the Fourier mechanism of §10, to a lattice box count lying in $[0, N]$, and the block vanishes unless its frequencies conserve ($\partial B \cdot \vec{m} = 0$). Three counts finish the proof. *Coefficient mass*: per block the number of (merge partition, cyclic permutation) pairs is at most $\text{Bell}(|B|)(|B| - 1)!$, and across blocks $\prod_B \text{Bell}(|B|) \leq \text{Bell}(b)$ (disjoint merges are particular merges), while $\sum_P \prod_B (|B| - 1)! = b!$ (permutations counted by cycle type); so the total coefficient mass is at most $\text{Bell}(b) \cdot b! \leq (b!)^2$. *Conservation dimension*: the $r - 1$ independent equality rows confine $\vec{m}$ to at most N^{b-r+1} lattice points (Lemma 10.3 makes the solution lattice coordinate-like; an elementary elimination also suffices). *Box factors*: each of the r blocks contributes a factor at most N . Assembling,

$$m_b(N) \leq \frac{N^{b-r+1} \cdot N^r}{N^{b+1}} \times (\text{coefficient mass}) \leq \text{Bell}(b) \cdot b! \leq (b!)^2.$$

(The archived exact values sit far below the bound — $m_{14} = 3141.5 \dots$ against $(14!)^2 \approx 7.6 \times 10^{21}$ — and the assembled coefficient mass itself is checked to lie below $(b!)^2$ with decreasing ratio for $b \leq 8$; the bound's only role is Carleman, for which factorial-squared growth is ample.) $\square$

Proposition 12.4 (Determinacy and weak convergence). *The sequence $(m_b)_{b \geq 0}$ is the moment sequence of a unique probability measure ν on $[0, \infty)$ (Stieltjes-determinate), and the expected spectral measures converge weakly, $\nu_N \Rightarrow \nu$.*

Proof. By Lemma 12.3, $\sum_b m_b^{-1/(2b)} \geq \sum_b ((b!)^2)^{-1/(2b)} \asymp \sum_b e/b = \infty$: the Stieltjes–Carleman condition holds, so the moment problem on $[0, \infty)$ has at most one solution. The measures ν_N are tight ($m_1(N) = 1$ exactly), and for every b the family x^b is uniformly ν_N -integrable ($m_{b+1}(N) \leq ((b+1)!)^2$ uniformly); by Lemma 10.1 every moment converges, $m_b(N) \rightarrow m_b$, so every subsequential weak limit of (ν_N) has moments (m_b) and by determinacy equals the same measure ν : the full sequence converges. $\square$

Proposition 12.5 (No atom at zero). $\nu(\{0\}) = 0$; indeed $\int \log^- x d\nu \leq 2 - \gamma$.

Proof. The finite- N log law of §8.6(vi) is exact: $\int \log x d\nu_N = H_N - 1 - \log N$. Since $\log^+ x \leq x$ and $m_1(N) = 1$,

$$\int \log^- x d\nu_N = \int \log^+ x d\nu_N - \int \log x d\nu_N \leq 1 + (1 + \log N - H_N) \leq 2 - \gamma$$

uniformly in N ($H_N - \log N \downarrow \gamma$). The function $\log^-$ is nonnegative and lower semicontinuous on $[0, \infty)$ (value $+\infty$ at 0), so by the Portmanteau theorem and Proposition 12.4, $\int \log^- x d\nu \leq \liminf_N \int \log^- x d\nu_N \leq 2 - \gamma < \infty$. An atom at 0 would make the integral infinite. $\square$

Lemma 12.6 (Christoffel limit [1, 16]). *For a determinate moment problem the Christoffel values decrease with k and converge to the point mass: $\lambda_k(0) \downarrow \nu(\{0\})$.*

Theorem 12.7 (Density one: almost all zeros are simple and on the critical line; grade of §14 layer (b)). *Unconditionally,*

$$\lim_{T \rightarrow \infty} \frac{N_0^s(T, 2T)}{N(T, 2T)} = 1, \quad \lim_{T \rightarrow \infty} \frac{N_d(T, 2T)}{N(T, 2T)} = 1,$$

and the same holds for the zeros of any fixed primitive Dirichlet L -function.

Proof. By Theorem 12.2, for every k , $\liminf_T N_0^s/N \geq 1 - 2\lambda_k(0)$. The left-hand side is a single number independent of k , so it dominates $\sup_k(1 - 2\lambda_k(0))$ — a supremum over countably many true statements, with no limit interchange (Remark 12.8). By Lemma 12.6 with Propositions 12.4 and 12.5, $\sup_k(1 - 2\lambda_k(0)) = 1 - 2\nu(\{0\}) = 1$. The identical argument applies to N_d with $1 - \lambda_k$, and to any fixed primitive Dirichlet L -function, whose chain is the same (§2, Theorem 12.2). $\square$

Remark 12.8 (Why no uniformity in k is required). The limit involves no interchange and no extrapolation, and it is worth displaying the logic in two lines. For each *fixed* k , the rung $\liminf_T N_0^s/N \geq 1 - 2\lambda_k(0)$ is a self-contained finite theorem: $T \rightarrow \infty$ is taken at fixed k , and its constants may depend on k (and are ineffective) without harm. The left-hand side is a single number independent of k ; it therefore dominates every rung, i.e. $\liminf_T N_0^s/N \geq \sup_k(1 - 2\lambda_k(0))$ — a supremum over countably many true statements, not a limit interchange, requiring no k -uniform error control. That this supremum equals 1 is Lemma 12.6 with Propositions 12.4 and 12.5 — determinacy, the exact logarithmic law, and classical moment theory — in which no finite-order numerical value appears: the archived moment tables play no role in the proof of Theorem 12.7, only in its effective instantiated rungs.

Remark 12.9 (What is and is not claimed). The theorem is a density statement with no effective rate: it exhibits no $k(\varepsilon)$ beyond the monotone scheme, and it does not assert that all zeros are simple or on the line. The effective schedule — how large k must be for $1 - \varepsilon$ — is governed by the decay of $\lambda_k(0)$, empirically $\lambda_k(0) \approx 0.33/k$ (§15); making that rate a theorem is the hard-edge programme, which the present section deliberately bypasses: determinacy, the exact log law, and reciprocity of the box-count polytopes are soft inputs, all proved above. The analytic layer consumed is that of §14 layer (b), and the theorem’s grade is exactly the grade of that layer; no hypothesis about ζ enters.

13 Numerical verification

The standalone pipeline (`pipeline/`; `numpy`, `scipy`) reruns every gate in a few minutes. Central height table (sixteen coprime modulus families):

T	2400	4800	9600	19200	38400	57600	76800	153600
dispersion (10^{-3})	8.39	5.19	4.16	4.06	3.87	1.84	1.29	1.04
tracking (%)	1.5	1.0	0.7	0.9	0.8	0.4	0.2	0.2

(full-range geometric mean 0.706 per doubling: the pure averaging law $2^{-1/2}$). Further faces: local closure exact (zero violations, seven configurations); Parseval exact; per-bin Siegel–Walfisz face falling tenfold over the range; Bell(5)/Bell(6) gates exact; five- and six-fold product faces green at all heights; well-spacing zero violations; hybrid size control $24\text{--}31 \times$ height-stable; consumption LP grid-stable. The registry: 571 pre-registered checks, 529 passed, 61 fired and

converted, all forward-recorded (the 31st–35th conversions belong to the corrected-constants pass, the Lean compiles and the spectral-route audit; the 36th–41st to the model-edge audit, the $\{2, 2, 2\}$ classification correction and its aftermath, and the C_8 candidate miss, §14(f), Appendix B, D19–D22).

The certification audit of August 16, 2026 re-ran: the full moment-chain gate suite at $T = 9600$ (ALL PASS, 72s: dispersion 4.163×10^{-3} , tracking median 0.73%, SW RMS 1.599×10^{-3} , tightness 24.8–31.3 $\times$, five-fold 0.06–1.29%, six-fold 0.01–0.73%, LP 0.79454/0.89727 at two grids under the since-retired endpoint constants; the corrected-constants rerun gives 0.79472/0.89736); the fourth-moment identity suite at $T = 9600$ (ALL PASS, 48s); the constant suite of §4 (`m1.suite.py validate`: 6/6 local factors, 100/100 vector sieve, anchors and cross-convention difference reproduced; core -0.00545 against archive -0.00544 ; α -truncation -0.00491); the sawtooth certification (`sawtooth.cert`: grid 1.8675, samples 1.8419/1.8494/1.8512, certificate 2.100); the continuum quadrature gates (`quadrature.cert`: geometric tail ratio 0.500 certified, $c_0 = -1.4228$ decade-stable); and the certificate suite of §3 (`certificate_verification.py` exact; `moment_lp_reopt` 0.0721; `mu3_ledger_gates`; `cumulant_engine` gates C_2 exact, $\int O_3 C_3 = 0$). An independent re-implementation of the consumption LP (different code, HiGHS backend) reproduced every rung as recorded in §11.

A second certification pass (same date) added: the exact pairing-layer integrals (`exact_t222.py`: all values of Proposition 8.2, dual-scheme identical rationals, the closed-form cross-checks, and the match of the conjectured rationals); the exact dual certificate (`certify_lp.py`: expansion, both band corners, final inequalities, all exact); the write-out faces (`writeouts_verification.py`: free-slot Fubini exact, window mass 1.00041, triple-pair product model median 0.9814, resolved Parseval 1.2×10^{-16} , resolved CRT contrast — one face fired during construction and was converted, the registry’s 30th); and the support-range probe behind Remark 11.8.

The corrected-constants pass re-ran: the grouped midpoint-grid ladders for C_5 , $\{4, 2\}$, $\{6\}$ (§8.6; checkpointed engines, $b = 4$ calibration exact against Proposition 8.2); the exact dual certificate at the corrected corners (`certify_lp.py`: expansion, both band corners, final inequalities, all exact, worst corner $w_0 = \frac{1153107070889}{11233957316589}$); the convolution trace gates (F-TRACE: $b = 2$ exact $\frac{1}{3}$, $b = 3, 4$ to 3×10^{-14}); and — for the first time in the program — a genuine Lean kernel compile of the identity layer on the maintainers’ toolchain (§14(f)).

The identification pass (August 17) ran the accelerated ladder engine (support pruning $171 \times$ at $b = 6$, term trie, fused kernels, negation symmetry — mathematically identical to the reference engines, validated gate-by-gate against them and bitwise across V100/consumer-GPU/256-thread-CPU paths) to seven Romberg rungs per constant: C_5 to $dv = 0.00125$ (3200^4 grid), $\{6\}$ to $dv = 0.0125$ (320^5), the $\{4, 2\}$ spectator-pair variant to $dv = 0.0015625$ (2560^4 ; total wall 234s), C_7 to $dv = 0.025$ (160^6). Three prediction gates pre-registered before the fine runs (F-RAT, from h^2 -model fits of the coarse rungs) were all confirmed — for $\{4, 2\}$ the fine ladder *hit the pre-registered fraction* $-\frac{23}{420}$ exactly; the exact anchors (0 at $b = 3$, $-\frac{1}{60}$ at $b = 4$) were reproduced by the identical pipeline at 10^{-18} in the same runs.

The symbolic-certification pass (August 17) closed N1: every term of the three constants lifted to a rational polytope volume and evaluated by the exact facet recursion of §8.5, 1310 exact fractions archived per term, the full $\{6\}$ sum in 9.1 single-core hours. (The piecewise route, retained as the independent second method, projects to $\sim 10^8$ hours on the same $\{6\}$ term set — the polytope lift is a $\sim 10^9$ -fold algorithmic reduction at $b = 6$.) The pre-registered falsifier F-SYM-N1 — any mismatch between a symbolic sum and its identified fraction — did not fire: the symbolic values agree bitwise with the numerical reconstruction on all three

constants and on all three $\{4, 2\}$ placement components.

Remark 13.1 (Effective faces, ineffective theorem). The faces are finite-height measurements with effective constants; the theorem’s constants are ineffective (Siegel). The faces verify the exact-identity layer, the model identifications, and the qualitative averaging behaviour the proof predicts; they cannot and do not verify the asymptotics. The two kinds of evidence are kept distinct throughout.

14 Verification status

Claim grade: certified-candidate. The chain = classical inputs (Siegel–Walfisz [7, 10], Vaughan, singular-series truncation, [12] for §4) + archived certified lemmas (bridge, glue, aggregation, D1 truncation, consumer interface, anchor certificates: §4) + the machine-verified identities of §§3, 6–8 + the symbolically certified connected constants of §8.5 + the consumption layer of §2. Grades by layer:

Reading guide (headline theorems, what each consumes, at what grade).

- Theorem 1.1 (0.8493/0.9246): exact fourth moment (layer (b) analytic chain) + symbolically certified $k \leq 8$ constants (layer (a)) + consumption LP (exact). Independent of §4 and of §10.
- Theorem 10.8 (0.9104/0.9552): the above + the Toeplitz–trace moments $m_9..m_{14}$, now proved exact rationals (branch-equality and parity theorems, §10; finite tables at proof grade with dual implementations and pre-registered holdouts).
- Theorem 12.7 (density one): the order-uniform arithmetic devices (layer (b)) + the soft trio of §12 (determinacy, no atom, Christoffel limit — proved) + no numerical table and no rate.
- Theorem 4.1 (0.6916/0.8458): the one computer-assisted rung (Lemma 4.10 numerical-analysis step); quarantined — consumed by nothing outside §4.

(a) *Machine-exact or exact-rational, re-verified in the certification audit.* The certificate algebra (Lemma 3.1) and $\kappa(c)$ laws; the local closure (Proposition 6.1); the Parseval and product identities; the Bell ledgers and frozen-slot identities; the anchor-free 67/12 and 39/4; the 197/60 ledger arithmetic and the now fully exact two-bookkeeping reconciliation (Remark 3.6); the exact pairing-layer constants $t_{\text{adj}} = \frac{7}{60}$, $t_{\text{opp}} = \frac{1}{30}$, $\{2, 2, 2\} = \frac{32}{105}$ by dihedral placement class (Proposition 8.2, Appendix C); the symbolically certified connected constants $C_5 = \frac{1}{36}$, $\{4, 2\} = -\frac{23}{420}$, $\{6\} = -\frac{1}{126}$ (§8.5; register D18); the exact dual certificate (Lemma 11.4) — the consumption step involves no LP solver, no grid, and no support bound; the sawtooth constant $C_{\text{saw}} = 2.10$ (grid-plus-analytic certification).

(b) *Certified-candidate analytic layer (pre-Lean, pre-review).* The nine-class μ_3 ledger (two nonstandard steps written out in the archived memo); the truncation discharge; the spectral proof of Lemma D (Theorem 7.5) — whose minor/mixed-bin size bookkeeping is classical in kind but carried at write-out grade with face-checked budgets — and its factor-wise transports to $k = 5, 6$, including the written-out W1–W3 (§8.4), which are proved at this layer’s grade; the covered-zone/band/glue chain of §4. Five adversarial audit rounds precede the frozen chain (archive rounds r93, r95, r100, r102).

(c) *Write-out items.* Discharged: W1 (Proposition 8.7), W2 (Lemma 8.5), W3 (Lemma 8.6), each with a pre-registered face. W2 is complete (elementary); W1 and W3 are proved at the grade of layer (b), since they consume the same spectral transports.

(d) *Certified values.* N1 status: *closed*. All connected constants are proved rational (§8.6), and the three consumed constants are *proved* — each an exact-rational polytope-volume computation with per-term archives, dual-method fraction-by-fraction agreement, and all calibration gates passed (§8.5) — consumed at proof grade, the grade of Proposition 8.2. The numerical identification (machine-precision Romberg limits with adjacent-row drift 2×10^{-16} , uniqueness-guaranteed rational reconstruction, double exact-anchor calibration in the same pipeline, cross-architecture bitwise agreement, pre-registered h^2 prediction gates, and for $\{4, 2\}$ a pre-registration hit with componentwise identification) stands as an independent confirmation. The $k = 7, 8$ layer (§9) is graded item by item: $\{5, 2\}$, $\{2^4\}$, C_7 , C_8 at proof grade — $\{5, 2\}$, $\{2^4\}$ and C_7 with two independent confirmations each (pre-registration hits; 17/17 ladder agreement; exact hit of the registered identification), C_8 by the orbit-reduced symbolic run with model-side agreement to 10^{-4} ; the D_8 symmetry-premise audit is *complete* (six orbits, three members each, identical rationals), and the raw-term brute-force replication was deliberately stopped to reallocate compute to the ninth-moment layer — a recorded decision: C_8 ’s pedigree is the symbolic run plus the model-side A/B/C checks, not a second independent symbolic path); $\{2, 2, 4\}$, $\{4, 4\}$, $\{6, 2\}$ at proof grade with their second-path GPU ladders *returned and passed* (gate F-V-JOINTS: 22/22, 7/7, 4/4 orbits within the revised 3×10^{-6} band; a first comparison had mismatched two classes through ordinal orbit alignment, caught by its own pattern and repaired with canonical-form keys — the lesson, *content-canonical keys and asserted match counts for every cross-toolchain comparison*, is forward-recorded) — the constant layer of the $k \leq 8$ chain is complete at dual-path proof grade. The headline of Theorem 1.1 carries the weakest of these grades together with the layer-(b) grade of the W_∞ transport (Theorem 9.2, instances W4–W7); Corollary 11.5 consumes none of them and stands at the full $k \leq 6$ certification grade. The *certificate family* of Appendix C(A5) remains as the graded plug-in record. N2 (partially discharged): the quadrature’s slope is proved exactly 1, its constant is identified in closed form, $c_0 = \gamma - 2$ (Proposition C.1), and the geometric tail was already certified (r60); what remains computed is the mean oscillation $\overline{\text{osc}}$ inside c^* and the universality-collapse rate write-out. Only the conservative charge ε_{CL} is consumed anywhere (Lemma 4.10). Full certification of N2 upgrades Theorem 4.1 to 0.7031/0.8515.

(e) *Imports.* [5]: the framework of §2 (matrix, block structure, tail, counting, first two moments) — a publicly posted preprint of this program, not independently peer-reviewed; its own analytic imports are published literature. Directly consumed published literature: [20] and [10, Thm. 5.12] (explicit formula), [19, Thm. 9.2] (zero density), [13, 2] (prime-side pair-correlation evaluation, unconditional in the band), [12], [9], [7].

(f) *Formalization status.* The identity layer has now been *compiled and kernel-checked* on the maintainers’ machine: core Lean 4 (toolchain `leanprover/lean4:v4.33.0`, no `mathlib`), package `lean/RhGate/`, build log included. Compiled content: the anchor-free assembly identities $m_4 = \frac{13}{4}$, m_5 -part = $\frac{67}{12}$, m_6 pair-layer = $\frac{39}{4}$ as universally quantified rational identities (any anchors t_a, t_o — `Anchor.s.lean`); the consumption arithmetic $\frac{13}{18}, \frac{31}{36}$; and the coincidence-counting local law at $b = 4, p = 5$ by full kernel enumeration of all 125 lock classes (`LocalClosure.lean` — this theorem is *axiom-free*; the others use only `propext`, `Classical.choice`, `Quot.sound`). Zero `sorry`s. The compile itself fired a registered check and became the 33rd conversion: the seed files’ proofs invoked `ring` and `decide` on $\mathbb{Q}$, neither available in core Lean (`ring` is a `mathlib`

tactic; core `Rat` arithmetic is not kernel-reducible, so `decide` stalls on `instDecidableEqRat`); all five affected proofs were repaired with the core `grind` tactic, statements unchanged (Appendix B). A second same-machine compile (August 16, 20:09) covered the certificate layer: `Certificate.lean` (the denominator-cleared expansion identity of Lemma 11.4, nonnegativity, normalization, both corner evaluations of w_0 , the corner-selection sign — so the *choice* of binding corner is itself kernel-checked — the headline inequalities, the monic expansion of Appendix C(A6), the exact $\{2, 2, 2\}$ arithmetic and the instantiated anchor identities; a third same-machine compile (August 17) covers the previous revision — `Certificate.lean` at $(M_5 = \frac{101}{18}, M_6 = \frac{12809}{1260}, \text{headlines } \frac{7962}{10000} / \frac{8981}{10000})$, thirteen certificate-layer theorems, clean build, axiom audit as stated; the present revision’s certificate layer — the corrected $\{2, 2, 2\}$ bundle, $M_6 = \frac{640}{63}$, the rational kernel polynomials Q_3^*, Q_4^* and the exact headlines $\frac{2025}{2519}, \frac{138058329}{162540559}$ — is written (`Certificate84.lean`) and *queued for the same-machine recompile*) and `LocalClosure2.lean` (the local law at $b = 4, p = 7$ and $b = 5, p = 5$: 343 and 625 lock classes by kernel `decide` on $\mathbb{N}$ — both *axiom-free*). This compile fired the registry a second time (the 34th conversion): `grind`’s ring/linear core treats a square $(\cdot)^2$ as an opaque atom and cannot close perfect-square nonnegativity; the repair is the *core* ordered-ring lemma `Lean.Grind.OrderedRing.sq_nonneg` (term-mode, statement unchanged, still *mathlib-free*). The pre-compile seed `SOSCertificate.lean` — whose headline constants carried the retired endpoint-grid corner — is removed from the package to preclude coexisting 0.7946/0.7947 headline theorems; its still-valid content is absorbed into `Certificate.lean`, and the draft is preserved in the program archive. Every statement is proxy-verified in exact rational arithmetic (`certify_lp.py`) or full `Python` enumeration (zero deviations). What Lean certifies is the rational-arithmetic layer alone: the moment pinning, the measure-theoretic Chebyshev–Markov step and the analytic chain remain graded as in layers (b)–(e). The truncation, spectral and consumption *analysis* remains unformalized; formalization of the analytic layer — which relies on the framework of [5] — is graded pending. The formalization package accompanying [5] is separate from this repository; the compiled coverage of this paper is exactly as stated above.

(g) *The gate.* Per the program’s verification gate, the record claim follows (i) the symbolic verification of the N1 fractions and the $\{4, 2\}$ closure — *now discharged* (§8.5, (d) above) — and optionally N2, (ii) compiled Lean formalization of the identity and certificate layers with the queue’s Q1–Q6 (both are now compiled, §14(f); the analytic layer remains), and (iii) external review. No Riemann Hypothesis, zero-density input, pair-correlation conjecture, or Kloosterman technology is used anywhere. Nothing here bears on RH itself.

15 Outlook: rates, structure, and the wall

With the tower run to its supremum (Theorem 12.7) the outlook is no longer about higher proportions — that question is closed. What remains is of three kinds: *effectivity* (the rate at which the instantiated rungs approach the limit), *structure* (closed forms for the measure ν and its constants), and *the wall* (what separates all of this from the Riemann Hypothesis, mapped precisely at the end of item (c)). We record each with the quantitative evidence in hand.

(a) *Closing the present chain.* Of the designated closures, the first — symbolic verification of the three connected fractions, moving item N1 from identification to proof — is carried out in this paper (§8.5); what remains is the formalization of the analytic layer and external review. These change grades, not constants.

(b) *Closing the $k = 8$ grades, and the ninth and tenth moments.* The $k = 8$ rung is closed

at the grades stated in §14(d); what remains there is the Lean extension and review. The reconnaissance of the ninth and tenth layers is kept below as it stood — it illustrates the program’s measurement-before-proof pipeline — and the paragraph closes with where those layers now stand. On the exact side, five of the eleven singleton-free classes of the ninth- and tenth-order ledgers are certified:

$$\{2^5\} = \frac{10531}{13860}, \quad \{4, 2, 2, 2\} = -\frac{208771}{498960}, \quad \{2, 2, 5\} = \frac{2263}{5040}, \quad \{5, 4\} = -\frac{257}{10080}, \quad \{6, 2, 2\} = -\frac{120139}{498960},$$

the first now carrying a full per-orbit second-path ladder (79/79 orbits within 2.1×10^{-6} ; its largest orbit is the blind-predicted $\frac{73}{11088}$ of §8.6(v)), the second returned identically by two hosts with different parallelizations, the third still on a single certified path (its ladder is queued). The closed form of §8.6(v) also *pre-registers* two rational targets in the twelfth and fourteenth ledgers — the sequential orbits $\frac{523}{192192}$ and $\frac{119}{102960}$, the family’s first denominators divisible by 13, exactly as the B-spline mechanism of (e) predicts — awaiting independent facet-recursion confirmation (gates F-SEQ-12/14, forward-recorded). On the measurement side, a high-statistics run of the exact finite model ($N = 128$, 2×10^5 CUE samples; the six known central moments $\Sigma_2, \dots, \Sigma_8$ all reproduced within 0.4 standard errors, the pre-assigned Σ_7 gate at +0.1) gives

$$\Sigma_9 = 0.172618(61), \quad \Sigma_{10} = 0.400778(108),$$

and the re-centring closed forms $m_9 = \frac{495107}{6720} + \Sigma_9$, $m_{10} = \frac{199427}{1344} + 10\Sigma_9 + \Sigma_{10}$ are reproduced by the independent Bell(9)/Bell(10) ledger assemblies exactly (gate F-PRED-RAT). The Toeplitz–trace engine has since closed all six further layers at proof grade — the values, evidence matrices, the $k \leq 14$ certificates and the branch-equality and parity theorems that ground them are the content of §10; the delegated dual-implementation points $m_{13}(4)$ and $m_{14}(3)$ have landed bitwise. What remains on this front is redundancy, not proof: the C_9 facet recursion has landed on its pre-registered target $\frac{27649}{302400}$ bitwise (§10), giving Σ_9 a second independent proof; the tenth- and twelfth-layer joint classes remain as further second methods for already-proved constants.

(c) *The tower’s horizon.* The moment route cannot stop at any fixed proportion — level repulsion forces $w_0(k) \rightarrow 0$ as all moments are pinned — but the marginal gain per rung is now governed by the mismatch between the LP-extremal’s implied moments and the sine-model moments, which the exact data now quantify precisely: the exact Christoffel sequence gives $n\lambda_n(0)$ increasing through $7\lambda_7 = 0.3134$ with two-point extrapolant $A \approx 0.330$, $B \approx 0.114$ in $n\lambda_n \approx A - B/n$ — a value in tension with the earlier pool-measured $A \approx 0.30$ at $n = 10..15$, most plausibly because those measurements carried the finite- N moment bias that is now an exact theorem; the refreshed band-level pricing is $k = 16 \rightarrow \approx 0.921$, $k = 18 \rightarrow \approx 0.930$, with 0.95 near $k \approx 26$ and 0.99 near $k \approx 120$ on the extrapolated law. Since Theorems 10.4 and 10.5 hold at *every* b , any future rung is born at proof grade: it needs only $\lfloor (b+1)/2 \rfloor + 1$ exact finite- N evaluations — polynomial cost — rather than a new structural argument. With Theorem 12.7 the tower’s destination is no longer in question — only its schedule: what remains open here is the *rate* $\lambda_k(0) \approx 0.33/k$, i.e. the hard-edge density of ν at 0 and the Máté–Nevai–Totik constant, whose proof route is the resummation programme below. A structural hint, recorded for that programme: the logarithmic spectral law integrates to the cumulant generating function $\Lambda(s) = (1+s)\log(1+s) - s - \log \Gamma(2+s)$, and $e^{\Lambda(b)} = (1+b)^{1+b}e^{-b}/(1+b)!$ is exactly the size-biased total-progeny (Borel) sequence of the *critical* Poisson Galton–Watson tree — the log-spectral side of ν is a critical-branching object, suggesting an exact identification of ν within that family.

The interpolation map, and what this paper does not do. It is worth recording exactly where Theorem 12.7 sits relative to the Riemann Hypothesis, because the framework makes the distance

precise. The method has two dials, the bandwidth λ of the consumed prime-side data (in Montgomery’s normalization) and the moment order k : ($\lambda=1, k=2$) yields $2/3$; ($\lambda=1, k=14$) yields 0.9104 ; ($\lambda=1, k=\infty$) is Theorem 12.7 — density one, the *supremum of the band-limited averaged world*, now attained; and Weil’s criterion (positivity for *all* test functions, individually) is the Riemann Hypothesis itself. The remaining distance is not a percentage but a change in the kind of input, and the framework locates it exactly. A zero at distance $\delta > 0$ off the line enters band-limited forms through its functional-equation mirror pair, and a three-lines (log-convexity) constraint suppresses its effect in single-scale $g \star \bar{g}$ forms; its full exponential weight $T^{2\lambda\delta}$ is recovered only by *two-scale* test structures — which are precisely the $\alpha > 1$ regime of the pair-correlation function $F(\alpha)$, i.e. Hardy–Littlewood prime-pair information. Under the height averaging that makes such data unconditionally evaluable, the off-line signal is *linear* in the number of off-line zeros: absence of signal at bandwidth λ translates into zero-density exponents of the shape $N_{\text{off}}(\delta, T) \ll T^{1-2\lambda\delta+\varepsilon}$, never into emptiness at any finite λ — rarity weakens the averaged signal proportionally, so no sparsity leverage exists inside the averaged world. The Riemann Hypothesis is exactly the $\lambda \rightarrow \infty$ corner, equivalently per-height (unaveraged) prime data at sub-spacing resolution. No statement in this paper crosses that wall, and no within-band improvement can; we record the map so that progress inside the band — including the density-one capstone — is not misread as progress toward it.

The per- k bookkeeping is now automated on *both* sides: the arithmetic side once and for all by Theorem 9.2 (every further order is an instance), the model side by the finite- N route of §10 at polynomial cost. A historical note, kept because it locates a still-open structural question: the position-space orbit-reduced integrator, which produced the $b \leq 8$ constants, has cost growing $\sim 11\times$ per level and reached its economic boundary near $b = 10$; the intrinsic complexity parameter of a class there is the *crossing width* of its diagram (non-crossing sectors collapse to the closed forms of §8.6(v), while 87% of the C_8 terms have width ≥ 4). That cost wall has been *replaced*, not climbed: Theorems 10.4 and 10.5 route every order through $\lfloor (b+1)/2 \rfloor + 1$ exact finite- N evaluations, and the position-space route survives as an independent second method (as instantiated by the C_9 facet recursion). The crossing-width analysis retains interest solely for the closed-form question of item (e).

(d) *Beyond moments.* Moments are global functionals, while the mass at the origin is killed by a local property (level repulsion: spacing density $\sim x^2$). A consumption mechanism accepting localized test functions — a Christoffel-type dual at non-polynomial functions with arithmetically computable values — would argue below the resolution at which the averaged world operates; in the language of the interpolation map of (c), it would be a step off the (λ, k) plane rather than along it. Whether the compressed-matrix framework admits one is, in our view, the most consequential open question raised by this work. Any future mechanism of this kind that argues at the level of *individual* zeros must pass a discriminator we record here as a permanent pre-registered gate: it must demonstrably fail for the Davenport–Heilbronn function (which shares the functional equation and zero symmetries but violates the analogue of RH) — an argument that cannot distinguish ζ from Davenport–Heilbronn consumes no arithmetic and proves nothing.

(e) *A structural question.* Every constant certified or identified by the program — $\frac{1}{3}, \frac{7}{60}, \frac{1}{30}, -\frac{1}{60}, \frac{32}{105}, \frac{1}{36}, -\frac{23}{420}, -\frac{1}{126}, -\frac{17}{360}$ — is a rational with denominator dividing $2520 = \text{lcm}(1..9)$; at level 8 the law fails ($3780 \nmid 2520$ for $\{2^4\}$, orbit denominators up to 90720), and at level 10 the prime 11 appears ($\{2^5\} = \frac{10531}{13860}$, $\{4, 2, 2, 2\} = -\frac{208771}{498960}$, both with $11 \mid$ denominator) — so the correct invariant is level-dependent. Corollary 10.7(iii) now gives the proved ceiling

$\text{den}(m_b) \mid (b+1)!$ at every order; the remaining question is why the observed denominators are so much smaller than that ceiling $((15)!/\text{den}(m_{14}) = 48$, but the per-class denominators are smaller still). The non-crossing calculus of §8.6(v) explains the mechanism: the building blocks are central B-spline values ($\int \text{sinc}^{2m} = \beta_{2m-1}(0)$ and relatives), whose factorial-family denominators grow with the level exactly as observed. A structure-probe on the nine exact moments (over-determined P-recursive and low-degree algebraic fits, all excluded; Hankel determinants without product structure; λ -denominators splitting into large primes) rules out the cheapest integrable families for the spectral measure itself. The dimension theorem of §8.6 explains rationality but not the small denominators; a closed-form theory of these polytope volumes (producing all C_b at once) would collapse the model side of every future rung.

(f) *The diagonal limit, achieved.* Earlier revisions of this program recorded here a reduction: “one hundred percent requires every rung at proof grade plus the origin mass of ν .” That reduction has been discharged — and, instructively, not along the route it anticipated. The anticipated route priced the limit through the cost wall of the orbit integrator; the actual proof (§12) needed *no* computed constant at all: determinacy (Lemma 12.3, without any compactness assumption on ν), the exact logarithmic law, and the classical Christoffel limit close the supremum softly, while the finite rungs remain as its effective instantiations. What survives of the old item is its boundary sentence, which we restate: Theorem 12.7 is a density-one statement, not the Riemann Hypothesis, and not a statement about every zero; the precise location of the remaining wall is the interpolation map of item (c).

Acknowledgements

The mathematical development of this paper — the framework of §2, the certificate calculus, the spectral proof of Lemma D, the moment identifications, the convolution calculus, and the reproduction and Lean packages — was carried out by Claude (Anthropic), a large language model, working within a research program directed by the authors under the falsifier-first verification discipline described in §14 and Appendix B. The authors specified objectives, arbitrated decisions, compiled and audited the Lean package on their own toolchain, and take full academic responsibility for the content. This paper is built on the two-thirds theorem [5], a preprint of the same research program.

A The comb identification, with a worked pair

In the τ -frame the model of $\rho(\tau)$ is assembled from twin singular series: with $\mathfrak{S}$ the Hardy–Littlewood series $\mathfrak{S}(h) = 2C_2 \prod_{p|h, p>2} \frac{p-1}{p-2}$ (h even), window overlap counts V_m, V_n , and exact diagonal pieces,

$$\text{Model}(\tau) = \sum_{\substack{t \equiv 0 \pmod{b_2} \\ t \equiv \tau \pmod{b_1} \\ t \notin \{0, \tau\}}} \mathfrak{S}\left(\frac{t}{b_2}\right) \mathfrak{S}\left(\frac{t-\tau}{b_1}\right) V_m\left(\frac{t}{b_2}\right) V_n\left(\frac{t-\tau}{b_1}\right) + \text{diagonals}.$$

The CRT progression $t \equiv 0 \pmod{b_2}$, $t \equiv \tau \pmod{b_1}$ is the *complete* orthogonality of the joint frame: a shift t contributes iff it is simultaneously a b_2 -lattice twin shift on the m -side and, offset by τ , a b_1 -lattice twin shift on the n -side; no incomplete sums occur because the τ -transform runs

over the full period. The local factors of $\mathfrak{S}$ at the primes $p \mid b_1 b_2$ are exactly the κ -values of Proposition 6.1 at $b = 4$; for $p \nmid b_1 b_2$ they are the twin factors.

Worked pair $(b_1, b_2) = (5, 3)$: $M = 15$, and for $\tau \equiv \tau_0$ the surviving shifts are the single residue class $t \equiv t_0(\tau)$ (15) with $t_0 = 3 \cdot (2\tau \bmod 5) \bmod 15$ (here $3^{-1} \equiv 2 \pmod{5}$). For example $\tau = 1$ gives $t \equiv 6 \pmod{15}$: the model sums $\mathfrak{S}(t/3)\mathfrak{S}((t-1)/5)V_m V_n$ over $t = 6, 21, 36, \dots$ and the negative branch, plus the diagonal $t = \tau$ term when $3 \mid \tau$ and $t = 0$ when $5 \mid \tau$ (neither here). This is precisely what `face_dispersion` constructs; its agreement with the measured profile at 0.2–4% across all sixteen families and eight heights is the numerical face of Lemma 7.4. At cycle lengths 5 and 6 the same complete-orthogonality logic gives the product-model combs of §8 (aggregate faces green at 0.1–2.5%).

B Provenance, research record, and discrepancy register

The research record

The mathematics of this paper was developed in an interactive research program (August 2026) whose working discipline is falsifier-first: every structural claim was registered with a numerical check before adoption, every firing was converted on the forward record, and the complete trail — round logs, memos, claim ledgers, correction history, scripts — is preserved in the program archive that accompanies the reproduction package. The four stages of the chain (§3, §4, §§5–7, §§8–11) were each closed under adversarial audit rounds before the certification audit of §14 froze the chain presented here.

Conversions that shaped this text

Six of the registry’s fired-and-converted checks shaped the mathematics directly: the evaluation-transfer repair (envelopes are never transported to limits; r61); the model-subtraction correction of the reduced covariance (r80); the hybrid size budget (two firings, r88); the τ -frame lesson (fixed-cutoff arc statistics never self-average; the Parseval identity entered here; r90); the incomplete-sum slip of the progression frame (r95, Remark 7.6); and the removal of an unproven sixth-moment cap from an internal ladder rung (the intermediate values quoted in Theorem 1.1 consume exactly what their chains prove). The 29th conversion (r103) is the Lean-draft inverse omission of §6.

Discrepancy register of the certification audit (August 16, 2026)

- D1. Earlier draft of §4, assembly paragraph: printed $\Lambda_2(0) = 0.156293$; the closed form at $\bar{R} = 0.0066$ gives 0.154193, and the theorem constants 0.6916/0.8458 are consistent with the latter. Corrected in §4; theorem unaffected. *Impact on the main theorem: none.*
- D2. An earlier draft attributed the bands 5.573 ± 0.113 , 10.02 ± 0.27 to measurement “on 2×10^7 zeta zeros”; the reproduction package derives exactly these numbers from GUE double-window Richardson extrapolation (`gue_bands.py`, `m5m6_bands.py`), and no zeta-zero dataset of that size exists in the record. This paper attributes the bands to the sine-model/GUE measurement (3.2). If a zeta-zero measurement exists it is not in the record and is not consumed. *Impact: attribution only; no constant changes.*

- D3. Consumption LP band model: correlated vs anti-correlated δ -accounting (0.79454 vs 0.78972); the correlated form is correct and is what `face_lp_full` implements; both on record (§11). *Impact: the anti-correlated variant is doubly conservative; headline unchanged.*
- D4. The m_4+m_5 rung: an earlier stage quoted 0.7295/0.8647; the audit’s independent LP gives 0.73144/0.86572 under the same constraints at bounded support. Superseded by D6: the rung is demoted regardless. *Impact: none on the main chain (the demoted rung is not consumed).*
- D5. Formalization language: earlier program documents described [5] as “Lean-formalized”; that statement concerns the package accompanying [5], of which no compiled artifact exists in this repository (§14(f)). The compiled coverage of this paper is exactly as graded in (f). *Impact: grading language only.*
- D6. The intermediate m_4+m_5 rung (0.7295/0.8647, an earlier stage) is unsound on unbounded support without sixth-moment information: the audit’s LP degrades to 13/18 as the support range grows (0.73144 $\rightarrow$ 0.72257 at $R = 8 \rightarrow 200$), and a dual certificate for a fifth-moment lower bound necessarily consumes x^6 -information. Rung demoted and repaired (Remark 11.8); the main chain is unaffected ($y_6 > 0$ in Lemma 11.4). *Impact: removes an intermediate rung; headline unaffected.*
- D7. The archived four-cycle anchor $t_{\text{adj}} = 0.11717$ (rounds 69–71; quoted in earlier drafts and `m1.suite`) carries a +0.43% rectangle-rule grid bias: the exact value is $\frac{7}{60} = 0.11667$ (Proposition 8.2). Downstream totals are unaffected (the m_5/m_6 rational layers are anchor-free); the saturated-split triple 0.2677/0.0156/−0.0177 becomes exactly $\frac{4}{15}/\frac{1}{60}/-\frac{1}{60}$ (Remark 3.6), in better agreement with the cumulant engine’s independent −0.01663. *Impact: none on the rational layers (anchor-free); corrects the archived anchor.*
- D8. The rational identifications $T_0, \dots, T_3 = \frac{3}{70}, \frac{1}{90}, \frac{1}{180}, \frac{1}{70}$, recorded in the audit ledger as a conjecture and not consumed, are now proven (Proposition 8.2); $\{2, 2, 2\} = \frac{131}{420}$ is consumed exactly and its band contribution retired. *Impact: tightens m_6 ; contributes to the headline.*
- D9. The continuum-quadrature constant $c_0 = -1.4228$ (r60, “computed, not fitted”) is identified in closed form: $c_0 = \gamma - 2$ (Proposition C.1). The audit located a $d = 2$ discrepancy between the archive’s tensor-sieve γ -array and the raw Euler closed form ($\tilde{\gamma}_2 = C_2 - 1$ vs $\gamma_2 = C_2$, shifting the constant by exactly -2); this is the archive’s parity-mean convention (the $q = 2$ arc block booked in $\bar{D}$), not an error — the raw-convention constant is γ , and both conventions are now documented with the conversion. The slope-1 law of the partial sums, previously a measurement (0.99996), is proved exactly. *Impact: N2 partial discharge; main theorem unaffected.*
- D10. The endpoint-grid ladders behind the previous connected constants fired their refinement face: at $b \geq 6$ and coarse steps the endpoint evaluation *sign-flips* ($b = 6$ at $dv = 0.4$: +0.036 against the true −0.010), and the archived adoptions $C_5 = 0.0275(5)$, $\{4, 2\} = -0.0548(10)$, $\{6\} = -0.0100(12)$ inherited a systematic shift. Grouped *midpoint*-grid ladders (§8.6), calibrated exactly against Proposition 8.2 at $b = 4$, settle monotonically and give the corrected $C_5 = 0.0278(1)$, $\{4, 2\} = -0.0552(8)$, $\{6\} = -0.0078(8)$, since certified exactly (§8.5) (registry’s 32nd conversion; the 31st is D7’s anchor correction,

- independently forced by the same audit). *Impact: direct — these are the constants the headline consumes; the theorem value moved $0.7946 \rightarrow 0.7947$ under the correction.*
- D11. The Lean draft files claimed core-decidability of their $\mathbb{Q}$ -statements (“`decide on  $\mathbb{Q}$` ”) and used `ring`; the maintainers’ compile showed both unavailable in core Lean 4.33 (`ring` is `mathlib`; `instDecidableEqRat` does not kernel-reduce). Five proofs repaired with `grind`, statements unchanged; build then succeeded (registry’s 33rd conversion; §14(f)). *Impact: proofs only; statements unchanged.*
- D12. The second compile fired on `cert_nonneg`: `grind` assigns the square $(\cdot)^2$ as an opaque atom (its diagnostics guess the atom’s value) and cannot derive perfect-square nonnegativity; repaired with core `Lean.Grind.OrderedRing.sq_nonneg` (registry’s 34th conversion). The audit also found the superseded seed `SOSCertificate.lean` silently outside the build roots while declaring identically named headline theorems at the retired constants; the file is removed from the package (content absorbed; the draft is preserved in the program archive) and the build now covers every `.lean` file in the directory. *Impact: proofs and packaging only; statements unchanged.*
- D13. The citation relation to [5] was reversed twice on the forward record: an intermediate revision absorbed the framework into the text with proofs (when [5] was not yet publicly posted), and the present revision restores [5] as the cited foundation and removes the duplicated proofs, keeping statements only (§2). *Impact: attribution and exposition; no mathematical content changed.*
- D14. A spectral identification of the connected constants as cumulants of the sharp-window counting statistic (prolate spectrum route) was falsified by its pre-registered gate before adoption ($\kappa_2 = 0.3442$ vs $\frac{1}{3}$; $\kappa_3 \neq 0$) — the 35th conversion. The corrected structure (the chain-graph family over the partition-cyclic terms, with the $b = 2$ dictionary closing exactly at $1 - \int \text{sinc}^4 = \frac{1}{3}$) informs the accelerated engine’s design. *Impact: none on constants (caught pre-adoption); structural understanding gained.*
- D15. The identification run: the accelerated ladder engine (equivalent-transform speedups only: exact-zero support pruning, term-trie sharing, kernel fusion, negation symmetry; validated against the reference engine on five gates and bitwise across three architectures) extended every ladder to seven Romberg rungs and identified $C_5 = \frac{1}{36}$, $\{6\} = -\frac{1}{126}$ (and $C_7 = -\frac{17}{360}$ one rung weaker), all inside the previous bands. The paper’s constants move $0.0278(1) \rightarrow \frac{1}{36}$, $-0.0078(8) \rightarrow -\frac{1}{126}$ and the theorem $0.7947 \rightarrow 0.7957$. *Impact: direct — band collapse by eleven orders of magnitude on two of the three consumed constants.*
- D16. The $\{4, 2\}$ spectator-pair variant (specified with three coarse reference rungs and the pre-registered candidate $-\frac{23}{420}$) was ported to the accelerated engine, passed its four gates (C_4 mechanics; bitwise reference agreement; archive cross-check $-0.0588/-0.0558$; band), and its seven-rung ladder identified $\{4, 2\} = -\frac{23}{420}$, hitting the pre-registered candidate, with components $-\frac{2}{315}$, $-\frac{1}{1260}$, $-\frac{1}{252}$. All three N1 constants are now identified; the theorem moves $0.7957 \rightarrow 0.7960$. *Impact: direct — the last constant band is retired; only the transport allowance remains.*
- D17. The residual m_6 allowance (0.0002), formerly labelled a frozen-slot transport allowance, is traced in the audit ledger to the adopted band of the mixed odd-block class $\{3, 3\}$; the

class is written out as vanishing (identically zero on every grid under both protocols, all placements — the model face of the arithmetic odd-block vanishing of §8) and the allowance is retired. M_6 becomes the exact rational $\frac{12809}{1260}$ and the theorem moves $0.7960 \rightarrow 0.7962$. *Impact: direct — the consumption inputs are now exact rationals end to end.*

- D18. The symbolic certification of N1: the direct 4D/5D generalization of the piecewise integrator closed $\{4, 2\}$ term by term but projected to centuries per term at $b = 6$ (multiplicative piece-count growth); the polytope lift of §8.5 — each term one rational polytope volume, exact facet recursion, a $\sim 10^9$ -fold reduction at $b = 6$ — evaluated all 1310 terms as exact fractions, summing to $\frac{1}{36}$, $-\frac{23}{420}$, $-\frac{1}{126}$ exactly. All gates passed ($\Phi_4 = -\frac{1}{60}$ by both methods; 3D regression matching Proposition 8.2; per-term dual-method agreement; componentwise $\{4, 2\}$ recombination); the pre-registered falsifier F-SYM-N1 did not fire; symbolic and numerical values agree bitwise. *Impact: grades only — N1 moves from identification to proof; the constants and the theorem are unchanged.*
- D19. The $\{2, 2, 2\}$ classification correction (the 38th conversion; Remark 8.3): the consumed $\frac{131}{420}$ classified the 15 hexagon matchings by crossing number; the nested crossing-zero class had never been separately evaluated and equals $\frac{17}{420} \neq \frac{3}{70}$ (three independent exact paths). Corrected bundle $\{2, 2, 2\} = \frac{32}{105}$; $M_6 = \frac{640}{63}$. Caught by the pre-registered $b = 6$ calibration gate of the eighth-moment pairing run; independently detected by two model-side estimators ($55\text{--}110\sigma$ and $11\text{--}21\sigma$ misses against the old value, $\sim 1\sigma$ against the new). *Impact: direct and favourable — the upper constraint M_6 decreases and the theorem rises ($0.7962 \rightarrow \frac{2025}{2519}$); the published $0.7962/0.8981$ consumed a weaker-but-true bound and remains valid. Lesson: crossing number is not a complete invariant of cyclic chord diagrams; every bundle is now evaluated by dihedral placement class.*
- D20. A ledger-enumerator deduplication error (the 39th conversion): a head-fixing quotient valid only between equal-size blocks was applied across unequal sizes, undercounting $\{5, 2\}$ ($21 \rightarrow 15$) and others; caught before any value was computed by the invariant *orbit size divides group order* ($5 \nmid |D_7| = 14$), now emitted with every enumeration report. A second, independent enumerator confirms all inventories. *Impact: none on constants (caught pre-computation); the check is institutionalized.*
- D21. Ladder-band calibration (the 40th conversion): the pre-registered h^4 acceptance band for the $\{2^4\}$ ladder (3×10^{-7}) was optimistic by a factor 3 (realized deviation 8.1×10^{-7} against the exact values); the criterion is revised to 3×10^{-6} for all future F-V gates. In the same audit the level-6 denominator law (2520) was found to fail at level 8 ($3780 \nmid 2520$), explaining the failure of small-denominator reconstruction there. *Impact: verification-calibration only; all 17 exact values confirmed within the revised band.*
- D22. The C_8 candidate miss (the 41st conversion): the pre-registered $C_8 = \frac{7}{180}$ (from small-denominator reconstruction of the model-side Σ_8) was missed by the exact run, $C_8 = \frac{157}{4032}$, deviation 4.96×10^{-5} — inside the tolerance gate, outside the candidate. The corner moved and the certificate re-optimized in closed form ($\frac{17230713}{20284508} \rightarrow \frac{138058329}{162540559}$, i.e. $0.849452 \rightarrow 0.849378$); the companion C_7 run hit its registered value exactly. *Impact: headline -7×10^{-5} ; lesson — at level 8 small-denominator reconstruction is unreliable (cf. D21), and only symbolic runs settle candidates.*

- D23. A claimed flat-window minimality lemma for Σ_2 was falsified by its own first-order probe (mild tapers reduce m_2) before entering any draft; λ -optimality of the flat window stands as an empirical statement only. *Lesson: optimality claims get a variational probe before a name.*
- D24. Truncated decimals had been written as rounded in eight places ($1-2\lambda_7 = 0.910460410551\dots$ typeset as “0.910460411...”, making the displayed inequality false at that precision); corrected to truncation with explicit ellipsis throughout. *Lesson: quoted decimals declare rounding or truncation; inequalities are safe only in the truncation direction.*
- D25. A campaign memo misstated the $b = 14$ table extent (11 points, surplus 3) where the archive holds 9 points ($N=2..10$), surplus 1; corrected forward, conclusions unaffected. Separately, an over-merged symmetry fold in the lifted-system scan (the reflection permutes blocks, hence must permute the c -columns) was caught by count reconciliation against the independent enumeration (11/61 vs the true 13/73) and repaired by dropping the fold. *Lesson: symmetry deduplication must be re-proved whenever coordinates are extended; falsification scans may overcount, never undercount.*
- D26. A cost calibration for the $b = 7$ lifted-system scan underestimated by a factor of seven: the probe took the *first* three matrices of the enumeration order, which front-loads structurally simple partitions; separately, a dispatch reported 100 cores while launching 24. Re-sharded and re-run. *Lesson: calibration probes must be randomly sampled — “the first N ” is not a sample — and reported core counts must be reconciled against the actual dispatch.*
- D27. A compute order specified an assembly gate “with the Lemma 10.2 coefficients” without an executable formula; the compute side’s query caught the gap before any wrong result. The exact coefficient law (merge partitions, cyclic class orders, sign $\prod_B (-1)^{|Q_B|-1}$, Fubini term counts 75/541/4683 at $b = 4/5/6$) was then derived, verified bitwise on both sides, and recorded in Lemma 10.2. *Lesson: “with the standard coefficients” is not a specification; every assembly order carries an explicit formula and pre-registered self-check values.*
- D28. During final packaging a receipt table listed a $b = 5$ lifted-system scan as PASS while its artifact file was absent from the package (the run predated a session switch and its output was never written to disk); the scan was re-run and archived before acceptance. *Lesson: a number in a log or a receipt is not evidence; every table row must resolve to an archived artifact — and this discipline must be enforced against receipts themselves.*

The registry’s 30th conversion occurred during the certification audit: the first version of the W1 resolved-comb face used raw multiples of b_1b_2 as the on-progression and fired; the corrected face implements the per-coordinate CRT classes (§8.4). Forward-recorded in `writeouts_verification.py`. The 31st–35th are D7, D10, D11, D12 and D14 above.

Reproduction

The complete package — paper, reproduction pipelines, certification layer, GPU engines, and the Lean sources with their build log — is public at

<https://github.com/JoshuaHKU/zeta-density-one-reproduction>.

A consolidated reproduction checklist accompanies this paper (`REPRODUCTION.md`): environment, one command per suite, expected outputs with the August 16, 2026 rerun values, and the

mapping from every quoted constant to the script that regenerates it. The present revision’s computational assets are organized as a self-contained package (**repro/**: engines, per-constant orbit inventories and values under content-canonical keys with per-value provenance and second-path fields, measurement raw data, one executable script per falsifier gate with asserted match counts, and an append-only log archive; specification in **compute/REPRO_SPEC.md**). The package is frozen as v1: thirteen constants with per-orbit records, 126 files under a hashed manifest, eleven gate scripts all green, and a recorded two-sided acceptance audit (**repro/ACCEPTANCE_r141.md**); the subsequent freezes v3 (trace route) and v4 (branch-equality campaign) extend this append-only to the present suite of eighteen gates (**repro/REPRO_V4_FREEZE.md**). Total runtime of the full set is under fifteen minutes on one CPU (the session package’s optional long gates under thirty; the exact $\{2, 2, 2\}$ certification under two hours; the N1 polytope certification: $\{4, 2\}$ and C_5 in minutes, $\{6\}$ in 9.1 single-core hours, term-parallel).

C The exact certification layer

A1. The exact-integration algorithm

Every pairing-layer constant is an integral $\int_{[-1,1]^d} (1 - \text{spread}(p_1, \dots, p_b))_+ \prod_i |x_i| dx$ where the walk positions p_j are $\{0, 1\}$ -combinations of the chord variables and 0 is a position (so the support forces $|x_i| \leq 1$ and the $\min(|\cdot|, 1)$ caps of the C_2 factors never bind). The integrand is piecewise polynomial; its kink set is contained in the hyperplane family

$$\mathcal{H} = \{p_i - p_j = c\}_{i < j, c \in \{0, \pm 1\}} \cup \{x_i = 0, \pm 1\},$$

all with coefficients in $\{-1, 0, 1\}$ and unit leading coefficients in the innermost variable. Iterated integration: the inner integral is computed exactly between the enumerated inner breakpoints; because leading coefficients are units, the breakpoints are themselves small-integer linear forms in the outer variables, and the piece structure of a partial integral changes only on (i) direct kinks, (ii) collisions of two inner breakpoints, (iii) breakpoint–boundary collisions — an explicitly enumerated resultant family at each level. On every kink-free piece the (partial) integrand is a polynomial of degree at most $2/4/8$ in the inner/middle/outer variable, integrated exactly by closed Newton–Cotes rules of order $5/7/9$ at rational nodes, in exact rational arithmetic throughout. *Self-check*: each piece is integrated by the same rule on the whole piece and on its two halves; for a polynomial of degree below the rule order both are the exact integral, so the two rationals must be *identical*, and a kink missed by the enumeration forces a discrepancy on the containing piece. No discrepancy occurred at any level of any run. Independent cross-checks: closed-form orthant evaluations (A2), and the float grid ladders of the archive (10^{-6} -agreement with the certified rationals).

A2. Closed-form cross-checks

T_0 (positions $\{0, v, w, u\}$): in the all-positive orthant $\int_{[0,1]^3} (1 - \max(v, w, u))vwu = \frac{1}{56}$ (order and integrate); a one-negative orthant gives $\int_0^1 t(1-t)^5/20 dt = \frac{1}{840}$; summing the orbit, $T_0 = 2 \cdot \frac{1}{56} + 6 \cdot \frac{1}{840} = \frac{3}{70}$. t_{opp} (positions $\{0, v, v+w, w\}$): the same decomposition gives $\frac{1}{30}$ exactly, confirming the archived exact anchor. t_{adj} : the two same-sign quadrants give $\frac{1}{20}$ each and the two mixed quadrants $\frac{1}{120}$ each: $t_{\text{adj}} = \frac{14}{120} = \frac{7}{60}$ (not the archived grid value 0.11717; register D7).

A3. The certified values

$$t_{\text{adj}} = \frac{7}{60}, \quad t_{\text{opp}} = \frac{1}{30}, \quad T_0 = \frac{3}{70}, \quad T_1 = \frac{1}{90}, \quad T_2 = \frac{1}{180}, \quad T_3 = \frac{1}{70}, \quad T'_0 = \frac{17}{420} \text{ (nested; D19),}$$

$$\{2, 2, 2\} = 2T_0 + 3T'_0 + 6T_1 + 3T_2 + T_3 = \frac{32}{105},$$

with $\Phi_4 = -\frac{1}{60}$, saturated $\{2, 2\} = \frac{4}{15}$, plateau $= \frac{1}{60}$, connected model $= -\frac{1}{60}$ (Remark 3.6). Code: `certification/exact_t222.py` (classes run separately; dual-scheme identical-rational check at every piece).

A4. The continuum-quadrature constant in closed form (N2, partial discharge)

Proposition C.1 ($c_0 = \gamma - 2$). *Let γ_d denote the free coefficients of the partial-sum identity (Lemma 4.7) at the (1, 1) class. Then γ_d has the exact Euler closed form*

$$\gamma_{2m} = C_2 \prod_{p|m} \frac{1}{p(p-2)} \quad (m \text{ odd squarefree}), \quad \gamma_d = 0 \text{ otherwise,}$$

and, in the raw convention, $\sum_{d \leq M} d\gamma_d = \log M + \gamma + o(1)$; in the archive convention (the $q = 2$ arc block booked inside the parity mean, $\tilde{\gamma}_2 = \gamma_2 - \beta_2 = C_2 - 1$),

$$c_0 = \gamma - 2 = -1.4227843 \dots$$

Proof. (i) The closed form: from the arc identity (Proposition 4.2) at $b_1 = b_2 = 1$, $\beta_q = \mu(q)^2/\varphi(q)^2$, and $\gamma_d = \sum_e \mu(e)\beta_{de}$ factorizes over primes with local weights $f_0(p) = 1 - (p-1)^{-2}$, $f_1(p) = (p-1)^{-2}$, $f_{v \geq 2} = 0$; $f_0(2) = 0$ forces $2 \parallel d$ and the product assembles to the display (C_2 the twin-prime constant). (ii) The constant: $\sum_{d \leq M} d\gamma_d = 2C_2 \sum_{m \leq M/2} \prod_{p|m} \frac{1}{p-2}$, whose Dirichlet series is $\zeta(s+1)E(s)$ with $E(0) = \prod_p (1 - \frac{1}{p}) \prod_{p>2} (1 + \frac{1}{p-2}) = \frac{1}{2C_2}$ (slope $2C_2E(0) = 1$ exactly) and $E'(0)/E(0) = \sum_p \frac{\log p}{p-1} - \sum_{p>2} \frac{\log p}{p-1} = \log 2$ (the p -sums cancel pairwise except $p = 2$); the double-pole residue gives $2C_2[E(0)(\log \frac{M}{2} + \gamma) + E'(0)] = \log M + \gamma$. (iii) The convention shift: the archive tensor sieve (`tail.bound`) and the closed form agree at *every* d except $d = 2$, where the archived class- d replacement books $\tilde{\gamma}_2 = C_2 - 1$ (the $q = 2$ block moved to the parity mean); this shifts $\sum d\gamma_d$ by exactly -2 . Numerical verification (`n2.c0.certification.py`): the raw sum converges to γ with deviations $9.4 \times 10^{-4}/6.3 \times 10^{-5}/9.2 \times 10^{-6}$ at $M = 10^4/10^5/2 \times 10^7$; the archived measurement -1.42279 at $M = 1.6 \times 10^7$ matches $\gamma - 2$ to 10^{-5} ; the partial-sum identity closes with its $d > M$ tail (LHS $-$ RHS $-$ tail $= -0.010$ at $M = 2 \times 10^5$, $D = 2 \times 10^7$, consistent with the remaining truncation). $\square$

Consequently the “computed constants” of Lemma 4.10 reduce to: slope $= 1$ (proved), $c_0 = \gamma - 2$ (closed form), the certified geometric tail T_Γ (archive r60, ratio 0.500), and the single remaining computed piece $\overline{\text{osc}}$ inside c^* . N2 is thereby partially discharged; what remains is the $\overline{\text{osc}}$ identification (or certified enclosure) and the universality-collapse rate write-out.

A5. The N1 certificate family and its payoff curve

The exact dual certificate re-optimizes against any hypothetical certified enclosures $C_5 \in [c_-, c_+]$, $\{4, 2\} \leq u_{42}$, $\{6\} \leq u_6$ (correlated accounting: the single unknown C_5 enters m_5 and m_6 together; the worst case is a corner of the affine family). Exact rational certificates (`certificate_family.py`, atoms re-optimized per scenario):

scenario (graded by the information consumed)	$1 - 2w_0 \geq$	$1 - w_0 \geq$
$k \leq 6$ certified exactly (Corollary 11.5)	2025/2519	2272/2519
$k = 8$ layer at the stated grades (Lemma 11.6)	0.84937	0.92468
enclosure width 0.01 (historical)	0.77418	0.88709
enclosure width 0.04 (historical)	0.74299	0.87149
signs only: $C_5 \geq 0, \{4, 2\} \leq 0, \{6\} \leq 0$	0.74031	0.87016
no connected information	13/18	31/36

Sensitivity of the shipped certificate (exact): $\partial(\text{headline})/\partial c_- = +7.32$, $\partial/\partial c_+ = -5.61$, $\partial/\partial(u_{42}+u_6) = -0.93$. In particular a *sign-only* certification of the three connected constants — far coarser than an enclosure — already certifies 0.740/0.870, and the exact computation — now carried out, landing on the top row — realizes the full payoff. (The level-by-level cost model that placed the direct piecewise runs at cluster scale is what the polytope lift of §8.5 circumvented: the certified runs completed on a single workstation core in hours.)

A6. The exact dual certificates and the provenance of the historical atoms

The degree-6 certificate of Lemma 11.4 is $P = (Q_3^*)^2$ with the rational kernel coefficients displayed there; its corner evaluation $\lambda_3 = \frac{247}{2519}$ and the degree-8 corner $\lambda_4 = \frac{12241115}{162540559}$ of Lemma 11.6 are finite rational computations (`certification/certify84.py`; Lean queued). The numerator of Q_3^* is proportional to $1932x^3 - 7368x^2 + 8232x - 2519$, whose roots $0.50151\dots$, $1.27817\dots$, $2.03397\dots$ are the certificate “atoms” of the Christoffel extremal measure — never needed by the proof. *Provenance of the historical atoms*: at the superseded $M_6 = \frac{12809}{1260}$ the same construction returns roots $0.53234\dots$, $1.31220\dots$, $2.05855\dots$, whose four-place rationalizations are exactly the atoms $(\frac{5323}{10000}, \frac{6561}{5000}, \frac{10293}{5000})$ of the earlier revision — the LP extraction had been finding the kernel-polynomial roots. For the record, that earlier certificate expands (numerator, monic) as

$$x^6 - \frac{39031}{5000}x^5 + \frac{2422533313}{100000000}x^4 - \frac{4746141355593}{125000000000}x^3 + \frac{39293368568783721}{125000000000000}x^2 - \frac{20200560698400422913}{156250000000000000}x + (abc)^2,$$

with $(abc)^2 = \frac{129222147277358919747441}{625000000000000000000}$, so $y_0 = 1$, $y_5 = -\frac{39031}{5000}(abc)^{-2} < 0$, $y_6 = (abc)^{-2} > 0$. At the superseded corner it gave $w_0 = \frac{829278553005924403328783}{8140995278473611944088783}$ and the published headlines $\frac{7962}{10000}/\frac{8981}{10000}$ — valid then and now (it consumed a weaker-but-true M_6), superseded by the exact optima of Lemma 11.4 and Lemma 11.6. Under the Christoffel formulation the solver, grid, atom extraction and rationalization steps no longer exist in any form; the entire consumption layer is the rational linear algebra of Lemma 11.1. Verification: `certification/certify84.py`; Lean: `lean/RhGate/Certificate.lean` (compiled, previous corner) and `Certificate84.lean` (queued): expansion identities, nonnegativity, corner evaluations, sign pattern, headlines; zero sorrys.

References

- [1] N. I. Akhiezer, *The Classical Moment Problem and Some Related Questions in Analysis*, Oliver & Boyd, 1965.
- [2] S. A. C. Baluyot, D. A. Goldston, A. I. Suriajaya, and C. L. Turnage-Butterbaugh, *Pair correlation of zeros of the Riemann zeta function I: proportions of simple zeros and critical zeros*, arXiv:2501.14545 (2025).

- [3] D. A. Goldston and A. I. Suriajaya, *Zeta zeros on the critical line*, arXiv:2511.20059 (2025).
- [4] H. M. Bui, J. B. Conrey, and M. P. Young, *More than 41% of the zeros of the zeta function are on the critical line*, Acta Arith. **150** (2011), 35–64.
- [5] Claude (Anthropic), *More than two thirds of the zeros of the Riemann zeta function lie on the critical line*, preprint, August 2026, <https://www-cdn.anthropic.com/564f962e60643842f5fcb4a17c9dbc8f608f1c37.pdf>.
- [6] J. B. Conrey, *More than two fifths of the zeros of the Riemann zeta function are on the critical line*, J. reine angew. Math. **399** (1989), 1–26.
- [7] H. Davenport, *Multiplicative Number Theory*, 3rd ed., Springer GTM 74, 2000.
- [8] S. Feng, *Zeros of the Riemann zeta function on the critical line*, J. Number Theory **132** (2012), 511–542.
- [9] H. Halberstam and H.-E. Richert, *Sieve Methods*, Academic Press, 1974.
- [10] H. Iwaniec and E. Kowalski, *Analytic Number Theory*, AMS Colloquium Publ. 53, 2004.
- [11] N. Levinson, *More than one third of zeros of Riemann’s zeta-function are on $\sigma = 1/2$* , Adv. Math. **13** (1974), 383–436.
- [12] K. Matomäki, M. Radziwiłł, and T. Tao, *Correlations of the von Mangoldt and higher divisor functions I. Long shift ranges*, Proc. London Math. Soc. (3) **118** (2019), 284–350.
- [13] H. L. Montgomery, *The pair correlation of zeros of the zeta function*, Proc. Sympos. Pure Math. **24** (1973), 181–193.
- [14] H. L. Montgomery, *Distribution of the zeros of the Riemann zeta function*, Proc. ICM Vancouver 1974, vol. 1, 379–381.
- [15] K. Pratt, N. Robles, A. Zaharescu, and D. Zeindler, *More than five-twelfths of the zeros of ζ are on the critical line*, Res. Math. Sci. **7** (2020), no. 2.
- [16] B. Simon, *The classical moment problem as a self-adjoint finite difference operator*, Adv. Math. **137** (1998), 82–203.
- [17] A. Soshnikov, *Determinantal random point fields*, Russian Math. Surveys **55** (2000), 923–975.
- [18] A. Selberg, *On the zeros of Riemann’s zeta-function*, Skr. Norske Vid.-Akad. Oslo I (1942), no. 10.
- [19] E. C. Titchmarsh, *The Theory of the Riemann Zeta-Function*, 2nd ed. (revised by D. R. Heath-Brown), Oxford University Press, 1986.
- [20] A. Weil, *Sur les “formules explicites” de la théorie des nombres premiers*, Comm. Sémin. Math. Univ. Lund (1952), 252–265.